


# CSCI 5451: Introduction to Parallel Computing

## Lecture 3: Parallel Architectures (cont'd)



UNIVERSITY OF MINNESOTA  
**Driven to Discover®**

# Lecture Overview

- ❑ Recap
- ❑ Multi-Stage Networks (in-depth)
- ❑ Network Topology
- ❑ Memory & Moving Bits
  - Memory Basics (Bandwidth, Caching, Latency, Prefetching, etc.)
  - Cache Coherence



# Lecture Overview

## ❑ **Recap**

❑ Multi-Stage Networks (in-depth)


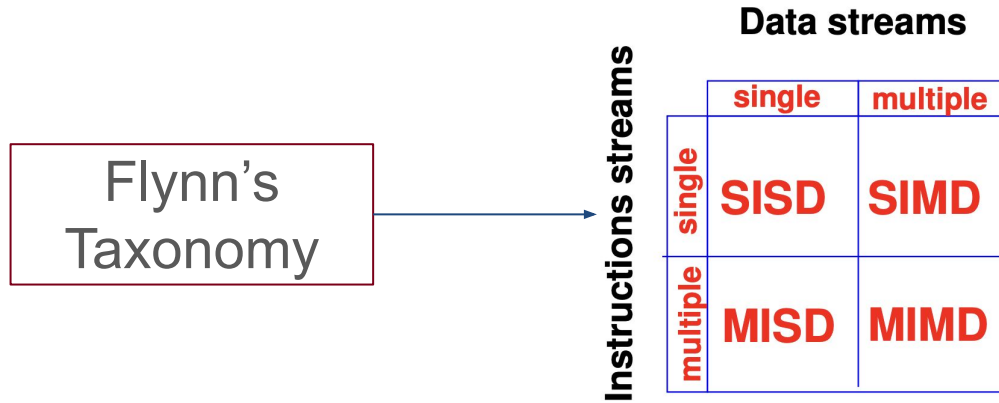
❑ Network Topology

❑ Memory & Moving Bits


- Memory Basics (Bandwidth, Caching, Latency, Prefetching, etc.)
- Cache Coherence



# Logical Parallel Hardware Models



Source: Kaminsky/Parallel Java




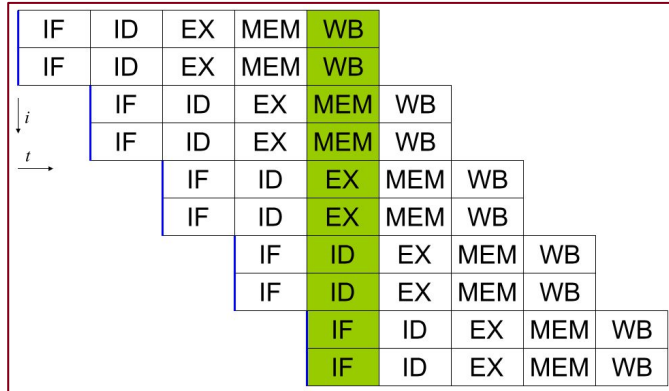
Source: Kaminsky/Parallel Java



# Superscalar Execution (Implicit Parallelism)

Recap

Pipelining



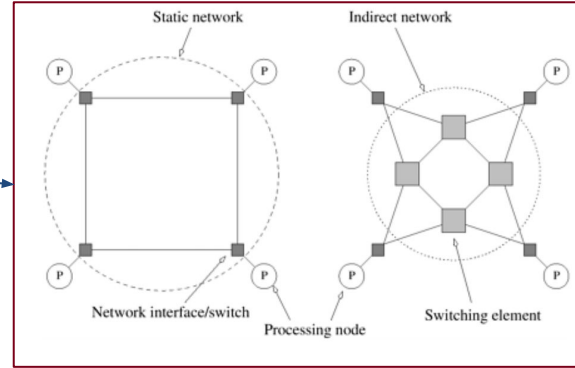
2-Way Superscalar Execution




# Network Topology Choices

Recap


Static vs. Dynamic



Bus



Crossbar



# Lecture Overview

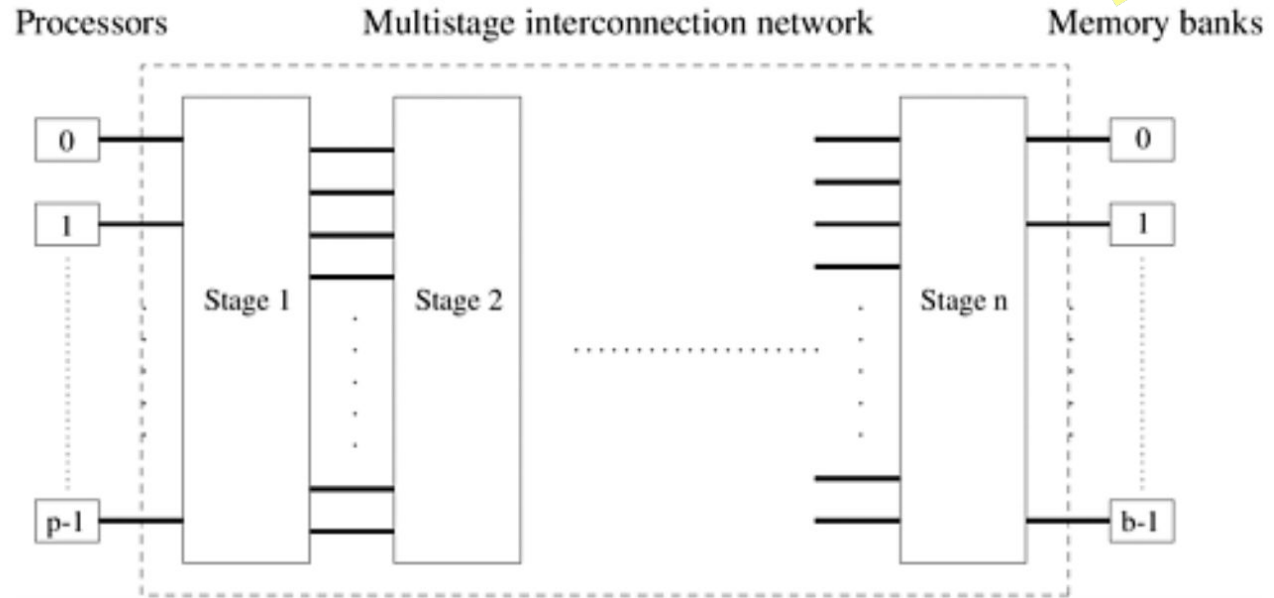
- ❑ Recap
- ❑ **Multi-Stage Networks (in-depth)**
- ❑ Network Topology
- ❑ Memory & Moving Bits
  - Memory Basics (Bandwidth, Caching, Latency, Prefetching, etc.)
  - Cache Coherence



# Multi-Stage Networks

Recap

- ❑ Shuffles connections at each stage
- ❑ Worse than Bus on price, better than Crossbar
- ❑ Worse than Crossbar on scalable performance, better than Bus
- ❑ Uses  $\log(p)$  stages






# Omega Network

$p$  processors

Processors




# Omega Network

Memory Banks

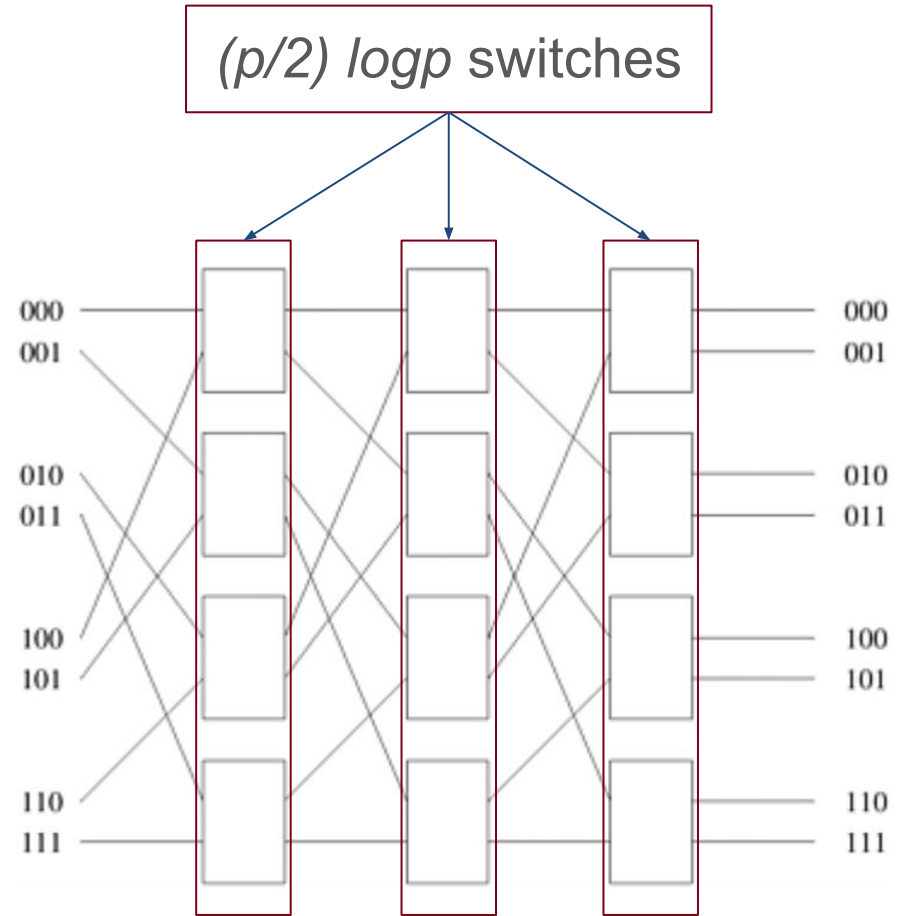
We make the simplifying assumption that the number of banks is equal to the number of processors

$b$  banks



# Omega Network



Switches



# Omega Network

## Switch Structure

Switches have 2 inputs, 2 outputs → Inputs can be routed in either direction of the outputs




# Omega Network

## Physical Link Structure

For all but the final layer of links, connecting links follows the below function (where each switch contains 2 inputs).  $j$  is the output index and  $i$  is the input


$$j = \begin{cases} 2i, & 0 \leq i \leq p/2 - 1 \\ 2i + 1 - p, & p/2 \leq i \leq p - 1 \end{cases}$$



# Omega Network


This interconnection pattern is also called a

***Perfect Shuffle***



# Omega Network (Example)

Note: Only one possible path for A or B to go from their respective processors to the appropriate memory banks



In Class Example:  
4 processes to 4 banks Omega Network





# Lecture Overview


- ❑ Recap
- ❑ Multi-Stage Networks (in-depth)
- ❑ **Network Topology**
- ❑ Memory & Moving Bits
  - Memory Basics (Bandwidth, Caching, Latency, Prefetching, etc.)
  - Cache Coherence




How to connect processors  
to one another?



# Processor Connections



Fully Connected




Star

# Processor Connections




Linear Array  
w/o Wraparound




Linear Array  
w/ Wraparound


# Processor Connections



2-d Mesh  
w/o Wraparound



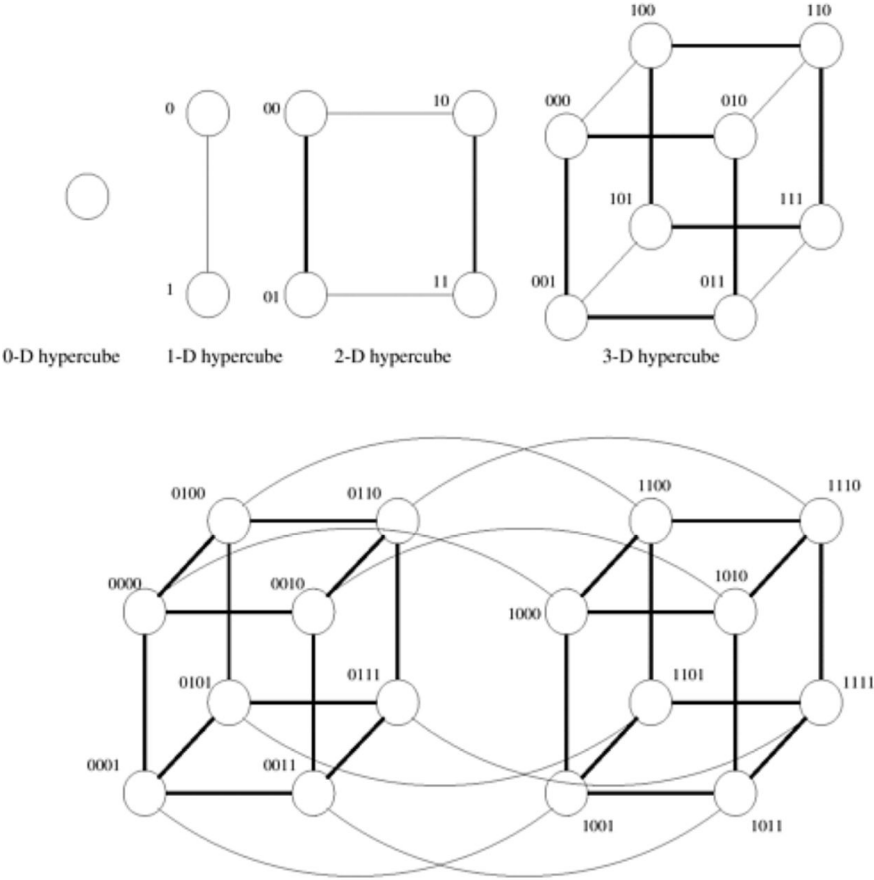
2-d Mesh  
w/ Wraparound  
(2-d torus)



3-d Mesh  
w/o Wraparound


# Processor Connections

## Hypercube Construction



# Processor Connections

Fat Tree



# Processor Connections in the Limit

- ❑ What are the limits of a parallel computer? How complex do networks get in reality?
- ❑ Organizations (Companies/Governments)
  - Modern Supercomputers (Frontier (DoE), Aurora (DoE), Hyperion (Meta), Dojo (Tesla)) will not fit cleanly into the above categories
  - Nodes/switches/links/memory will be added/removed over time
  - Energy/Application requirements may change
- ❑ Distributed Computing (SETI@Home, PrimeIntellect, etc.)
  - Perform jobs at massive scale with consumer hardware
  - Users can volunteer their own hardware
  - Topologies are wildly different and change incredibly frequently






# Network Evaluation

- ❑ Diameter → The maximum distance between any two nodes in the network, where the distance between any two nodes is defined as the shortest path between them
- ❑ Connectivity → Minimum number of links which must be cut to split the network in two
- ❑ Bisection width → Minimum number of links which must be cut to ensure that the network is partitioned into two equal halves.
- ❑ Cost (no of links) → How many links are required for this topology?



# Example : Fully Connected

Network	Diameter	Bisection Width	Arc Connectivity	Cost (No. of links)
Completely-connected	1	$p^2/4$	$p - 1$	$p(p - 1)/2$




# Example Linear Array w/o Wraparound

Network	Diameter	Bisection Width	Arc Connectivity	Cost (No. of links)
Linear array	$p - 1$	1	1	$p - 1$




# Example: 2-d Mesh w/o Wraparound

Network	Diameter	Bisection Width	Arc Connectivity	Cost (No. of links)
2-D mesh, no wraparound	$2(\sqrt{p} - 1)$	$\sqrt{p}$	2	$2(p - \sqrt{p})$




# Example: 2-d Mesh w/Wraparound

Network	Diameter	Bisection Width	Arc Connectivity	Cost (No. of links)
2-D wraparound mesh	$2\lfloor\sqrt{p}/2\rfloor$	$2\sqrt{p}$	4	$2p$



# Example: Hypercube

Network	Diameter	Bisection Width	Arc Connectivity	Cost (No. of links)
Hypercube	$\log p$	$p/2$	$\log p$	$(p \log p)/2$



# Static Network Evaluations

❑ Feel free to explore further on your own

**Table 2.1** A summary of the characteristics of various static network topologies connecting  $p$  nodes.

Network	Diameter	Bisection Width	Arc Connectivity	Cost (No. of links)
Completely-connected	1	$p^2/4$	$p - 1$	$p(p - 1)/2$
Star	2	1	1	$p - 1$
Complete binary tree	$2 \log((p + 1)/2)$	1	1	$p - 1$
Linear array	$p - 1$	1	1	$p - 1$
2-D mesh, no wraparound	$2(\sqrt{p} - 1)$	$\sqrt{p}$	2	$2(p - \sqrt{p})$
2-D wraparound mesh	$2\lfloor \sqrt{p}/2 \rfloor$	$2\sqrt{p}$	4	$2p$
Hypercube	$\log p$	$p/2$	$\log p$	$(p \log p)/2$
Wraparound $k$ -ary $d$ -cube	$d\lfloor k/2 \rfloor$	$2k^{d-1}$	$2d$	$dp$



# Lecture Overview

- ❑ Recap
- ❑ Multi-Stage Networks (in-depth)
- ❑ Network Topology
- ❑ **Memory & Moving Bits**
  - **Memory Basics (Bandwidth, Caching, Latency, Prefetching, etc.)**
  - Cache Coherence






With processors connected, how can we most intelligently handle memory?




# Bandwidth

- ❑ The rate data can be pumped from memory to the processor
- ❑ Analogous to the number of lanes on a highway which has no traffic
- ❑ Is not increasing nearly as fast as processor speeds → Many bottlenecks in computing are memory-bound now



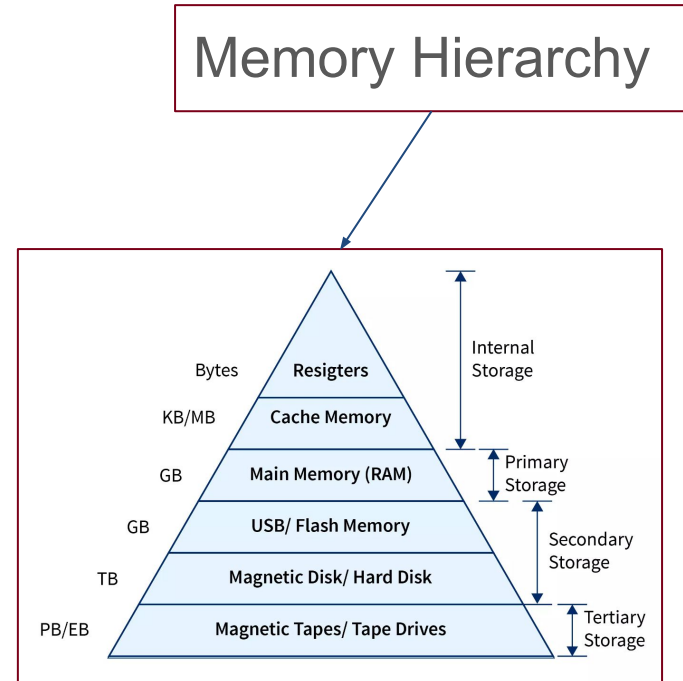
# Latency

- ❑ How long it takes for a memory request to actually retrieve bits from memory
- ❑ Analogous to the time required to travel on a highway with no traffic between two exact points



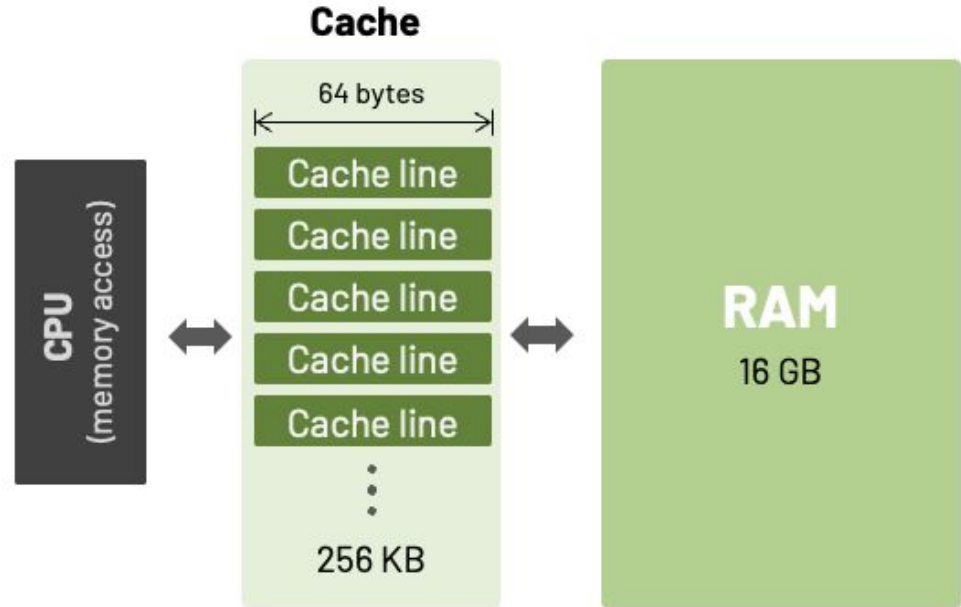
# Caching

- ❑ Processor speed + DRAM speed mismatches have led to increasing use & complexity of caching
- ❑ Place smaller + faster memory closer/on the Processor
- ❑ Low-latency + high bandwidth storage
- ❑ Sits in between processor and main memory
- ❑ Typically multiple levels (L1, L2, L3)



# Cache Lines

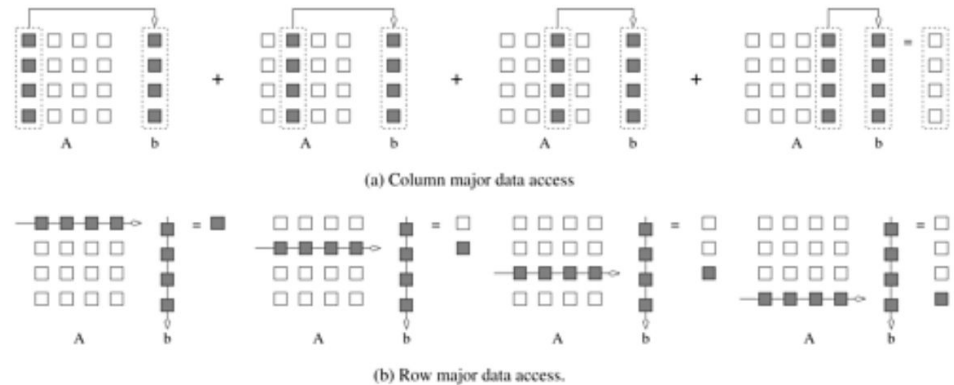
- ❑ Caches do not load single elements from memory
- ❑ They load *lines* from memory
- ❑ Typical lines are the same size across each level of the hierarchy
- ❑ Sizes ~64Bytes
- ❑ Loading a variable means pulling the whole line that variable is on from memory



# Exploiting Cache


The main method for exploiting the cache is to increase the hit ratio (i.e. The proportion of time)

- Temporal Locality → Repeatedly reference same data element multiple time in a small window
- Spatial Locality → Repeatedly reference nearby elements in memory (makes use of cache lines)



# Hiding Latency

- ❑ Prefetching - If we know in advance what we are going to be loading, then load it into cache earlier.
- ❑ Multithreading - Parallelize memory accesses (more feasible today as DRAM today usually permits a relatively large number of concurrent reads)



# Lecture Overview

- ❑ Recap
- ❑ Multi-Stage Networks (in-depth)
- ❑ Network Topology
- ❑ **Memory & Moving Bits**
  - Memory Basics (Bandwidth, Caching, Latency, Prefetching, etc.)
  - **Cache Coherence**






How do we manage states in shared memory systems so that each processes' cache retains the same value?




# Cache Coherence

- ❑ There are three copies across three locations.
  - 1 in main memory
  - 1 in P0 cache
  - 1 in P1 cache
- ❑ They must all stay the same in the course of program execution.
- ❑ We ignore multiple levels of cache in this simple case here
- ❑ The problem is made more complex with multi-level cache



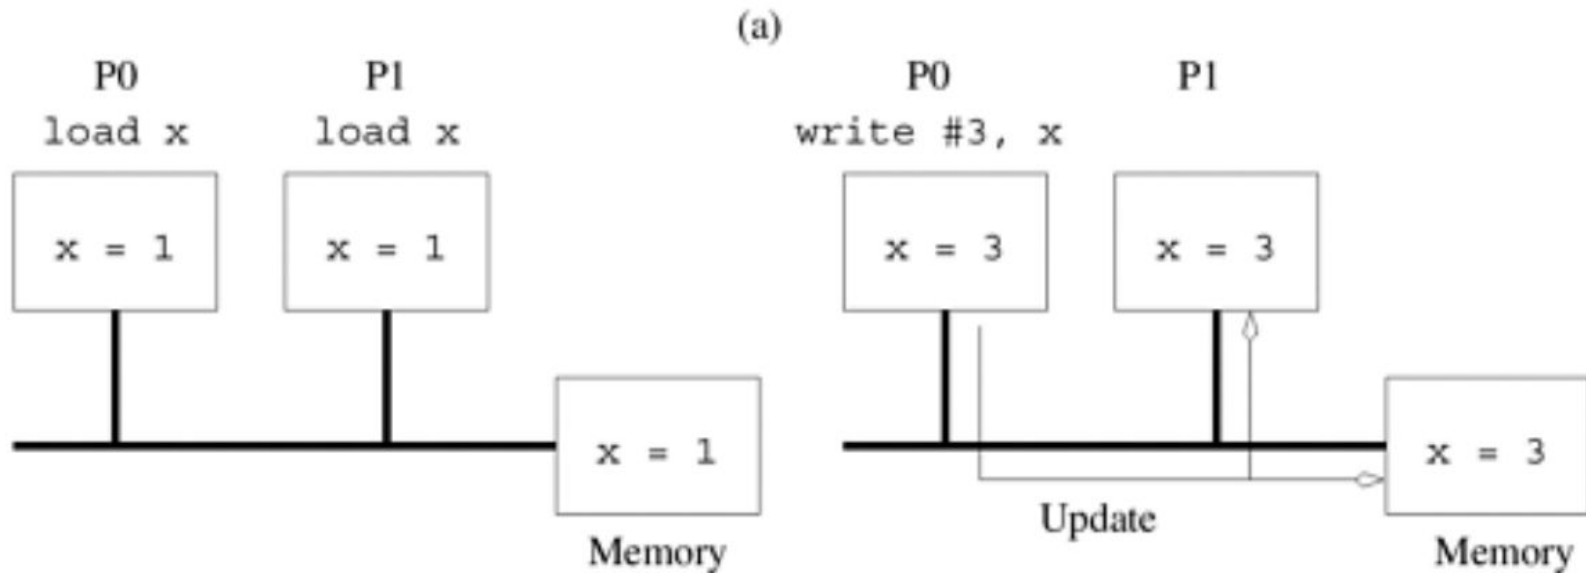
# Invalidate Protocol

Similar to lazy loading → Cache lines are marked as invalid on other processors, but not updated until referenced



# Update Protocol

Similar to eager loading → Cache lines are updated on all other processors as soon as they are updated on one




# General Three State Protocol

❑ All cache lines have an additional state attached to them on each processor

- Shared → All variables the same
- Dirty → Processor which updates the variable
- Invalid → All other processors must mark their line as invalid

❑ Solid Lines → Processor actions

❑ Dashed Lines → Coherence actions



# False Sharing

- ❑ Different variables are updated on the same cache line
- ❑ Can lead to excessive overhead even though different things are updated
- ❑ KEY TIP: Try to have your threads update different cache lines if possible on CPUs

