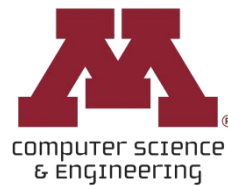


# CSCI 5451: Introduction to Parallel Computing

## Lecture 13: MPI Examples



UNIVERSITY OF MINNESOTA  
**Driven to Discover®**

# Announcements (10/15)

- ❑ HW1 → Due later today ([Canvas](#))
- ❑ HW 2 released on Monday, Oct 20
- ❑ Group Formation due Sunday, Oct 19 ([Canvas](#))



# Lecture Overview

## ❑ Leftover MPI Functions from (10/13)

- MPI\_Alltoallv
- MPI\_Comm\_split & MPI\_Cart\_sub

## ❑ MPI Examples (Whiteboard walkthrough)

- Row-Wise Matrix-Vector Multiplication
- Column-Wise Matrix-Vector Multiplication
- 2-D Matrix-Vector Multiplication
- Dijkstra's Single Source Shortest Path
- Sample Sort



# Lecture Overview

## ❑ Leftover MPI Functions from (10/13)

- **MPI\_Alltoallv**
- MPI\_Comm\_split & MPI\_Cart\_sub

## ❑ MPI Examples (Whiteboard walkthrough)

- Row-Wise Matrix-Vector Multiplication
- Column-Wise Matrix-Vector Multiplication
- 2-D Matrix-Vector Multiplication
- Dijkstra's Single Source Shortest Path
- Sample Sort



# MPI\_Alltoallv

```
int MPI_Alltoallv(const void *sendbuf,  
                 const int *sendcounts, const int *sdispls,  
                 MPI_Datatype sendtype, void *recvbuf,  
                 const int *recvcounts, const int *rdispls,  
                 MPI_Datatype recvtype,  
                 MPI_Comm comm);
```

Each process sends variable-sized data to all others and receives variable-sized data from all others.

- **sendbuf**: Starting address of send buffer.
- **sendcounts**: Array specifying number of elements sent to each rank.
- **sdispls**: Array of offsets in **sendbuf**.
- **sendtype**: Type of send data.
- **recvbuf**: Starting address of receive buffer.
- **recvcounts**: Array specifying number of elements received from each rank.
- **rdispls**: Array of offsets in **recvbuf**.
- **recvtype**: Type of receive data.
- **comm**: Communicator handle.



# MPI\_Alltoallv

```
int MPI_Alltoallv(const void *sendbuf,  
    const int *sendcounts, const int *sdispls,  
    MPI_Datatype sendtype, void *recvbuf,  
    const int *recvcounts, const int *rdispls,  
    MPI_Datatype recvtype,  
    MPI_Comm comm);
```

Each process sends variable-sized data to all others and receives variable-sized data from all others.

```
int sendbuf[10];  
for (int i=0; i<10; i++) sendbuf[i] = rank*10 + i;  
int sendcounts[4] = {1,2,3,4};  
int sdispls[4] = {0,1,3,6};  
  
int recvcounts[4] = {rank+1, rank+1, rank+1, rank+1};  
int rdispls[4] = {0, (rank+1), (rank+1)*2, (rank+1)*3};  
int recvbuf[16];  
  
MPI_Alltoallv(sendbuf, sendcounts, sdispls, MPI_INT, recvbuf,  
    recvcounts, rdispls, MPI_INT, MPI_COMM_WORLD);  
  
printf("Rank %d received:", rank);  
for (int i=0; i<(rank+1)*4; i++)  
    printf(" %d", recvbuf[i]);  
printf("\n");
```



# MPI\_Alltoallv

```
int MPI_Alltoallv(const void *sendbuf,  
    const int *sendcounts, const int *sdispls,  
    MPI_Datatype sendtype, void *recvbuf,  
    const int *recvcounts, const int *rdispls,  
    MPI_Datatype recvtype,  
    MPI_Comm comm);
```

Each process sends variable-sized data to all others and receives variable-sized data from all others.

```
int sendbuf[10];  
for (int i=0; i<10; i++) sendbuf[i] = rank*10 + i;  
int sendcounts[4] = {1,2,3,4};  
int sdispls[4] = {0,1,3,6};  
  
int recvcounts[4] = {rank+1, rank+1, rank+1, rank+1};  
int rdispls[4] = {0, (rank+1), (rank+1)*2, (rank+1)*3};  
int recvbuf[16];  
  
MPI_Alltoallv(sendbuf, sendcounts, sdispls, MPI_INT, recvbuf,  
    recvcounts, rdispls, MPI_INT, MPI_COMM_WORLD);  
  
printf("Rank %d received:", rank);  
for (int i=0; i<(rank+1)*4; i++)  
    printf(" %d", recvbuf[i]);  
printf("\n");
```

Rank 0 received: 0 10 20 30  
Rank 1 received: 1 2 11 12 21 22 31 32  
Rank 2 received: 3 4 5 13 14 15 23 24 25 33 34 35  
Rank 3 received: 6 7 8 9 16 17 18 19 26 27 28 29 36 37 38 39



# Lecture Overview

## ❑ Leftover MPI Functions from (10/13)

- MPI\_Alltoallv
- **MPI\_Comm\_split & MPI\_Cart\_sub**

## ❑ MPI Examples (Whiteboard walkthrough)

- Row-Wise Matrix-Vector Multiplication
- Column-Wise Matrix-Vector Multiplication
- 2-D Matrix-Vector Multiplication
- Dijkstra's Single Source Shortest Path
- Sample Sort





# MPI\_Comm\_split

```
int MPI_Comm_split(MPI_Comm comm,  
                  int color, int key,  
                  MPI_Comm *newcomm);
```

Creates new communicator by splitting an existing communicator into subgroups based on **color**.

Processes with the same **color** are grouped together; within each group, ranks are ordered by **key**



# MPI\_Comm\_split

```
int MPI_Comm_split(MPI_Comm comm,  
                  int color, int key,  
                  MPI_Comm *newcomm);
```

Creates new communicator by splitting an existing communicator into subgroups based on **color**.

Processes with the same **color** are grouped together; within each group, ranks are ordered by **key**

- **comm**: Existing communicator to be split (e.g., **MPI\_COMM\_WORLD**).
- **color**: Integer identifier for subgroup membership.
  - All processes with the same **color** form a new communicator.
  - A process can set **color = MPI\_UNDEFINED** to be excluded entirely.
- **key**: Determines ordering of ranks within the new communicator.
  - Lower keys get lower ranks.
  - Equal keys are ordered according to ordering in previous communicator
- **newcomm**: Output handle to the newly created communicator.



# MPI\_Comm\_split

```
int MPI_Comm_split(MPI_Comm comm,  
                  int color, int key,  
                  MPI_Comm *newcomm);
```

```
int key = 1;  
if (rank <= 2)    color = 0;  
else if (rank <= 6) color = 1;  
else             color = 2;  
  
MPI_Comm newcomm;  
MPI_Comm_split(MPI_COMM_WORLD,  
              color, key, &newcomm);
```



# MPI\_Comm\_split

```
int MPI_Comm_split(MPI_Comm comm,  
                  int color, int key,  
                  MPI_Comm *newcomm);
```

```
int key = 1;  
if (rank <= 2)    color = 0;  
else if (rank <= 6) color = 1;  
else             color = 2;
```

```
MPI_Comm newcomm;  
MPI_Comm_split(MPI_COMM_WORLD,  
              color, key, &newcomm);
```

process	0	1	2	3	4	5	6	7
color	0	0	0	1	1	1	1	2
key	1	1	1	1	1	1	1	1

Which processes will be grouped together?




# MPI\_Comm\_split

```
int MPI_Comm_split(MPI_Comm comm,  
                  int color, int key,  
                  MPI_Comm *newcomm);
```

```
int key = 1;  
if (rank <= 2)    color = 0;  
else if (rank <= 6) color = 1;  
else             color = 2;
```

```
MPI_Comm newcomm;  
MPI_Comm_split(MPI_COMM_WORLD,  
              color, key, &newcomm);
```



# MPI\_Cart\_sub

```
int MPI_Cart_sub(MPI_Comm comm,  
                 const int *remain_dims,  
                 MPI_Comm *newcomm);
```

Creates a lower-dimensional subcommunicator from an existing Cartesian topology communicator by selecting which dimensions to keep.



# MPI\_Cart\_sub

```
int MPI_Cart_sub(MPI_Comm comm,  
                const int *remain_dims,  
                MPI_Comm *newcomm);
```


Creates a lower-dimensional subcommunicator from an existing Cartesian topology communicator by selecting which dimensions to keep.

- **comm**: Input communicator with Cartesian topology (created via **MPI\_Cart\_create**).
- **remain\_dims**: Logical array of length equal to the number of dimensions.
  - **remain\_dims[i] = 1** → keep this dimension.
  - **remain\_dims[i] = 0** → collapse (drop) this dimension.
- **newcomm**: Output communicator corresponding to the subgrid of the original topology.



# MPI\_Cart\_sub

```
int MPI_Cart_sub(MPI_Comm comm,  
    const int *remain_dims,  
    MPI_Comm *newcomm);
```



```
int dims[3]    = {2, 4, 7};    // grid dimensions  
int periods[3] = {0, 0, 0};    // non-periodic in all dimensions  
int reorder    = 0;            // keep original ranks
```

```
MPI_Comm cart_comm;  
MPI_Cart_create(MPI_COMM_WORLD, 3, dims, periods,  
    reorder, &cart_comm);
```


```
int keep_dims[3] = {1, 0, 1};  
MPI_Comm sub_comm;  
MPI_Cart_sub(cart_comm, keep_dims, &sub_comm);
```





# MPI\_Cart\_sub

```
int MPI_Cart_sub(MPI_Comm comm,  
                const int *remain_dims,  
                MPI_Comm *newcomm);
```



```
int dims[3]   = {2, 4, 7};    // grid dimensions  
int periods[3] = {0, 0, 0};    // non-periodic in all dimensions  
int reorder   = 0;            // keep original ranks
```

```
MPI_Comm cart_comm;
```

```
MPI_Cart_create(MPI_COMM_WORLD, 3, dims, periods,  
                reorder, &cart_comm);
```

```
int keep_dims[3] = {1, 0, 1};
```


```
MPI_Comm sub_comm;
```

```
MPI_Cart_sub(cart_comm, keep_dims, &sub_comm);
```



# MPI\_Cart\_sub


```
int MPI_Cart_sub(MPI_Comm comm,  
                const int *remain_dims,  
                MPI_Comm *newcomm);
```



```
int dims[3]    = {2, 4, 7};    // grid dimensions  
int periods[3] = {0, 0, 0};    // non-periodic in all dimensions  
int reorder    = 0;            // keep original ranks
```


```
MPI_Comm cart_comm;  
MPI_Cart_create(MPI_COMM_WORLD, 3, dims, periods,  
                reorder, &cart_comm);
```

```
int keep_dims[3] = {1, 0, 1};  
MPI_Comm sub_comm;  
MPI_Cart_sub(cart_comm, keep_dims, &sub_comm);
```



# MPI\_Cart\_sub

```
int MPI_Cart_sub(MPI_Comm comm,  
                const int *remain_dims,  
                MPI_Comm *newcomm);
```



```
int dims[3]   = {2, 4, 7};    // grid dimensions  
int periods[3] = {0, 0, 0};    // non-periodic in all dimensions  
int reorder   = 0;            // keep original ranks
```


```
MPI_Comm cart_comm;  
MPI_Cart_create(MPI_COMM_WORLD, 3, dims, periods,  
                reorder, &cart_comm);
```

```
int keep_dims[3] = {0, 0, 1};  
MPI_Comm sub_comm;  
MPI_Cart_sub(cart_comm, keep_dims, &sub_comm);
```



# MPI\_Cart\_sub


```
int MPI_Cart_sub(MPI_Comm comm,  
                const int *remain_dims,  
                MPI_Comm *newcomm);
```



```
int dims[3]    = {2, 4, 7};    // grid dimensions  
int periods[3] = {0, 0, 0};    // non-periodic in all dimensions  
int reorder    = 0;            // keep original ranks
```

```
MPI_Comm cart_comm;  
MPI_Cart_create(MPI_COMM_WORLD, 3, dims, periods,  
                reorder, &cart_comm);
```

```
int keep_dims[3] = {0, 0, 1};  
MPI_Comm sub_comm;  
MPI_Cart_sub(cart_comm, keep_dims, &sub_comm);
```



# Lecture Overview

## ❑ Leftover MPI Functions from (10/13)

- MPI\_Alltoallv
- MPI\_Comm\_split & MPI\_Cart\_sub

## ❑ MPI Examples (Whiteboard walkthrough)

- **Row-Wise Matrix-Vector Multiplication**
- Column-Wise Matrix-Vector Multiplication
- 2-D Matrix-Vector Multiplication
- Dijkstra's Single Source Shortest Path
- Sample Sort



# Row-Wise Matrix-Vector Multiplication

```
void RowMatrixVectorMultiply(int n, double *a, double *b, double *x, MPI_Comm comm)
{
    int i, j;
    int nlocal; /* Number of locally stored rows of A */
    double *fb; /* Buffer that stores the entire vector b */
    int npes, myrank;
    MPI_Status status;

    /* Get information about the communicator */
    MPI_Comm_size(comm, &npes);
    MPI_Comm_rank(comm, &myrank);

    /* Allocate memory to store the entire vector b */
    fb = (double *)malloc(n * sizeof(double));

    nlocal = n / npes;
```



# Row-Wise Matrix-Vector Multiplication

```
/* Gather the entire vector b on each processor using MPI's ALLGATHER operation */
MPI_Allgather(b, nlocal, MPI_DOUBLE, fb, nlocal, MPI_DOUBLE, comm);

/* Perform the matrix-vector multiplication involving the locally stored submatrix */
for (i = 0; i < nlocal; i++) {
    x[i] = 0.0;
    for (j = 0; j < n; j++)
        x[i] += a[i * n + j] * fb[j];
}

free(fb);
}
```



# Lecture Overview

## ❑ Leftover MPI Functions from (10/13)

- MPI\_Alltoallv
- MPI\_Comm\_split & MPI\_Cart\_sub

## ❑ MPI Examples (Whiteboard walkthrough)

- Row-Wise Matrix-Vector Multiplication
- **Column-Wise Matrix-Vector Multiplication**
- 2-D Matrix-Vector Multiplication
- Dijkstra's Single Source Shortest Path
- Sample Sort





# Column-Wise Matrix-Vector Multiplication

```
void ColMatrixVectorMultiply(int n, double *a, double *b, double *x, MPI_Comm comm)
{
    int i, j;
    int nlocal;
    double *px;
    double *fx;
    int npes, myrank;
    MPI_Status status;

    /* Get identity and size information from the communicator */
    MPI_Comm_size(comm, &npes);
    MPI_Comm_rank(comm, &myrank);

    nlocal = n / npes;
```



# Column-Wise Matrix-Vector Multiplication

```
/* Allocate memory for arrays storing intermediate results. */
px = (double *)malloc(n * sizeof(double));
fx = (double *)malloc(n * sizeof(double));

/* Compute the partial dot products that correspond to the local columns of A. */
for (i = 0; i < n; i++) {
    px[i] = 0.0;
    for (j = 0; j < nlocal; j++)
        px[i] += a[i * nlocal + j] * b[j];
}

/* Sum up the results by performing an element-wise reduction operation */
MPI_Reduce(px, fx, n, MPI_DOUBLE, MPI_SUM, 0, comm);

/* Redistribute fx in a fashion similar to that of vector b */
MPI_Scatter(fx, nlocal, MPI_DOUBLE, x, nlocal, MPI_DOUBLE, 0, comm);

free(px);
free(fx);
}
```



# Lecture Overview

## ❑ Leftover MPI Functions from (10/13)

- MPI\_Alltoallv
- MPI\_Comm\_split & MPI\_Cart\_sub

## ❑ MPI Examples (Whiteboard walkthrough)

- Row-Wise Matrix-Vector Multiplication
- Column-Wise Matrix-Vector Multiplication
- **2-D Matrix-Vector Multiplication**
- Dijkstra's Single Source Shortest Path
- Sample Sort



# 2-D Matrix-Vector Multiplication

```
void MatrixVectorMultiply_2D(int n, double *a, double *b, double *x, MPI_Comm comm)
{
    int ROW = 0, COL = 1;
    int i, j, nlocal;
    double *px;
    int npes, dims[2], periods[2], keep_dims[2];
    int myrank, my2drank, mycoords[2];
    int other_rank, coords[2];
    MPI_Status status;
    MPI_Comm comm_2d, comm_row, comm_col;

    MPI_Comm_size(comm, &npes);
    MPI_Comm_rank(comm, &myrank);

    dims[ROW] = dims[COL] = sqrt(npes);
    nlocal = n / dims[ROW];
    px = malloc(nlocal * sizeof(double));
```



```
periods[ROW] = periods[COL] = 1;
MPI_Cart_create(MPI_COMM_WORLD, 2, dims, periods, 1, &comm_2d);
MPI_Comm_rank(comm_2d, &my2drank);
MPI_Cart_coords(comm_2d, my2drank, 2, mycoords);

keep_dims[ROW] = 0; keep_dims[COL] = 1;
MPI_Cart_sub(comm_2d, keep_dims, &comm_row);

keep_dims[ROW] = 1; keep_dims[COL] = 0;
MPI_Cart_sub(comm_2d, keep_dims, &comm_col);

if (mycoords[COL] == 0 && mycoords[ROW] != 0) {
    coords[ROW] = mycoords[ROW];
    coords[COL] = mycoords[ROW];
    MPI_Cart_rank(comm_2d, coords, &other_rank);
    MPI_Send(b, nlocal, MPI_DOUBLE, other_rank, 1, comm_2d);
}

if (mycoords[ROW] == mycoords[COL] && mycoords[ROW] != 0) {
    coords[ROW] = mycoords[ROW];
    coords[COL] = 0;
    MPI_Cart_rank(comm_2d, coords, &other_rank);
    MPI_Recv(b, nlocal, MPI_DOUBLE, other_rank, 1, comm_2d, &status);
}
```



# 2-D Matrix-Vector Multiplication

```
coords[0] = mycoords[COL];
MPI_Cart_rank(comm_col, coords, &other_rank);
MPI_Bcast(b, nlocal, MPI_DOUBLE, other_rank, comm_col);

for (i = 0; i < nlocal; i++) {
    px[i] = 0.0;
    for (j = 0; j < nlocal; j++)
        px[i] += a[i * nlocal + j] * b[j];
}

coords[0] = 0;
MPI_Cart_rank(comm_row, coords, &other_rank);
MPI_Reduce(px, x, nlocal, MPI_DOUBLE, MPI_SUM, other_rank, comm_row);

MPI_Comm_free(&comm_2d);
MPI_Comm_free(&comm_row);
MPI_Comm_free(&comm_col);
free(px);
}
```




# Lecture Overview

- ❑ Leftover MPI Functions from (10/13)
  - MPI\_Alltoallv
  - MPI\_Comm\_split & MPI\_Cart\_sub
- ❑ MPI Examples (Whiteboard walkthrough)
  - Row-Wise Matrix-Vector Multiplication
  - Column-Wise Matrix-Vector Multiplication
  - 2-D Matrix-Vector Multiplication
  - **Dijkstra's Single Source Shortest Path**
  - Sample Sort



# Dijkstra's Single-Source Shortest Path (Review)

- Find the shortest path from a single node in a graph to all other nodes in a graph
- Serial Example Walkthrough





# Dijkstra's Single-Source Shortest Path

```
void SingleSource(int n, int source, int *wgt, int *lengths, MPI_Comm comm)
{
    int i, j;
    int nlocal;    /* The number of vertices stored locally */
    int *marker;    /* Used to mark the vertices belonging to Vo */
    int firstvtx;    /* The index number of the first vertex that is stored locally */
    int lastvtx;    /* The index number of the last vertex that is stored locally */
    int u, udist;
    int lminpair[2], gminpair[2];
    int npes, myrank;
    MPI_Status status;

    MPI_Comm_size(comm, &npes);
    MPI_Comm_rank(comm, &myrank);

    nlocal = n / npes;
    firstvtx = myrank * nlocal;
    lastvtx = firstvtx + nlocal - 1;
```



# Dijkstra's Single-Source Shortest Path

```
/* Set the initial distances from source to all the other vertices */  
for (j = 0; j < nlocal; j++)  
    lengths[j] = wgt[source * nlocal + j];  
  
/* This array is used to indicate if the shortest path to a vertex has been found */  
marker = (int *)malloc(nlocal * sizeof(int));  
for (j = 0; j < nlocal; j++)  
    marker[j] = 1;  
  
/* The process that stores the source vertex marks it as seen */  
if (source >= firstvtx && source <= lastvtx)  
    marker[source - firstvtx] = 0;
```



```

/* The main loop of Dijkstra's algorithm */
for (i = 1; i < n; i++) {
    /* Step 1: Find the local vertex at the smallest distance from source */
    lminpair[0] = MAXINT; /* architecture dependent large number */
    lminpair[1] = -1;

    for (j = 0; j < nlocal; j++) {
        if (marker[j] && lengths[j] < lminpair[0]) {
            lminpair[0] = lengths[j];
            lminpair[1] = firstvtx + j;
        }
    }

    /* Step 2: Compute the global minimum vertex and insert it into Vc */
    MPI_Allreduce(lminpair, gminpair, 1, MPI_2INT, MPI_MINLOC, comm);
    udist = gminpair[0];
    u = gminpair[1];

    /* The process that stores the minimum vertex marks it as seen */
    if (u == lminpair[1])
        marker[u - firstvtx] = 0;

    /* Step 3: Update the distances given that u got inserted */
    for (j = 0; j < nlocal; j++) {
        if (marker[j] && udist + wgt[u * nlocal + j] < lengths[j])
            lengths[j] = udist + wgt[u * nlocal + j];
    }
}
free(marker);
}

```



# Lecture Overview

## ❑ Leftover MPI Functions from (10/13)

- MPI\_Alltoallv
- MPI\_Comm\_split & MPI\_Cart\_sub

## ❑ MPI Examples (Whiteboard walkthrough)

- Row-Wise Matrix-Vector Multiplication
- Column-Wise Matrix-Vector Multiplication
- 2-D Matrix-Vector Multiplication
- Dijkstra's Single Source Shortest Path
- **Sample Sort**



# Sample Sort

$P_0$								$P_1$								$P_2$							
22	7	13	18	2	17	1	14	20	6	10	24	15	9	21	3	16	19	23	4	11	12	5	8

Initial element  
distribution

$P_0$								$P_1$								$P_2$							
1	2	7	13	14	17	18	22	3	6	9	10	15	20	21	24	4	5	8	11	12	16	19	23

Local sort &  
sample selection

7	17	9	20	8	16
---	----	---	----	---	----

Sample combining

7	8	9	16	17	20
---	---	---	----	----	----

Global splitter  
selection

$P_0$								$P_1$								$P_2$							
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24

Final element  
assignment



# Sample Sort

```
int *SampleSort(int n, int *elmnts, int *nsorted, MPI_Comm comm)
{
    int i, j, nlocal, npes, myrank;
    int *sorted_elmnts, *splitters, *allpicks;
    int *scounts, *sdispls, *rcounts, *rdispls;

    /* Get communicator-related information */
    MPI_Comm_size(comm, &npes);
    MPI_Comm_rank(comm, &myrank);

    nlocal = n / npes;

    /* Allocate memory for the arrays that will store the splitters */
    splitters = (int *)malloc(npes * sizeof(int));
    allpicks = (int *)malloc(npes * (npes - 1) * sizeof(int));

    /* Sort local array */
    qsort(elmnts, nlocal, sizeof(int), IncOrder);
```



# Sample Sort

```
/* Select local npes-1 equally spaced elements */
for (i = 1; i < npes; i++)
    splitters[i - 1] = elmnts[i * nlocal / npes];

/* Gather the samples in the processors */
MPI_Allgather(splitters, npes - 1, MPI_INT, allpicks, npes - 1, MPI_INT, comm);

/* Sort these samples */
qsort(allpicks, npes * (npes - 1), sizeof(int), IncOrder);

/* Select splitters */
for (i = 1; i < npes; i++)
    splitters[i - 1] = allpicks[i * npes - (int)ceil((double)npes/2)];
splitters[npes - 1] = MAXINT;

/* Compute the number of elements that belong to each bucket */
counts = (int *)malloc(npes * sizeof(int));
for (i = 0; i < npes; i++)
    counts[i] = 0;
for (j = i = 0; i < nlocal; i++) {
    if (elmnts[i] < splitters[j])
        counts[j]++;
    else
        counts[++j]++;
}
```



# Sample Sort

```
/* Determine starting locations of each bucket's elements */
sdispls = (int *)malloc(npes * sizeof(int));
sdispls[0] = 0;
for (i = 1; i < npes; i++)
    sdispls[i] = sdispls[i - 1] + counts[i - 1];

/* Inform all processes about receive counts */
rcounts = (int *)malloc(npes * sizeof(int));
MPI_Alltoall(counts, 1, MPI_INT, rcounts, 1, MPI_INT, comm);

/* Compute receive displacements */
rdispls = (int *)malloc(npes * sizeof(int));
rdispls[0] = 0;
for (i = 1; i < npes; i++)
    rdispls[i] = rdispls[i - 1] + rcounts[i - 1];

*nsorted = rdispls[npes - 1] + rcounts[npes - 1];
sorted_elmnts = (int *)malloc((*nsorted) * sizeof(int));

/* Exchange elements between processors */
MPI_Alltoallv(elmnts, counts, sdispls, MPI_INT,
               sorted_elmnts, rcounts, rdispls, MPI_INT, comm);

/* Final local sort */
qsort(sorted_elmnts, *nsorted, sizeof(int), IncOrder);
```





# Sample Sort

```
free(splitters);  
free(allpicks);  
free(scounts);  
free(sdispls);  
free(rcounts);  
free(rdispls);  
  
return sorted_elmnts;  
}
```

