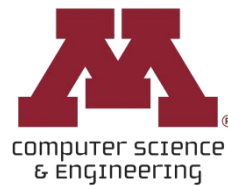


CSCI 5451: Introduction to Parallel Computing

Lecture 14/15: MPI Examples (cont'd)



UNIVERSITY OF MINNESOTA
Driven to Discover®

Announcements (10/20)

- ❑ HW 2 released (due Nov 2)
 - o [Canvas](#)
 - o Currently uploaded to [course site](#)
 - o Canvas autograder currently running
 - o No hidden tests
- ❑ Group Formation due yesterday → Sunday, Oct 19 ([Canvas](#))



Lecture Overview

□ MPI Examples (Whiteboard walkthrough)

- **2-D Matrix-Vector Multiplication**
- Dijkstra's Single Source Shortest Path
- Sample Sort



2-D Matrix-Vector Multiplication

```
void MatrixVectorMultiply_2D(int n, double *a, double *b, double *x, MPI_Comm comm)
{
    int ROW = 0, COL = 1;
    int i, j, nlocal;
    double *px;
    int npes, dims[2], periods[2], keep_dims[2];
    int myrank, my2drank, mycoords[2];
    int other_rank, coords[2];
    MPI_Status status;
    MPI_Comm comm_2d, comm_row, comm_col;

    MPI_Comm_size(comm, &npes);
    MPI_Comm_rank(comm, &myrank);

    dims[ROW] = dims[COL] = sqrt(npes);
    nlocal = n / dims[ROW];
    px = malloc(nlocal * sizeof(double));
```



```
periods[ROW] = periods[COL] = 1;
MPI_Cart_create(MPI_COMM_WORLD, 2, dims, periods, 1, &comm_2d);
MPI_Comm_rank(comm_2d, &my2drank);
MPI_Cart_coords(comm_2d, my2drank, 2, mycoords);

keep_dims[ROW] = 0; keep_dims[COL] = 1;
MPI_Cart_sub(comm_2d, keep_dims, &comm_row);

keep_dims[ROW] = 1; keep_dims[COL] = 0;
MPI_Cart_sub(comm_2d, keep_dims, &comm_col);

if (mycoords[COL] == 0 && mycoords[ROW] != 0) {
    coords[ROW] = mycoords[ROW];
    coords[COL] = mycoords[ROW];
    MPI_Cart_rank(comm_2d, coords, &other_rank);
    MPI_Send(b, nlocal, MPI_DOUBLE, other_rank, 1, comm_2d);
}

if (mycoords[ROW] == mycoords[COL] && mycoords[ROW] != 0) {
    coords[ROW] = mycoords[ROW];
    coords[COL] = 0;
    MPI_Cart_rank(comm_2d, coords, &other_rank);
    MPI_Recv(b, nlocal, MPI_DOUBLE, other_rank, 1, comm_2d, &status);
}
```



2-D Matrix-Vector Multiplication

```
coords[0] = mycoords[COL];
MPI_Cart_rank(comm_col, coords, &other_rank);
MPI_Bcast(b, nlocal, MPI_DOUBLE, other_rank, comm_col);

for (i = 0; i < nlocal; i++) {
    px[i] = 0.0;
    for (j = 0; j < nlocal; j++)
        px[i] += a[i * nlocal + j] * b[j];
}

coords[0] = 0;
MPI_Cart_rank(comm_row, coords, &other_rank);
MPI_Reduce(px, x, nlocal, MPI_DOUBLE, MPI_SUM, other_rank, comm_row);

MPI_Comm_free(&comm_2d);
MPI_Comm_free(&comm_row);
MPI_Comm_free(&comm_col);
free(px);
}
```



Lecture Overview


□ MPI Examples (Whiteboard walkthrough)

- 2-D Matrix-Vector Multiplication
- **Dijkstra's Single Source Shortest Path**
- Sample Sort



Dijkstra's Single-Source Shortest Path (Review)

- Find the shortest path from a single node in a graph to all other nodes in a graph
- Serial Example Walkthrough (Starting from A)



Dijkstra's Single-Source Shortest Path

```
void SingleSource(int n, int source, int *wgt, int *lengths, MPI_Comm comm)
{
    int i, j;
    int nlocal;    /* The number of vertices stored locally */
    int *marker;    /* Used to mark the vertices belonging to Vo */
    int firstvtx;   /* The index number of the first vertex that is stored locally */
    int lastvtx;    /* The index number of the last vertex that is stored locally */
    int u, udist;
    int lminpair[2], gminpair[2];
    int npes, myrank;
    MPI_Status status;

    MPI_Comm_size(comm, &npes);
    MPI_Comm_rank(comm, &myrank);

    nlocal = n / npes;
    firstvtx = myrank * nlocal;
    lastvtx = firstvtx + nlocal - 1;
```



Dijkstra's Single-Source Shortest Path

```
/* Set the initial distances from source to all the other vertices */  
for (j = 0; j < nlocal; j++)  
    lengths[j] = wgt[source * nlocal + j];  
  
/* This array is used to indicate if the shortest path to a vertex has been found */  
marker = (int *)malloc(nlocal * sizeof(int));  
for (j = 0; j < nlocal; j++)  
    marker[j] = 1;  
  
/* The process that stores the source vertex marks it as seen */  
if (source >= firstvtx && source <= lastvtx)  
    marker[source - firstvtx] = 0;
```



```

/* The main loop of Dijkstra's algorithm */
for (i = 1; i < n; i++) {
    /* Step 1: Find the local vertex at the smallest distance from source */
    lminpair[0] = MAXINT; /* architecture dependent large number */
    lminpair[1] = -1;

    for (j = 0; j < nlocal; j++) {
        if (marker[j] && lengths[j] < lminpair[0]) {
            lminpair[0] = lengths[j];
            lminpair[1] = firstvtx + j;
        }
    }

    /* Step 2: Compute the global minimum vertex and insert it into Vc */
    MPI_Allreduce(lminpair, gminpair, 1, MPI_2INT, MPI_MINLOC, comm);
    udist = gminpair[0];
    u = gminpair[1];

    /* The process that stores the minimum vertex marks it as seen */
    if (u == lminpair[1])
        marker[u - firstvtx] = 0;

    /* Step 3: Update the distances given that u got inserted */
    for (j = 0; j < nlocal; j++) {
        if (marker[j] && udist + wgt[u * nlocal + j] < lengths[j])
            lengths[j] = udist + wgt[u * nlocal + j];
    }
}
free(marker);
}

```




Lecture Overview

□ MPI Examples (Whiteboard walkthrough)


- 2-D Matrix-Vector Multiplication
- Dijkstra's Single Source Shortest Path
- **Sample Sort**




Sample Sort




Initial element distribution




Local sort & sample selection



Sample combining



Global splitter selection



Final element assignment



Sample Sort

```
int *SampleSort(int n, int *elmnts, int *nsorted, MPI_Comm comm)
{
    int i, j, nlocal, npes, myrank;
    int *sorted_elmnts, *splitters, *allpicks;
    int *scounts, *sdispls, *rcounts, *rdispls;

    /* Get communicator-related information */
    MPI_Comm_size(comm, &npes);
    MPI_Comm_rank(comm, &myrank);

    nlocal = n / npes;

    /* Allocate memory for the arrays that will store the splitters */
    splitters = (int *)malloc(npes * sizeof(int));
    allpicks = (int *)malloc(npes * (npes - 1) * sizeof(int));

    /* Sort local array */
    qsort(elmnts, nlocal, sizeof(int), IncOrder);
```



Sample Sort

```
/* Select local npes-1 equally spaced elements */
for (i = 1; i < npes; i++)
    splitters[i - 1] = elmnts[i * nlocal / npes];

/* Gather the samples in the processors */
MPI_Allgather(splitters, npes - 1, MPI_INT, allpicks, npes - 1, MPI_INT, comm);

/* Sort these samples */
qsort(allpicks, npes * (npes - 1), sizeof(int), IncOrder);

/* Select splitters */
for (i = 1; i < npes; i++)
    splitters[i - 1] = allpicks[i * npes - (int)ceil((double)npes/2)];
splitters[npes - 1] = MAXINT;

/* Compute the number of elements that belong to each bucket */
counts = (int *)malloc(npes * sizeof(int));
for (i = 0; i < npes; i++)
    counts[i] = 0;
for (j = i = 0; i < nlocal; i++) {
    if (elmnts[i] < splitters[j])
        counts[j]++;
    else
        counts[++j]++;
}
```



Sample Sort

```
/* Determine starting locations of each bucket's elements */
sdispls = (int *)malloc(npes * sizeof(int));
sdispls[0] = 0;
for (i = 1; i < npes; i++)
    sdispls[i] = sdispls[i - 1] + counts[i - 1];

/* Inform all processes about receive counts */
rcounts = (int *)malloc(npes * sizeof(int));
MPI_Alltoall(counts, 1, MPI_INT, rcounts, 1, MPI_INT, comm);

/* Compute receive displacements */
rdispls = (int *)malloc(npes * sizeof(int));
rdispls[0] = 0;
for (i = 1; i < npes; i++)
    rdispls[i] = rdispls[i - 1] + rcounts[i - 1];

*nsorted = rdispls[npes - 1] + rcounts[npes - 1];
sorted_elmnts = (int *)malloc((*nsorted) * sizeof(int));

/* Exchange elements between processors */
MPI_Alltoallv(elmnts, counts, sdispls, MPI_INT,
               sorted_elmnts, rcounts, rdispls, MPI_INT, comm);

/* Final local sort */
qsort(sorted_elmnts, *nsorted, sizeof(int), IncOrder);
```



Sample Sort

```
free(splitters);  
free(allpicks);  
free(scounts);  
free(sdispls);  
free(rcounts);  
free(rdispls);  
  
return sorted_elmnts;  
}
```

