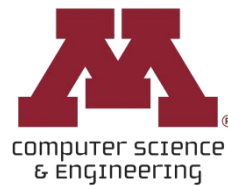


CSCI 5451: Introduction to Parallel Computing

Lecture 16: Analytical Modeling



UNIVERSITY OF MINNESOTA
Driven to Discover®

Announcements (10/27)

Project Groups → If you are not a part of a project group, or only have 1/2 people, you must form a group of 3-5

- ❑ If this is not done by Wednesday, we will form larger groups as needed



Lecture Overview

- ❑ Background
- ❑ Overheads
- ❑ Definitions
 - Serial Runtime/Parallel Runtime/Parallel Overhead
 - Speedup
 - Efficiency
 - Cost
- ❑ Granularity



Lecture Overview

☐ Background

☐ Overheads

☐ Definitions



- Serial Runtime/Parallel Runtime/Parallel Overhead
- Speedup
- Efficiency
- Cost

☐ Granularity



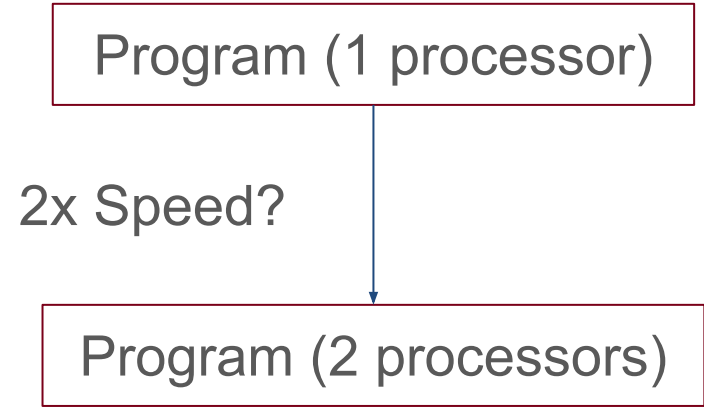
Analysis of Parallel Programs

- ❑ Serial programs can make use of asymptotic runtimes to define the relative 'goodness' of different algorithms
- ❑ Parallel programs introduce the use of potentially many processors at once
- ❑ We need to define terms which show the relative 'goodness' of different parallel algorithms beyond just runtimes



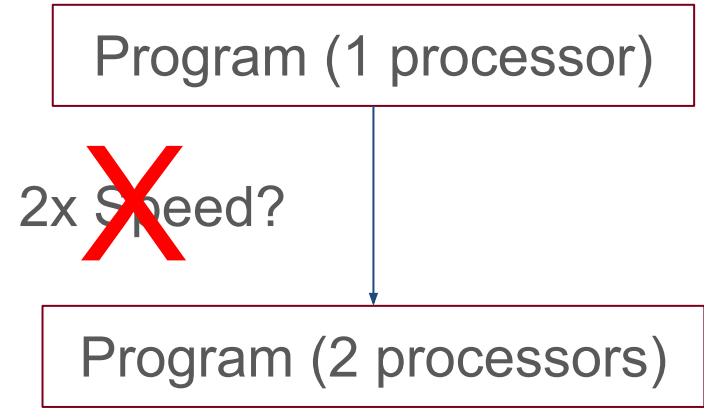
Desired Speedups

- ❑ Suppose we have some parallel program which we have chosen to run with twice as much hardware as was run in the serial case
- ❑ We would hope that this will lead to twice the speedups



Desired Speedups

- ❑ Suppose we have some parallel program which we have chosen to run with twice as much hardware as was run in the serial case
- ❑ We would hope that this will lead to twice the speedups



The overheads of running a program will usually result in less than the desired speedups




Lecture Overview

- ❑ Background
- ❑ **Overheads**
- ❑ Definitions
 - Serial Runtime/Parallel Runtime/Parallel Overhead
 - Speedup
 - Efficiency
 - Cost
- ❑ Granularity




Overhead In Programs (Idling)

- ❑ Load imbalances → If your program distributes work unevenly, then some processes will finish early & wait (e.g. sparse matrix multiplication)
- ❑ Uneven hardware → Uneven processing speeds/networking speeds/memory/memory locality
- ❑ Resource Contention on locks





Overhead In Programs (Communication)

- ❑ Distributed Memory → Any communication among processors will take excess time away from actual computation
- ❑ Shared Memory → Even though there may be shared memory, each thread will likely be running on separate cores - each with their own separate cache. If they need to share information, it must still be updated here.



Overhead In Programs (Redundant Computation)

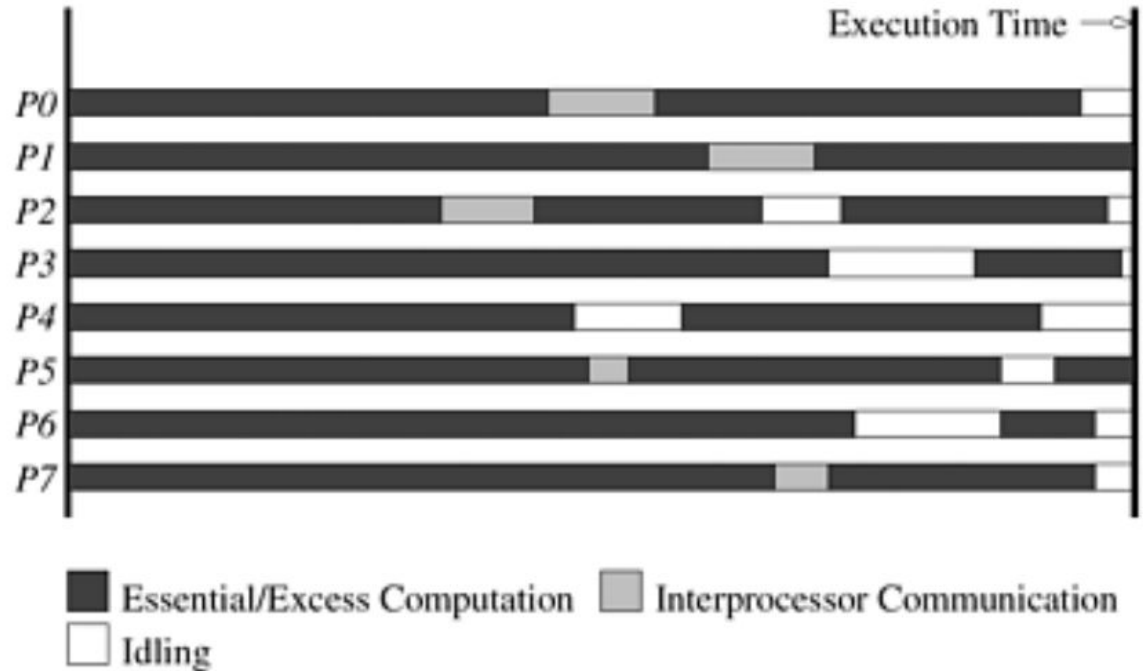
- ❑ Sometimes, we can accept performing some degree of excess computation in order to reduce extra communication or idling
- ❑ Parallel variants of sorting, graph traversal, FFT all make use of this
- ❑ Implies that we are computing more things than in the serial case



Overhead In programs

Overheads are not necessarily always a bad thing - they are a necessary component of parallel programming.

We will discuss how to asymptotically define our programs in terms of overhead



Lecture Overview

- ❑ Background
- ❑ Overheads
- ❑ **Definitions**
 - **Serial Runtime/Parallel Runtime/Parallel Overhead**
 - Speedup
 - Efficiency
 - Cost
- ❑ Granularity



Serial Runtime

- ❑ The time elapsed between the beginning & ending of a serial program
- ❑ You should ***always*** use the fastest algorithm to solve the given problem you are examining
- ❑ We will compare this speed to that of a hypothetical parallel program - ***don't*** compare to a slower algorithm, compare to the best serial algorithm

Serial Runtime: T_s



Parallel Runtime

- ❑ Analogous to the serial runtime (T_s)
- ❑ The time that elapses between the start of parallel computation to the moment the last processing element completes execution
- ❑ Note that this term is a function of the number of processes p

Parallel Runtime: T_p



Parallel Overhead

- ❑ Gathers all the previously discussed overheads into a single term
- ❑ Defines the excess processing time taken up across all processes

$$T_o = pT_p - T_s$$



Parallel Overhead

You want to minimize this term. More overheads means wasting processing power

- ❑ Gathers all the previously discussed overheads into a single term
- ❑ Defines the excess processing time taken up across all processes

$$T_o = pT_p - T_s$$



Lecture Overview

- ❑ Background
- ❑ Overheads
- ❑ **Definitions**
 - Serial Runtime/Parallel Runtime/Parallel Overhead
 - **Speedup**
 - Efficiency
 - Cost
- ❑ Granularity



Speedup

- ❑ Defines the ratio of time taken in the serial program to the parallel program
- ❑ If there are p processes, we want S to be closer to p
- ❑ That is, multiplying the hardware by p should get us close to p times speedups
- ❑ Assume that the serial program uses the same hardware as the parallel version

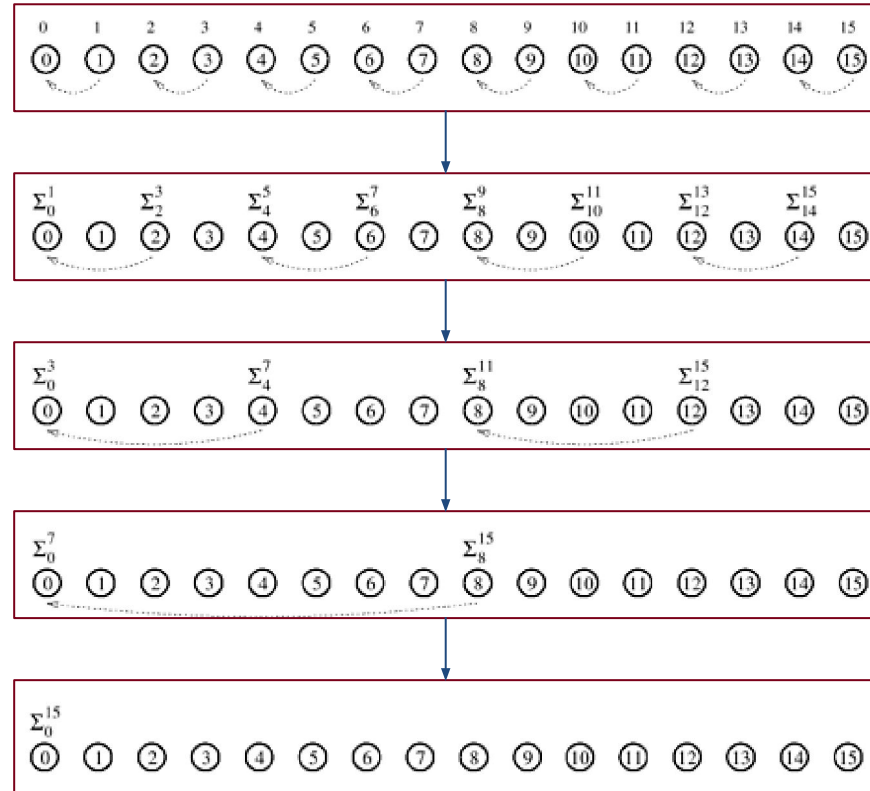
$$S = T_s / T_p$$



Speedup Example

How many processes can we use?

Sum up n numbers




Speedup Example

How many processes can we use?

$$p = n$$

Sum up n numbers



Speedup Example

How many processes can we use?


$$p = n$$

Big- Θ runtimes?

$$T_s = ?$$

$$T_p = ?$$

Sum up n numbers



Speedup Example

How many processes can we use?

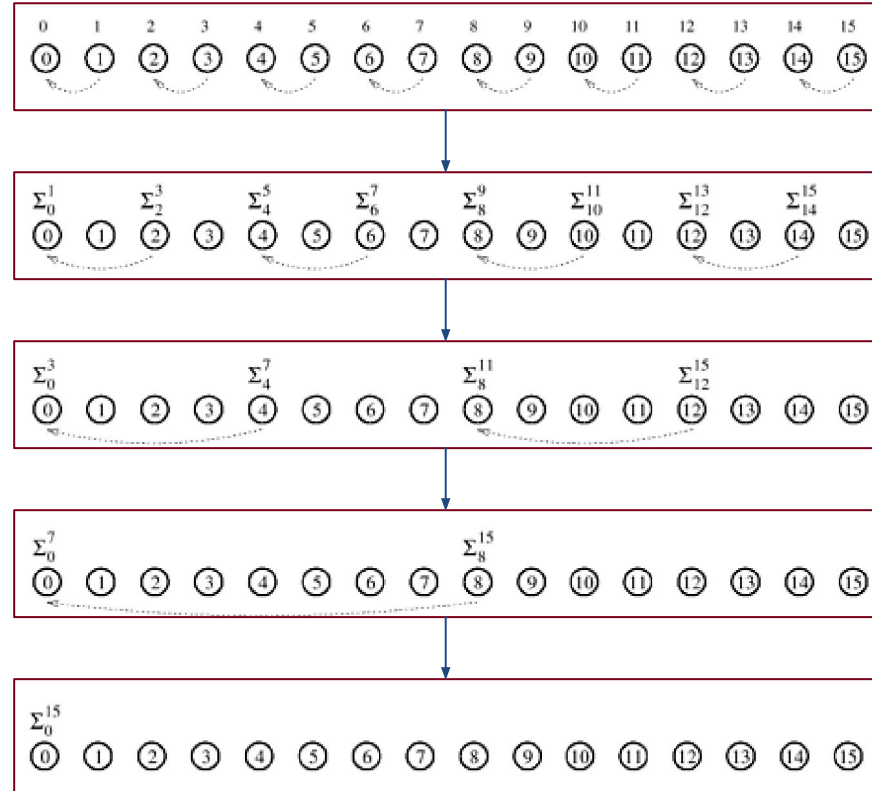
$$p = n$$

Big- Θ runtimes?

$$T_s = \Theta(n)$$

$$T_p = \Theta(\log n)$$

Sum up n numbers



Speedup Example

How many processes can we use?

$$p = n$$

Big- Θ runtimes?

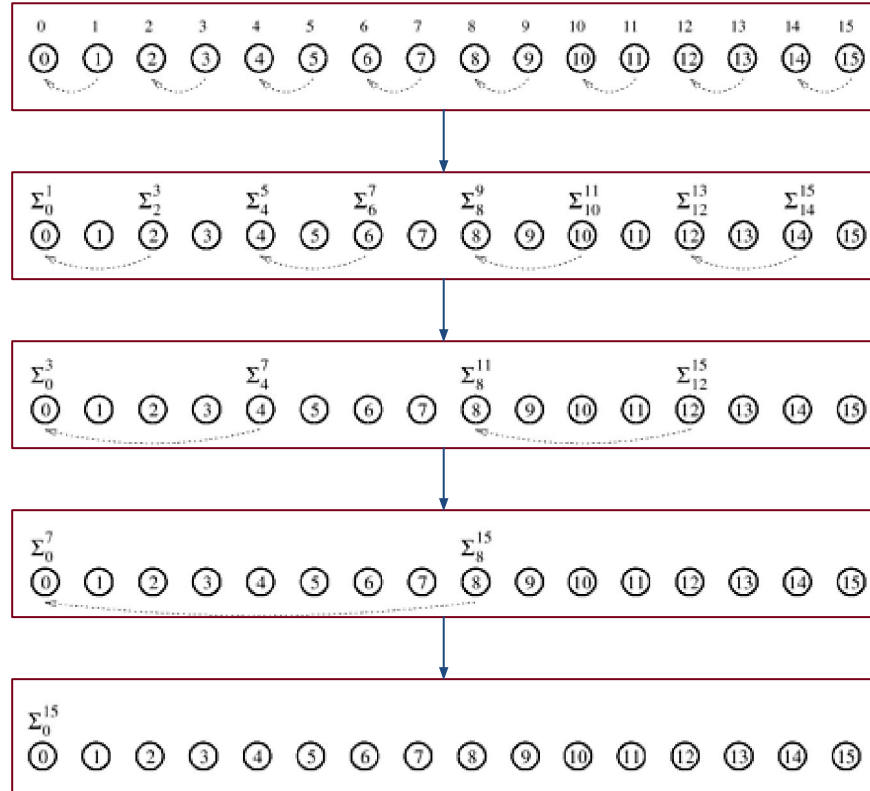
$$T_s = \Theta(n)$$

$$T_p = \Theta(\log n)$$

Speedup?

$$S = \Theta(?)$$

Sum up n numbers



Speedup Example

How many processes can we use?

$$p = n$$

Big- Θ runtimes?

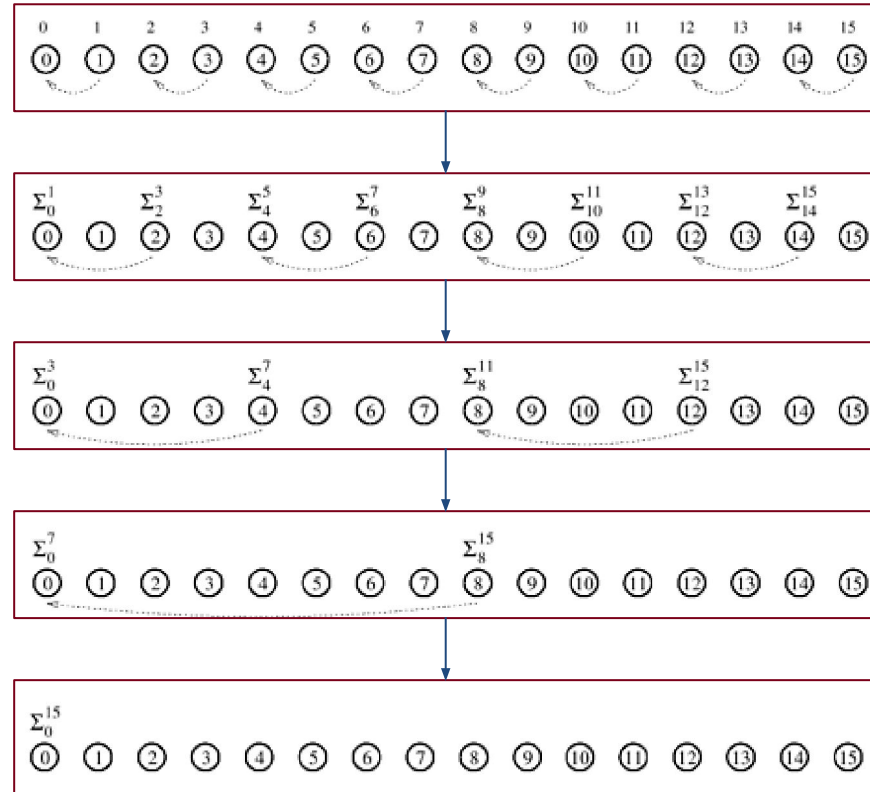
$$T_s = \Theta(n)$$

$$T_p = \Theta(\log n)$$

Speedup?

$$S = \Theta(n/\log n)$$

Sum up n numbers



Sorting Speedups Example

- ❑ Assume serial bubble sort takes 150 seconds
- ❑ Serial quicksort takes 30 seconds
- ❑ Parallel odd-even sort takes 40 seconds on 4 processes

$$T_p = ?$$

$$T_s = ?$$

$$S = ?$$

$$T_o = ?$$



Sorting Speedups Example

- ❑ Assume serial bubble sort takes 150 seconds
- ❑ Serial quicksort takes 30 seconds
- ❑ Parallel odd-even sort takes 40 seconds on 4 processes

$$T_p = 40$$

$$T_s = 30$$

$$S = .75$$

$$T_o = 130$$



Sorting Speedups Example

Always use best serial algorithm -
parallelizing in this example is pointless
as the best serial algorithm is better


- ❑ Assume serial bubble sort takes 150 seconds
- ❑ Serial quicksort takes 30 seconds
- ❑ Parallel odd-even sort takes 40 seconds on 4 processes

$$\begin{aligned}T_p &= 40 \\T_s &= 30 \\S &= .75 \\T_o &= 130\end{aligned}$$



Superlinear Speedups

- ❑ It is generally not possible to get greater than p speedup
- ❑ When this occurs, the program is said to exhibit superlinear speedup
- ❑ This is most commonly observed in exploration & caching



Caching Example

- ❑ Serial Version
 - Cache Latency: 2ns
 - DRAM latency: 100ns
 - 80% hit rate
 - Average access time:



Caching Example

- ❑ Serial Version
 - Cache Latency: 2ns
 - DRAM latency: 100ns
 - 80% hit rate
 - Average access time:
 $.8 * 2 + .2 * 100 = 21.6\text{ns}$
 - Assume program is memory bottlenecked & only performs one FLOP/memory access
 - 1 FLOP every 21.6ns $\rightarrow 1/(21.6 * 10^{-9}) = 46.3 \text{ Megaflops}$



Caching Example

Serial Version

- Cache Latency: 2ns
- DRAM latency: 100ns
- 80% hit rate
- Average access time:
 $.8 * 2 + .2 * 100 = 21.6\text{ns}$
- Assume program is memory bottlenecked & only performs one FLOP/memory access
- 1 FLOP every 21.6ns $\rightarrow 1/(21.6 * 10^{-9}) = 46.3 \text{ Megaflops}$

Parallel Version (2 threads)

- Cache Latency: 2ns
- DRAM latency: 100ns
- Remote DRAM latency: 400ns
- 90% cache hit, 8% DRAM, 2% Remote DRAM



Caching Example

Serial Version

- Cache Latency: 2ns
- DRAM latency: 100ns
- 80% hit rate
- Average access time:
 $.8 * 2 + .2 * 100 = 21.6\text{ns}$
- Assume program is memory bottlenecked & only performs one FLOP/memory access
- 1 FLOP every 21.6ns $\rightarrow 1/(21.6 * 1e-9) = 46.3$ Megaflops

Parallel Version (2 threads)

- Cache Latency: 2ns
- DRAM latency: 100ns
- Remote DRAM latency: 400ns
- 90% cache hit, 8% DRAM, 2% Remote DRAM

The cache hit rate can increase if (a) the total problem size is large enough to not fit in cache on 1 thread, but is closer when on 2 threads **and/or** (b) there is a highly irregular access pattern from memory that has improved locality on a larger number of threads



Caching Example

Serial Version

- Cache Latency: 2ns
- DRAM latency: 100ns
- 80% hit rate
- Average access time:
 $.8 * 2 + .2 * 100 = 21.6\text{ns}$
- Assume program is memory bottlenecked & only performs one FLOP/memory access
- 1 FLOP every 21.6ns $\rightarrow 1/(21.6 * 10^{-9}) = 46.3 \text{ Megaflops}$

Parallel Version (2 threads)

- Cache Latency: 2ns
- DRAM latency: 100ns
- Remote DRAM latency: 400ns
- 90% cache hit, 8% DRAM, 2% Remote DRAM



Caching Example

Serial Version

- Cache Latency: 2ns
- DRAM latency: 100ns
- 80% hit rate
- Average access time:
 $.8 * 2 + .2 * 100 = 21.6\text{ns}$
- Assume program is memory bottlenecked & only performs one FLOP/memory access
- 1 FLOP every 21.6ns $\rightarrow 1/(21.6 * 10^{-9}) = 46.3 \text{ Megaflops}$

Parallel Version (2 threads)

- Cache Latency: 2ns
- DRAM latency: 100ns
- Remote DRAM latency: 400ns
- 90% cache hit, 8% DRAM, 2% Remote DRAM
- Average access time:



Caching Example

Serial Version

- Cache Latency: 2ns
- DRAM latency: 100ns
- 80% hit rate
- Average access time:
 $.8 * 2 + .2 * 100 = 21.6\text{ns}$
- Assume program is memory bottlenecked & only performs one FLOP/memory access
- 1 FLOP every 21.6ns $\rightarrow 1 / (21.6 * 10^{-9}) = 46.3 \text{ Megaflops}$

Parallel Version (2 threads)

- Cache Latency: 2ns
- DRAM latency: 100ns
- Remote DRAM latency: 400ns
- 90% cache hit, 8% DRAM, 2% Remote DRAM
- Average access time:
 $.9 * 2 + .08 * 100 + .02 * 400 = 17.8\text{ns}$
- Assume program is memory bottlenecked & only performs one FLOP/memory access
- 1 FLOP every 17.8ns $\rightarrow 1 / (17.8 * 10^{-9}) = 56.18 \text{ Megaflops}$
- 2 threads $\rightarrow 112.36 \text{ MegaFlops}$



Caching Example

Serial Version

- Cache Latency: 2ns
- DRAM latency: 100ns
- 80% hit rate
- Average access time:
 $.8 * 2 + .2 * 100 = 21.6\text{ns}$
- Assume program is memory bottlenecked & only performs one FLOP/memory access
- 1 FLOP every 21.6ns $\rightarrow 1/(21.6 * 1e-9) = 46.3 \text{ Megaflops}$

More than 2x!


Parallel Version (2 threads)

- Cache Latency: 2ns
- DRAM latency: 100ns
- Remote DRAM latency: 400ns
- 90% cache hit, 8% DRAM, 2% Remote DRAM
- Average access time:
 $.9 * 2 + .08 * 100 + .02 * 400 = 17.8\text{ns}$
- Assume program is memory bottlenecked & only performs one FLOP/memory access
- 1 FLOP every 17.8ns $\rightarrow 1/(17.8 * 1e-9) = 56.18 \text{ Megaflops}$
- 2 threads $\rightarrow 112.36 \text{ MegaFlops}$




Exploratory Decomposition Example

- ❑ In search problems, sometimes the goal state can be expanded *much* more quickly in parallel
- ❑ Serial, depth-first search of the graph at right will take $14t_c$ where t_c is the cost of traversing one node ($T_s = 14t_c$)
- ❑ In parallel...



Exploratory Decomposition Example


- ❑ In search problems, sometimes the goal state can be expanded *much* more quickly in parallel
- ❑ Serial, depth-first search of the graph at right will take $14t_c$ where t_c is the cost of traversing one node ($T_s = 14t_c$)
- ❑ In parallel with $p=2$, we have $T_p = 5t_c$
- ❑ $S = ?$



Exploratory Decomposition Example

- ❑ In search problems, sometimes the goal state can be expanded *much* more quickly in parallel
- ❑ Serial, depth-first search of the graph at right will take $14t_c$ where t_c is the cost of traversing one node ($T_s = 14t_c$)
- ❑ In parallel with $p=2$, we have $T_p = 5t_c$
- ❑ $S = 14/5 > 2$

Superlinear



Lecture Overview

- ❑ Background
- ❑ Overheads
- ❑ **Definitions**
 - Serial Runtime/Parallel Runtime/Parallel Overhead
 - Speedup
 - **Efficiency**
 - Cost
- ❑ Granularity



Efficiency

- ❑ Measures the fraction of time a processor is usefully employed
- ❑ One of the more important metrics for allocating resources efficiently
- ❑ The most efficient parallel algorithms have E close to 1
- ❑ E should usually lie in $(0, 1)$

$$E = S/p$$



Efficiency (Adding n Numbers)

$$S = \Theta(n \log n)$$

$$p = n$$


$$E = \Theta(n \log n) / n = \Theta(1 \log n)$$



Efficiency (Edge Detection Example)


Edge-Detection

- ❑ Compute edge feature map
- ❑ Input is $n \times n$ image
- ❑ With 3×3 kernels, $9n^2$ computations (assume each computation takes t_c seconds)
- ❑ Assume column-wise distribution as in right-most figure
- ❑ Let t_s and t_w be the startup time and per-word transfer time for communication, respectively
- ❑ Let p be the total number of processes



-1	0	1
-2	0	2
-1	0	1

-1	-2	1
0	0	0
-1	2	1





Efficiency (Edge Detection Example)

$$T_s =$$

$$T_p =$$

$$S =$$

$$E =$$





Efficiency (Edge Detection Example)

$$T_s = 9t_c n^2$$

$$T_p =$$

$$S =$$

$$E =$$




Efficiency (Edge Detection Example)

$$T_s = 9t_c n^2$$

$$T_p = 9t_c \frac{n^2}{p} + 2(t_s + t_w n)$$


$$S =$$

$$E =$$



-1	0	1
-2	0	2
-1	0	1

-1	-2	1
0	0	0
-1	2	1




Efficiency (Edge Detection Example)

$$T_s = 9t_c n^2$$

$$T_p = 9t_c \frac{n^2}{p} + 2(t_s + t_w n)$$


$$S = \frac{9t_c n^2}{9t_c \frac{n^2}{p} + 2(t_s + t_w n)}$$

$$E =$$



-1	0	1
-2	0	2
-1	0	1

-1	-2	1
0	0	0
-1	2	1




Efficiency (Edge Detection Example)

$$T_s = 9t_c n^2$$

$$T_p = 9t_c \frac{n^2}{p} + 2(t_s + t_w n)$$


$$S = \frac{9t_c n^2}{9t_c \frac{n^2}{p} + 2(t_s + t_w n)}$$

$$E = \frac{1}{1 + \frac{2p(t_s + t_w n)}{9t_c n^2}}$$



-1	0	1
-2	0	2
-1	0	1

-1	-2	1
0	0	0
-1	2	1



Lecture Overview

- ❑ Background
- ❑ Overheads
- ❑ **Definitions**
 - Serial Runtime/Parallel Runtime/Parallel Overhead
 - Speedup
 - Efficiency
 - **Cost**
- ❑ Granularity



Cost

- ❑ Sum of time spent on the program across all processes
- ❑ We can reformulate efficiency as
$$E = S/p = (T_s/T_p)/p = T_s/(pT_p) = T_s/\text{Cost}$$
- ❑ Also sometimes called *work* or *processor-time product*
- ❑ A program is *cost-optimal* if it has the same big- Θ complexity as a function of input size as the fastest known sequential algorithm on a single processing element

$$\text{Cost} = pT_p$$



Cost of adding n numbers

□ Recall from our adding example that

- $p=n$
- $T_p = \log n$
- $T_s = n$

□ Cost = $\Theta(n \log n)$

□ Is this cost-optimal?



Cost of adding n numbers

□ Recall from our adding example that

- $p=n$
- $T_p = \log n$
- $T_s = n$

□ Cost = $\Theta(n \log n)$

□ Is this cost-optimal?

- No - cost grows asymptotically faster than the serial time of execution



Lecture Overview

- ❑ Background
- ❑ Overheads
- ❑ Definitions
 - Serial Runtime/Parallel Runtime/Parallel Overhead
 - Speedup
 - Efficiency
 - Cost
- ❑ **Granularity**




Granularity

- ❑ Using fewer processes than the maximum possible is often more practical given the overheads introduced by idling/communication
- ❑ If we choose an appropriate level of granularity, we can create cost-optimal programs with fewer processes




Cost Optimal Addition




(a)

Cost Optimal Addition




(a)




(b)


Cost Optimal Addition



(a)




(b)




(c)


Cost Optimal Addition



(a)



(b)



(c)



(d)




Cost Optimal Addition


$$T_p = ?$$

$$\text{Cost} = ?$$


$$T_s = ?$$



(a)



(b)



(c)



(d)




Cost Optimal Addition


$$T_p = \Theta(n/p + \log p)$$

$$\text{Cost} = \Theta(n + p \log p)$$


$$T_s = \Theta(n)$$



(a)



(b)



(c)



(d)




Cost Optimal Addition

$$T_p = \Theta(n/p + \log p)$$


$$\text{Cost} = \Theta(n + p \log p)$$

$$T_s = \Theta(n)$$


As long as $n = \Omega(p \log p)$, this program is cost-optimal



(a)



(b)



(c)



(d)

