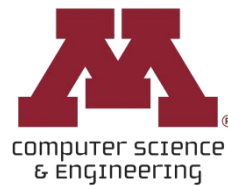


CSCI 5451: Introduction to Parallel Computing

Lecture 17: Advanced Analytical Modeling



UNIVERSITY OF MINNESOTA
Driven to Discover®

Lecture Overview

- ❑ Review of Asymptotic Function Growth Rates (Big-O, Ω , Θ)
- ❑ Scalability and the Definition of Work (W)
- ❑ Isoefficiency
- ❑ Optimal Computation Time
- ❑ Comparing Different Algorithms with Parallel Metrics



Lecture Overview

- ❑ **Review of Asymptotic Function Growth Rates (Big-O, Ω , Θ)**
- ❑ Scalability and the Definition of Work (W)
- ❑ Isoefficiency
- ❑ Optimal Computation Time
- ❑ Comparing Different Algorithms with Parallel Metrics



Review of Big-O, Ω , Θ

□ Big O (O)

- o Upper bound on function of interest, f
- o $f(n) = O(g(n)) \Rightarrow c \cdot g(n)$ is an upper bound for $f(n)$, $n > n_0$ for some c, n_0

□ Big Ω (Omega)

- o Lower bound on function of interest, f
- o $f(n) = \Omega(g(n)) \Rightarrow c \cdot g(n)$ is a lower bound for $f(n)$, $n > n_0$ for some c, n_0


□ Big Θ (Theta)

- o Tight bound on function of interest, f
- o $f(n) = \Theta(g(n)) \Rightarrow c_1 \cdot g(n)$ is a lower bound for $f(n)$, $c_2 \cdot g(n)$ an upper bound for $f(n)$, $n > n_0$ some c_1, c_2, n_0



Review of Big-O, Ω , Θ

- Big O (O)
 - Upper bound on function of interest, f
 - $f(n) = O(g(n)) \Rightarrow c \cdot g(n)$ is an upper bound for $f(n)$, $n > n_0$ for some c, n_0
- Big Ω (Omega)
 - Lower bound on function of interest, f
 - $f(n) = \Omega(g(n)) \Rightarrow c \cdot g(n)$ is a lower bound for $f(n)$, $n > n_0$ for some c, n_0
- Big Θ (Theta)
 - Tight bound on function of interest, f
 - $f(n) = \Theta(g(n)) \Rightarrow c_1 \cdot g(n)$ is a lower bound for $f(n)$, $c_2 \cdot g(n)$ an upper bound for $f(n)$, $n > n_0$ some c_1, c_2, n_0



Review of Big-O, Ω , Θ

Big O (O)


- Upper bound on function of interest, f
- $f(n) = O(g(n)) \Rightarrow c \cdot g(n)$ is an upper bound for $f(n)$, $n > n_0$ for some c, n_0

Big Ω (Omega)

- Lower bound on function of interest, f
- $f(n) = \Omega(g(n)) \Rightarrow c \cdot g(n)$ is a lower bound for $f(n)$, $n > n_0$ for some c, n_0

Big Θ (Theta)

- Tight bound on function of interest, f
- $f(n) = \Theta(g(n)) \Rightarrow c_1 \cdot g(n)$ is a lower bound for $f(n)$, $c_2 \cdot g(n)$ an upper bound for $f(n)$, $n > n_0$ some c_1, c_2, n_0



Review of Big-O, Ω , Θ

Big O (O)


- Upper bound on function of interest, f
- $f(n) = O(g(n)) \Rightarrow c \cdot g(n)$ is an upper bound for $f(n)$, $n > n_0$ for some c, n_0

Big Ω (Omega)

- Lower bound on function of interest, f
- $f(n) = \Omega(g(n)) \Rightarrow c \cdot g(n)$ is a lower bound for $f(n)$, $n > n_0$ for some c, n_0

Big Θ (Theta)

- Tight bound on function of interest, f
- $f(n) = \Theta(g(n)) \Rightarrow c_1 \cdot g(n)$ is a lower bound for $f(n)$, $c_2 \cdot g(n)$ an upper bound for $f(n)$, $n > n_0$ some c_1, c_2, n_0




Lecture Overview

- ❑ Review of Asymptotic Function Growth Rates (Big-O, Ω , Θ)
- ❑ **Scalability and the Definition of Work (W)**
- ❑ Isoefficiency
- ❑ Optimal Computation Time
- ❑ Comparing Different Algorithms with Parallel Metrics




Scalability

- ❑ Ultimately, we want to figure out how to go from programs that complete in minutes on dozens of processes to programs which take weeks on thousands of processes
- ❑ To do this, we must estimate how efficient our programs will be, as we will not be able to empirically determine their runtimes at such enormous scale



Scalability

- ❑ We have to consider how the efficiency, parallel runtime & speedups of our program change as we vary the size of the input problem and the number of processing elements
- ❑ [Right] Speedup comparison of several different parallel algorithms for Fast Fourier Transform



Revisiting Efficiency

We can rewrite efficiency in several different ways terms of T_o and T_s

$$E = \frac{S}{P}$$

$$E = \frac{T_s}{p T_p}$$

$$E = \frac{1}{1 + \left(\frac{T_o}{T_s} \right)}$$



Revisiting Efficiency

We can rewrite efficiency in several different ways terms of T_o and T_s

$$E = \frac{S}{P}$$

$$E = \frac{T_s}{p T_p}$$

$$E = \frac{1}{1 + \left(\frac{T_o}{T_s} \right)}$$

Increasing the number of processors

- T_o is an increasing function of p .
- Every program will have some serial portion taking up time t_{serial}
- All other $p-1$ processes will idle in this time
- Thus, in the best case, $T_o = \Omega(p)$
- Other costs (communication, excess computation) can further increase T_o
- If we fix the problem size, then T_s will remain constant
- **Increasing p will decrease efficiency E**



Revisiting Efficiency

We can rewrite efficiency in several different ways terms of T_o and T_s

$$E = \frac{S}{P}$$

$$E = \frac{T_s}{p T_p}$$

$$E = \frac{1}{1 + \left(\frac{T_o}{T_s}\right)}$$

Increasing the problem size

- Increasing the problem size will increase the size of T_s
- Increasing the problem size **can** increase the overhead T_o
- **A good parallel algorithm** should have an efficiency that increases as the problem size grows.
- If your algorithm does not exhibit this, it will not scale to larger problems



Revisiting Addition of N Numbers

Let's find some non-asymptotic metrics
when summing N numbers

- ❑ Sum n numbers on p processors
- ❑ Assume each addition takes 1 unit time
- ❑ Each communication takes 1 unit time
- ❑ What are the values of
 - $T_s = ?$
 - $T_p = ?$
 - $T_o = ?$
 - $S = ?$
 - $E = ?$



Revisiting Addition of N Numbers

Let's find some non-asymptotic metrics when summing N numbers

- ❑ Sum n numbers on p processors
- ❑ Assume each addition takes 1 unit time
- ❑ Each communication takes 1 unit time
- ❑ What are the values of
 - $T_s = ?$
 - $T_p = ?$
 - $T_o = ?$
 - $S = ?$
 - $E = ?$

$$T_s = n$$

$$T_p = \frac{n}{p} + 2 \log p$$

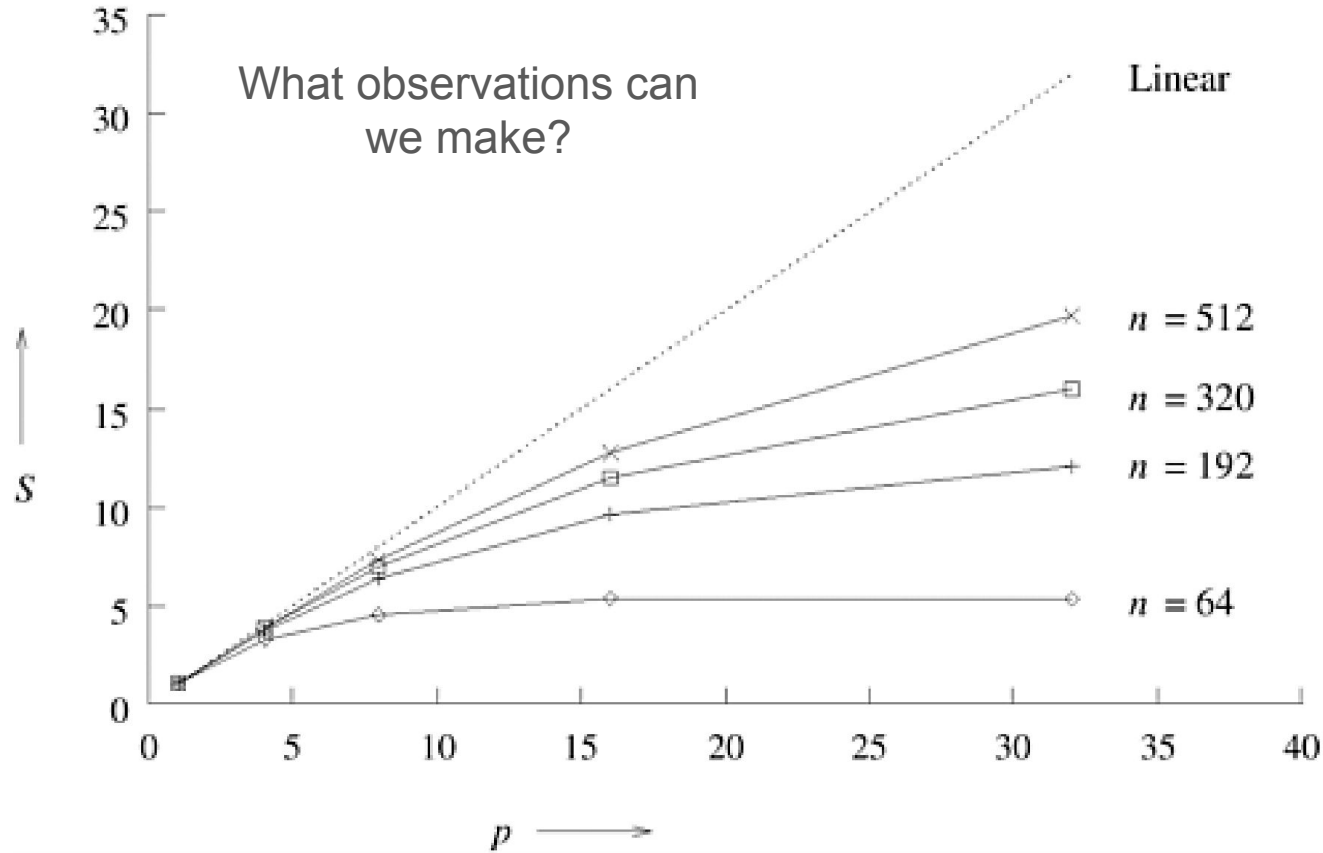
$$T_o = 2p \log p$$

$$S = \frac{n}{\frac{n}{p} + 2 \log p}$$

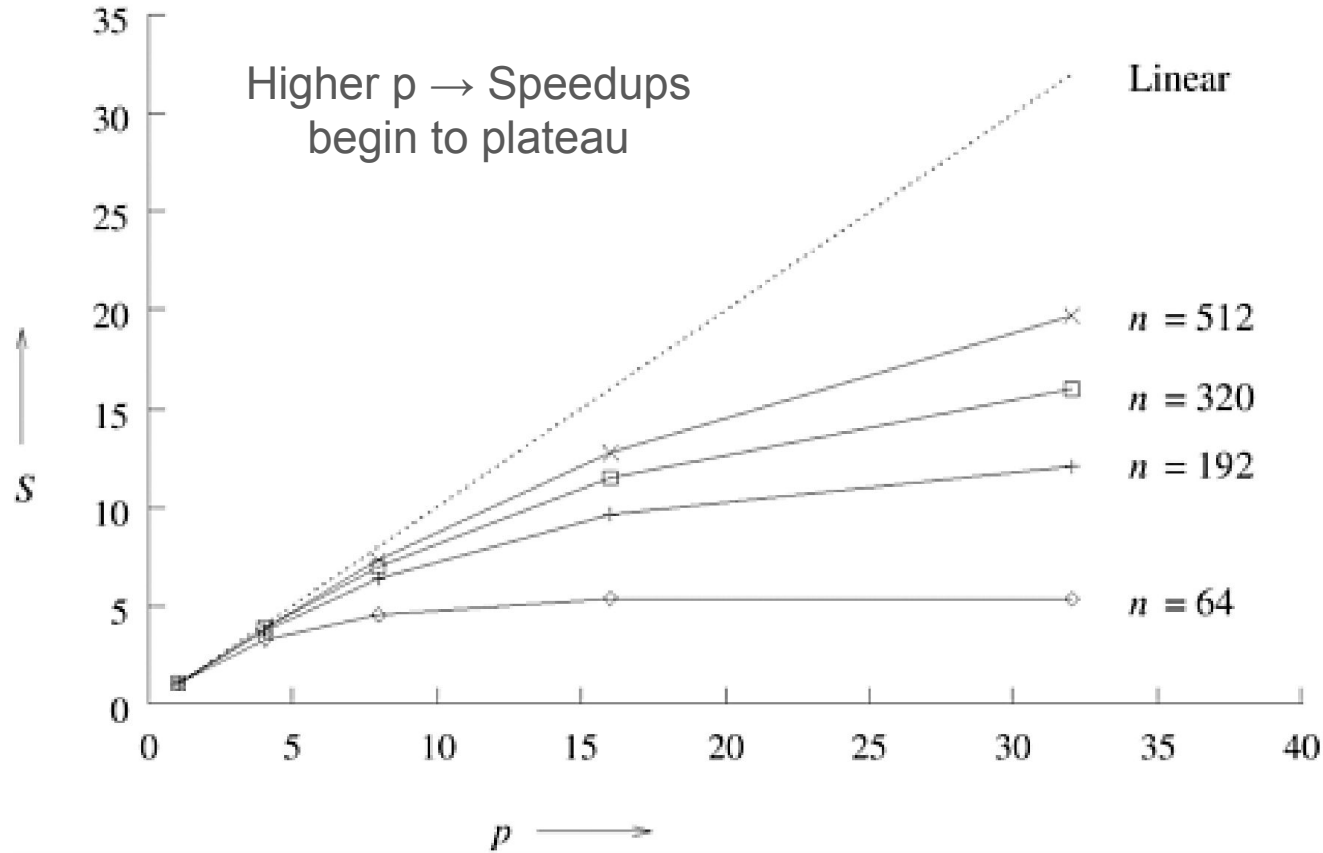
$$E = \frac{1}{1 + \left(\frac{2p \log p}{n} \right)}$$



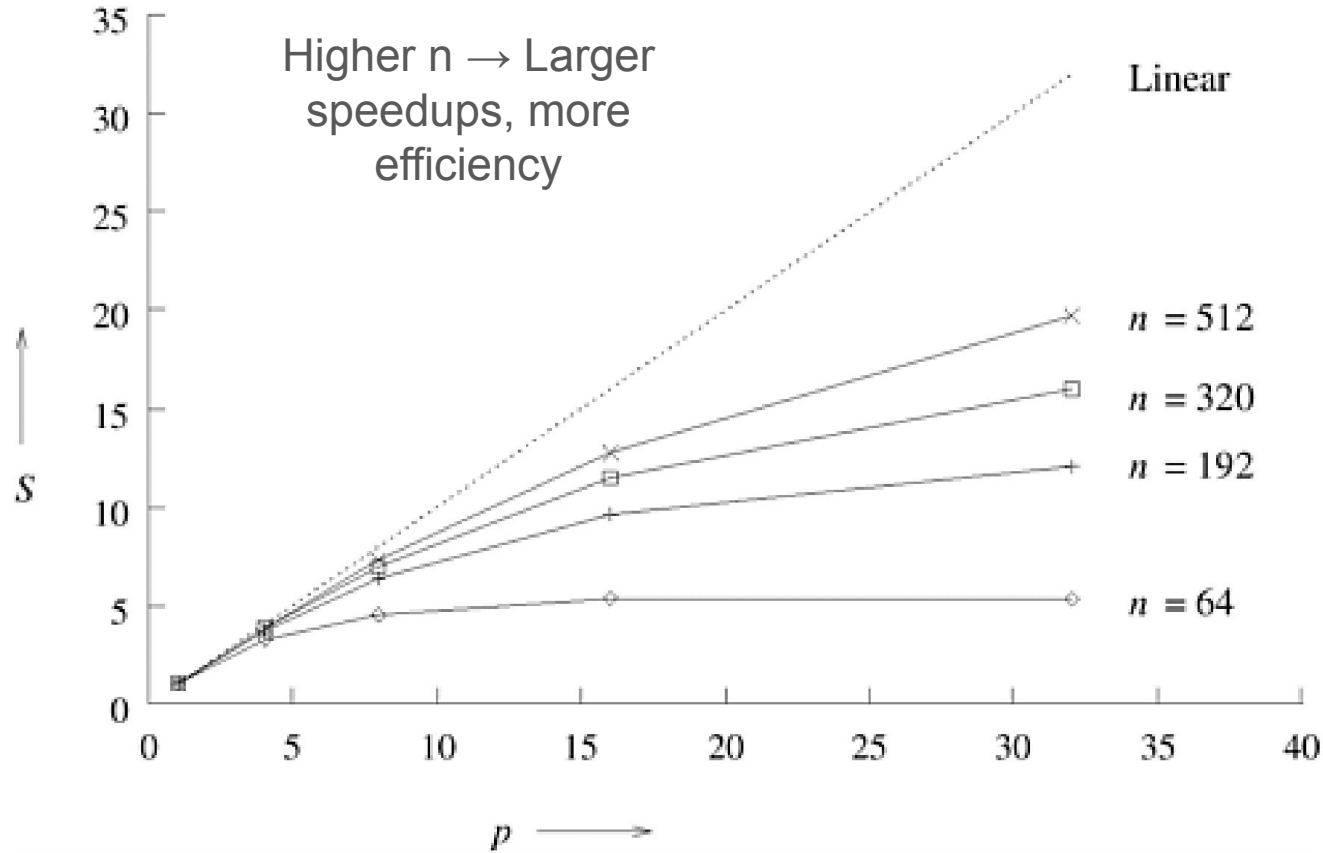
Revisiting Addition of N Numbers



Revisiting Addition of N Numbers



Revisiting Addition of N Numbers



Revisiting Addition of N Numbers

n	$p = 1$	$p = 4$	$p = 8$	$p = 16$	$p = 32$
64	1.0	0.80	0.57	0.33	0.17
192	1.0	0.92	0.80	0.60	0.38
320	1.0	0.95	0.87	0.71	0.50
512	1.0	0.97	0.91	0.80	0.62



Revisiting Addition of N Numbers

For a given number of processors, p , and efficiency, E , we can determine the problem size needed (in this case n)

n	$p = 1$	$p = 4$	$p = 8$	$p = 16$	$p = 32$
64	1.0	0.80	0.57	0.33	0.17
192	1.0	0.92	0.80	0.60	0.38
320	1.0	0.95	0.87	0.71	0.50
512	1.0	0.97	0.91	0.80	0.62



Revisiting Addition of N Numbers

IsoEfficiency

For a given number of processors, p , and efficiency, E , we can determine the problem size needed (in this case n)


n	$p = 1$	$p = 4$	$p = 8$	$p = 16$	$p = 32$
64	1.0	0.80	0.57	0.33	0.17
192	1.0	0.92	0.80	0.60	0.38
320	1.0	0.95	0.87	0.71	0.50
512	1.0	0.97	0.91	0.80	0.62



Revisiting Addition of N Numbers

General Observation #1:


With a fixed problem size, increasing the amount of processors **decreases** efficiency



Revisiting Addition of N Numbers


General Observation #2:

With a fixed number of processors, increasing the problem size **often increases** efficiency




Revisiting Addition of N Numbers

What would be more ideal versions of the below curves? In other words, how would we want efficiency to change w.r. to W, p ?



Revisiting Addition of N Numbers

What would be more ideal versions of the below curves? In other words, how would we want efficiency to change w.r. to W, p ?



Formally Defining Work W

- ❑ Up to this point, we have defined problem sizes & work in informal terms
- ❑ Naive way to define this → amount of input data n
- ❑ Problem with this?



Formally Defining Work W

- ❑ Up to this point, we have defined problem sizes & work in informal terms
- ❑ Naive way to define this → amount of input data n
- ❑ Problem with this?
 - Doesn't reflect runtime
 - Consider multiplying two $n \times n$ input matrices
 - Input is n^2 , runtime is n^3
- ❑ Problem Size (or Work) should reflect the serial execution complexity



Formally Defining Work W

- ❑ Up to this point, we have defined problem sizes & work in informal terms
- ❑ Naive way to define this → amount of input data n
- ❑ Problem with this?
 - Doesn't reflect runtime
 - Consider multiplying two $n \times n$ input matrices
 - Input is n^2 , runtime is n^3
- ❑ Problem Size (or Work) should reflect the serial execution complexity

Precise definition:

The problem size W is defined as the number of computation steps in the **best** sequential algorithm to solve the problem on a single processing element



Redefining Metrics in Terms of W

$$T_P = \frac{W + T_o(W, p)}{p}$$

$$S = \frac{W}{T_P} = \frac{Wp}{W + T_o(W, p)}$$

$$E = \frac{S}{p} = \frac{W}{W + T_o(W, p)} = \frac{1}{1 + \frac{T_o(W, p)}{W}}$$



Deriving W as a function of p, T_o, E

$$E = \frac{1}{1 + \frac{T_o(W, p)}{W}},$$

$$\frac{T_o(W, p)}{W} = \frac{1 - E}{E},$$

$$W = \frac{E}{1 - E} T_o(W, p),$$

$$W = K T_o(W, p)$$



Deriving W as a function of p, T_o, E

$$E = \frac{1}{1 + \frac{T_o(W, p)}{W}},$$

$$\frac{T_o(W, p)}{W} = \frac{1 - E}{E},$$

$$W = \frac{E}{1 - E} T_o(W, p),$$

$$W = K T_o(W, p)$$

For a constant efficiency, E , we can define the Work, W , as a function of the overhead T_o .



Deriving W as a function of p, T_o, E

$$E = \frac{1}{1 + \frac{T_o(W, p)}{W}},$$

$$\frac{T_o(W, p)}{W} = \frac{1 - E}{E},$$

$$W = \frac{E}{1 - E} T_o(W, p),$$

$$W = K T_o(W, p)$$

For a constant efficiency, E , we can define the Work, W , as a function of the overhead

IsoEfficiency



Lecture Overview

- ❑ Review of Asymptotic Function Growth Rates (Big-O, Ω , Θ)
- ❑ Scalability and the Definition of Work (W)
- ❑ **Isoefficiency**
- ❑ Optimal Computation Time
- ❑ Comparing Different Algorithms with Parallel Metrics



Isoefficiency

- This tells us if we change the number of processors, by what amount must we change the work to keep efficiency constant
- Usefulness → Understanding how well your program will scale when going from smaller problems to bigger ones

$$W = K T_o(W, p)$$



Comparing Isoefficiency Functions

For each of the two overhead functions at right

□ What are their isoefficiency functions?

$$(A) \quad T_o = \log p$$

$$(B) \quad T_o = p$$

□ Which has a better isoefficiency function and why?



Comparing Isoefficiency Functions

For each of the two overhead functions at right

□ What are their isoefficiency functions?

$$(A) \quad W = K \log p$$

$$(B) \quad W = Kp$$

$$(A) \quad T_o = \log p$$

$$(B) \quad T_o = p$$

□ Which has a better isoefficiency function and why?

- A has better scaling. If we increase p from 2 to 64, then maintaining efficiency on (A), requires increasing the work by 6x, whereas on (B) we must increase work by 32x



Isoefficiency (Adding n numbers)

- ❑ The overhead function for adding n numbers is

$$T_o = 2p \log p$$

- ❑ The isoefficiency function is

$$W = 2kp \log p$$

- ❑ Asymptotically this is $\Theta(p \log p)$
- ❑ In other words, if we increase the number of processors from p to p' , how much will we need to increase the amount of work by to maintain constant efficiency?



Isoefficiency (Adding n numbers)

- ❑ The overhead function for adding n numbers is

$$T_o = 2p \log p$$

- ❑ The isoefficiency function is

$$W = 2kp \log p$$

- ❑ Asymptotically this is $\Theta(p \log p)$
- ❑ In other words, if we increase the number of processors from p to p' , how much will we need to increase the amount of work by to maintain constant efficiency?

$$\frac{p' \log p'}{p \log p}$$



Isoefficiency (Complex Example)

- ❑ In the simple example of adding n numbers, the overhead function is simple & only a function of p
- ❑ Most overhead functions will be a function of both p and W
- ❑ It can be more difficult to define a closed form expression of W
- ❑ As such, we instead compute the isoefficiency function for each separate component - only considering the function that grows largest as a function of p


$$T_o = p^{3/2} + p^{3/4}W^{3/4}$$



Isoefficiency (Complex Example)

- ❑ In the simple example of adding n numbers, the overhead function is simple & only a function of p
- ❑ Most overhead functions will be a function of both p and W
- ❑ It can be more difficult to define a closed form expression of W
- ❑ As such, we instead compute the isoefficiency function for each separate component - only considering the function that grows largest as a function of p

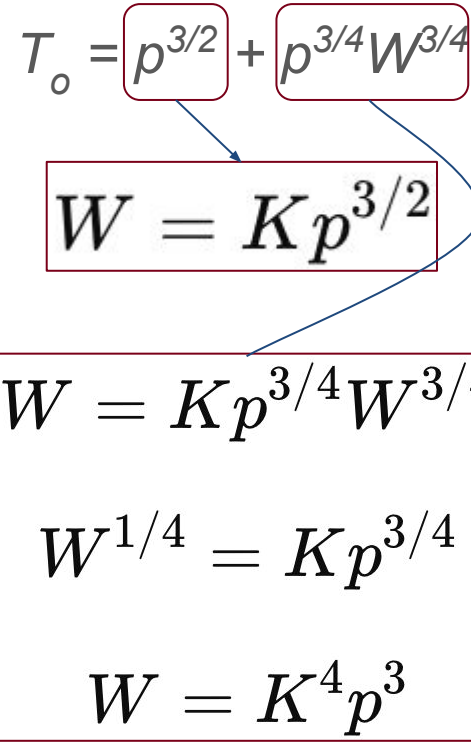
$$T_o = p^{3/2} + p^{3/4}W^{3/4}$$


$$W = Kp^{3/2}$$



Isoefficiency (Complex Example)

- ❑ In the simple example of adding n numbers, the overhead function is simple & only a function of p
- ❑ Most overhead functions will be a function of both p and W
- ❑ It can be more difficult to define a closed form expression of W
- ❑ As such, we instead compute the isoefficiency function for each separate component - only considering the function that grows largest as a function of p

$$T_o = p^{3/2} + p^{3/4} W^{3/4}$$


$$W = K p^{3/2}$$

$$W = K p^{3/4} W^{3/4}$$

$$W^{1/4} = K p^{3/4}$$


$$W = K^4 p^3$$



Isoefficiency (Complex Example)

Take largest

- ❑ In the simple example of adding n numbers, the overhead function is simple & only a function of p
- ❑ Most overhead functions will be a function of both p and W
- ❑ It can be more difficult to define a closed form expression of W
- ❑ As such, we instead compute the isoefficiency function for each separate component - only considering the function that grows largest as a function of p

$$T_o = p^{3/2} + p^{3/4}W^{3/4}$$


$$W = Kp^{3/4}W^{3/4}$$

$$W^{1/4} = Kp^{3/4}$$

$$W = K^4 p^3$$



Lecture Overview

- ❑ Review of Asymptotic Function Growth Rates (Big-O, Ω , Θ)
- ❑ Scalability and the Definition of Work (W)
- ❑ Isoefficiency
- ❑ **Optimal Computation Time**
- ❑ Comparing Different Algorithms with Parallel Metrics



Optimal Computation Time

- ❑ A common question we may ask is what the minimum parallel runtime (T_p) is and what the minimum number of processors are needed to achieve this minimum runtime?
- ❑ How can we achieve this?



Optimal Computation Time

- A common question we may ask is what the minimum parallel runtime (T_p) is and what the minimum number of processors are needed to achieve this minimum runtime?
- How can we achieve this?
 - Take the derivative of T_p with respect to p
 - Find p_o where this derivative is 0
 - Use this p as the input to T_p

$$\frac{d}{dp} T_P = 0$$



Optimal Computation Time (Adding n numbers)

$$T_p = \frac{n}{p} + 2 \log p$$



Optimal Computation Time (Adding n numbers)

$$T_p = \frac{n}{p} + 2 \log p$$

$$\frac{d}{dp} T_P = -\frac{n}{p_0^2} + \frac{2}{p_0} = 0,$$

$$-n + 2p_0 = 0,$$

$$p_0 = \frac{n}{2},$$

$$T_P^{min} = n \log n - n \log 2 + 2$$



Optimal Computation Time (Complex Example)

$$T_P = \frac{W}{p} + p^{1/2} + \frac{W^{3/4}}{p^{1/4}}$$



Optimal Computation Time (Complex Example)

$$\frac{d}{dp}T_P = -\frac{W}{p_0^2} + \frac{1}{2p_0^{1/2}} - \frac{W^{3/4}}{4p_0^{5/4}} = 0,$$

$$T_P = \frac{W}{p} + p^{1/2} + \frac{W^{3/4}}{p^{1/4}}$$



Optimal Computation Time (Complex Example)

$$\frac{d}{dp}T_P = -\frac{W}{p_0^2} + \frac{1}{2p_0^{1/2}} - \frac{W^{3/4}}{4p_0^{5/4}} = 0,$$

$$-W + \frac{1}{2}p_0^{3/2} - \frac{1}{4}W^{3/4}p_0^{3/4} = 0,$$

$$T_P = \frac{W}{p} + p^{1/2} + \frac{W^{3/4}}{p^{1/4}}$$



Optimal Computation Time (Complex Example)

$$\frac{d}{dp}T_P = -\frac{W}{p_0^2} + \frac{1}{2p_0^{1/2}} - \frac{W^{3/4}}{4p_0^{5/4}} = 0,$$

$$-W + \frac{1}{2}p_0^{3/2} - \frac{1}{4}W^{3/4}p_0^{3/4} = 0,$$

$$T_P = \frac{W}{p} + p^{1/2} + \frac{W^{3/4}}{p^{1/4}}$$

$$p_0^{3/4} = \frac{1}{4}W^{3/4} \pm \left(\frac{1}{16}W^{3/2} + 2W \right)^{1/2}$$



Optimal Computation Time (Complex Example)

$$\frac{d}{dp}T_P = -\frac{W}{p_0^2} + \frac{1}{2p_0^{1/2}} - \frac{W^{3/4}}{4p_0^{5/4}} = 0,$$

$$-W + \frac{1}{2}p_0^{3/2} - \frac{1}{4}W^{3/4}p_0^{3/4} = 0,$$

$$T_P = \frac{W}{p} + p^{1/2} + \frac{W^{3/4}}{p^{1/4}}$$

$$\begin{aligned} p_0^{3/4} &= \frac{1}{4}W^{3/4} \pm \left(\frac{1}{16}W^{3/2} + 2W \right)^{1/2} \\ &= \Theta(W^{3/4}), \end{aligned}$$



Optimal Computation Time (Complex Example)

$$\frac{d}{dp}T_P = -\frac{W}{p_0^2} + \frac{1}{2p_0^{1/2}} - \frac{W^{3/4}}{4p_0^{5/4}} = 0,$$

$$-W + \frac{1}{2}p_0^{3/2} - \frac{1}{4}W^{3/4}p_0^{3/4} = 0,$$

$$T_P = \frac{W}{p} + p^{1/2} + \frac{W^{3/4}}{p^{1/4}}$$

$$p_0^{3/4} = \frac{1}{4}W^{3/4} \pm \left(\frac{1}{16}W^{3/2} + 2W\right)^{1/2}$$

$$= \Theta(W^{3/4}),$$

$$p_0 = \Theta(W),$$



Optimal Computation Time (Complex Example)

$$\frac{d}{dp}T_P = -\frac{W}{p_0^2} + \frac{1}{2p_0^{1/2}} - \frac{W^{3/4}}{4p_0^{5/4}} = 0,$$

$$-W + \frac{1}{2}p_0^{3/2} - \frac{1}{4}W^{3/4}p_0^{3/4} = 0,$$

$$T_P = \frac{W}{p} + p^{1/2} + \frac{W^{3/4}}{p^{1/4}}$$

$$p_0^{3/4} = \frac{1}{4}W^{3/4} \pm \left(\frac{1}{16}W^{3/2} + 2W\right)^{1/2}$$

$$= \Theta(W^{3/4}),$$

$$p_0 = \Theta(W),$$

$$T_p^{min} = \Theta(W^{1/2}).$$





Lecture Overview

- ❑ Review of Asymptotic Function Growth Rates (Big-O, Ω , Θ)
- ❑ Scalability and the Definition of Work (W)
- ❑ Isoefficiency
- ❑ Optimal Computation Time
- ❑ **Comparing Different Algorithms with Parallel Metrics**



General Usefulness of these Metrics

- ❑ These metrics (E , S , isoefficiency, T_o , etc.) are useful tools for comparing different parallel approaches - we can determine which approaches are more scalable
- ❑ They expand on runtime analysis to show which algorithms will make **better** use of the processors we utilize, rather than just how much faster the algorithm itself will run



General Usefulness of these Metrics

Consider the below set of algorithms

Algorithm	A1	A2	A3	A4
p	n^2	$\log n$	n	\sqrt{n}
T_P	1	n	\sqrt{n}	$\sqrt{n} \log n$
S	$n \log n$	$\log n$	$\sqrt{n} \log n$	\sqrt{n}
E	$\frac{\log n}{n}$	1	$\frac{\log n}{\sqrt{n}}$	1
pT_P	n^2	$n \log n$	$n^{1.5}$	$n \log n$



General Usefulness of these Metrics

Which is fastest?

Algorithm	A1	A2	A3	A4
p	n^2	$\log n$	n	\sqrt{n}
T_P	1	n	\sqrt{n}	$\sqrt{n} \log n$
S	$n \log n$	$\log n$	$\sqrt{n} \log n$	\sqrt{n}
E	$\frac{\log n}{n}$	1	$\frac{\log n}{\sqrt{n}}$	1
pT_P	n^2	$n \log n$	$n^{1.5}$	$n \log n$



General Usefulness of these Metrics

Which is fastest? (A1)

Algorithm	A1	A2	A3	A4
p	n^2	$\log n$	n	\sqrt{n}
T_P	1	n	\sqrt{n}	$\sqrt{n} \log n$
S	$n \log n$	$\log n$	$\sqrt{n} \log n$	\sqrt{n}
E	$\frac{\log n}{n}$	1	$\frac{\log n}{\sqrt{n}}$	1
pT_P	n^2	$n \log n$	$n^{1.5}$	$n \log n$



General Usefulness of these Metrics

Should we use A1?

Algorithm	A1	A2	A3	A4
p	n^2	$\log n$	n	\sqrt{n}
T_P	1	n	\sqrt{n}	$\sqrt{n} \log n$
S	$n \log n$	$\log n$	$\sqrt{n} \log n$	\sqrt{n}
E	$\frac{\log n}{n}$	1	$\frac{\log n}{\sqrt{n}}$	1
pT_P	n^2	$n \log n$	$n^{1.5}$	$n \log n$



General Usefulness of these Metrics

Should we use A1? No -
requiring p processors is not
practical

Algorithm	A1	A2	A3	A4
p	n^2	$\log n$	n	\sqrt{n}
T_P	1	n	\sqrt{n}	$\sqrt{n} \log n$
S	$n \log n$	$\log n$	$\sqrt{n} \log n$	\sqrt{n}
E	$\frac{\log n}{n}$	1	$\frac{\log n}{\sqrt{n}}$	1
pT_P	n^2	$n \log n$	$n^{1.5}$	$n \log n$



General Usefulness of these Metrics

Which are cost optimal?

Algorithm	A1	A2	A3	A4
p	n^2	$\log n$	n	\sqrt{n}
T_P	1	n	\sqrt{n}	$\sqrt{n} \log n$
S	$n \log n$	$\log n$	$\sqrt{n} \log n$	\sqrt{n}
E	$\frac{\log n}{n}$	1	$\frac{\log n}{\sqrt{n}}$	1
pT_P	n^2	$n \log n$	$n^{1.5}$	$n \log n$



General Usefulness of these Metrics

Which are cost optimal? (A2, A4)

Algorithm	A1	A2	A3	A4
p	n^2	$\log n$	n	\sqrt{n}
T_P	1	n	\sqrt{n}	$\sqrt{n} \log n$
S	$n \log n$	$\log n$	$\sqrt{n} \log n$	\sqrt{n}
E	$\frac{\log n}{n}$	1	$\frac{\log n}{\sqrt{n}}$	1
pT_P	n^2	$n \log n$	$n^{1.5}$	$n \log n$



General Usefulness of these Metrics

What should we use and why?

Algorithm	A1	A2	A3	A4
p	n^2	$\log n$	n	\sqrt{n}
T_P	1	n	\sqrt{n}	$\sqrt{n} \log n$
S	$n \log n$	$\log n$	$\sqrt{n} \log n$	\sqrt{n}
E	$\frac{\log n}{n}$	1	$\frac{\log n}{\sqrt{n}}$	1
pT_P	n^2	$n \log n$	$n^{1.5}$	$n \log n$

