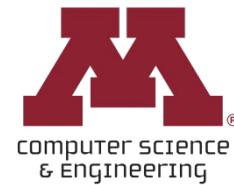


# CSCI 5451: Introduction to Parallel Computing

## Lecture 27: Practical Extensions & Real-World CUDA



# Overview

- ❑ Advanced CUDA Constructs
- ❑ Libraries
- ❑ Examples of Parallelism in Practice
- ❑ Tooling
- ❑ Alternative Hardware Architectures

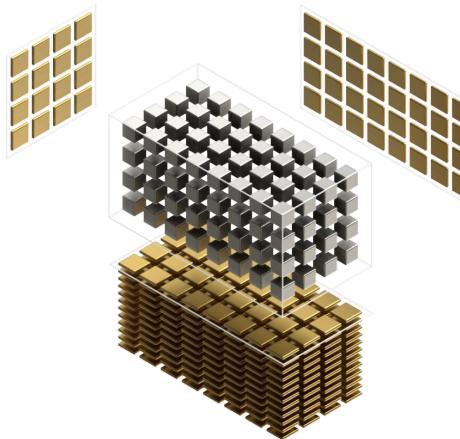
# Overview

- ❑ Advanced CUDA Constructs
- ❑ Libraries
- ❑ Examples of Parallelism in Practice
- ❑ Tooling
- ❑ Alternative Hardware Architectures



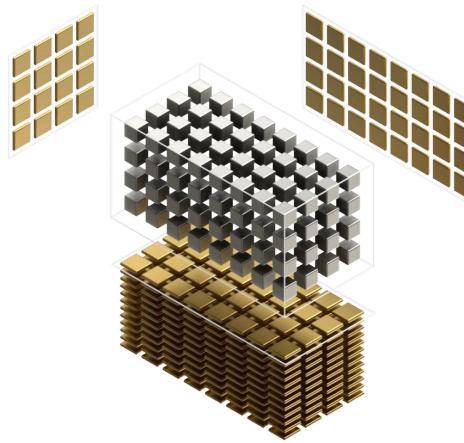
# Tensor Cores

- ❑ Specialized matrix-multiply-accumulate (MMA) units inside NVIDIA GPUs designed to accelerate dense linear algebra.
- ❑ Operate on small matrix tiles (e.g.,  $16 \times 16 \times 16$ ) at extremely high throughput.
- ❑ Perform  $\mathbf{D} = \mathbf{A} \times \mathbf{B} + \mathbf{C}$  in one hardware instruction (mma.sync).
- ❑ Accept low-precision inputs (FP16, BF16, TF32, INT8, FP8 on newer GPUs) and accumulate into higher precision (FP16/FP32).
- ❑ Integrated into each SM; a warp cooperatively issues a tensor operation.



# Tensor Cores

- ❑ A single warp (32 threads) collectively loads tile fragments of A, B, and C.
- ❑ The warp calls an MMA instruction → hardware performs the full matrix multiply & accumulate on the tile.  
Each instruction computes many FMA ops in parallel (e.g., 256 FMAs for a  $16 \times 16 \times 16$  tile).
- ❑ Outputs are stored into per-warp registers and then written to memory.
- ❑ The compiler maps WMMA/CUTLASS calls to optimized Tensor Core instructions.



# Tensor Cores

```
__global__ void wmma_example(const half *A, const half *B, float *C) {
    // Each warp computes one 16x16 tile
    wmma::fragment<wmma::matrix_a, 16,16,16, half, wmma::row_major> a_frag;
    wmma::fragment<wmma::matrix_b, 16,16,16, half, wmma::row_major> b_frag;
    wmma::fragment<wmma::accumulator,16,16,16, float> c_frag;

    wmma::fill_fragment(c_frag, 0.0f);

    // Load 16×16 tiles from global memory
    wmma::load_matrix_sync(a_frag, A, 16);
    wmma::load_matrix_sync(b_frag, B, 16);

    // Tensor Core fused-multiply-accumulate: C += A * B
    wmma::mma_sync(c_frag, a_frag, b_frag, c_frag);

    // Store result
    wmma::store_matrix_sync(C, c_frag, 16, wmma::mem_row_major);
}
```



# CUDA Streams

- ❑ A CUDA stream is a sequence of GPU operations executed in issue order.  
Operations in the same stream run *serially*; operations in different streams *may run concurrently*.
- ❑ Default stream (stream 0) synchronizes with all other work by default (“legacy” behavior).
- ❑ Streams enable concurrency:
  - Overlap kernel execution + memory copies
  - Run multiple independent kernels in parallel
  - Pipeline workloads for improved throughput
- ❑ Asynchronous operations
  - Most CUDA API calls have an async version (e.g., `cudaMemcpyAsync`) that takes a stream argument.

```
cudaStream_t s1, s2;
cudaStreamCreate(&s1);
cudaStreamCreate(&s2);

float *d_a, *d_b;
cudaMalloc(&d_a, N*sizeof(float));
cudaMalloc(&d_b, N*sizeof(float));

cudaMemcpyAsync(d_a, h_a, N*sizeof(float),
cudaMemcpyHostToDevice, s1);
cudaMemcpyAsync(d_b, h_b, N*sizeof(float),
cudaMemcpyHostToDevice, s2);

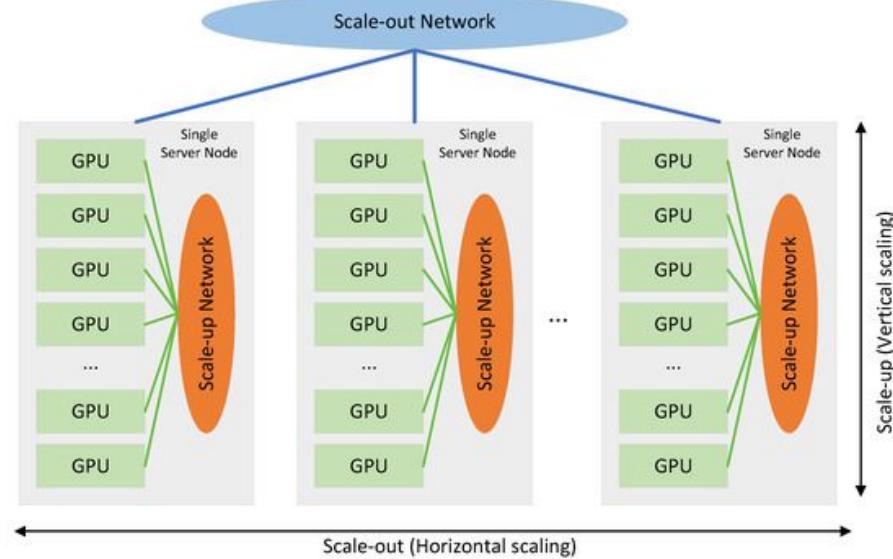
myKernel<<<grid, block, 0, s1>>>(d_a);
myKernel<<<grid, block, 0, s2>>>(d_b);

cudaStreamSynchronize(s1);
cudaStreamSynchronize(s2);
```



# Multiple GPUs

- ❑ Oftentimes, you will not have a single GPU, but many
- ❑ These gpus will be organized into a single node
  - o E.g. 4xA100 implies a node with 4 A100s
- ❑ A typical HPC setup will have many nodes connected to each other
- ❑ To use all GPUs in this configuration, we use primitives similar to MPI (NCCL)



# Multiple GPUs

- ❑ Oftentimes, you will not have a single GPU, but many
- ❑ These gpus will be organized into a single node
  - o E.g. 4xA100 implies a node with 4 A100s
- ❑ A typical HPC setup will have many nodes connected to each other
- ❑ To use all GPUs in this configuration, we use primitives similar to MPI (NCCL)

## Parallel Execution

```
size_t size = 1024 * sizeof(float);
cudaSetDevice(0);
float* p0;
cudaMalloc(&p0, size);
MyKernel<<<1000, 128>>>(p0);

cudaSetDevice(1);
float* p1;
cudaMalloc(&p1, size);
MyKernel<<<1000, 128>>>(p1);
```

# Multiple GPUs

- ❑ Oftentimes, you will not have a single GPU, but many
- ❑ These gpus will be organized into a single node
  - o E.g. 4xA100 implies a node with 4 A100s
- ❑ A typical HPC setup will have many nodes connected to each other
- ❑ To use all GPUs in this configuration, we use primitives similar to MPI (NCCL)

## Between GPU Communication

```
cudaSetDevice(0);
float* p0;
size_t size = 1024 * sizeof(float);
cudaMalloc(&p0, size);

cudaSetDevice(1);
float* p1;
cudaMalloc(&p1, size);

cudaSetDevice(0);
MyKernel<<<1000, 128>>>(p0);

cudaSetDevice(1);
cudaMemcpyPeer(p1, 1, p0, 0, size);
MyKernel<<<1000, 128>>>(p1);
```



# Dynamic Parallelism

- ❑ Dynamic Parallelism allows kernels to launch other kernels directly.
- ❑ Useful for irregular or adaptive workloads
- ❑ Reduces CPU–GPU synchronization by letting the GPU decide when to spawn work.
- ❑ Child kernels use the same `<<<grid, block>>>` syntax as host launches.
- ❑ Parent kernel must synchronize (or return) before child results are guaranteed.

```
__global__ void childKernel(int *out) {
    int i = threadIdx.x;
    out[i] = i * 2;
}

__global__ void parentKernel(int *out) {
    if (threadIdx.x == 0) {
        // Launch the child kernel from the device
        childKernel<<<1, 8>>>(out);
    }
}

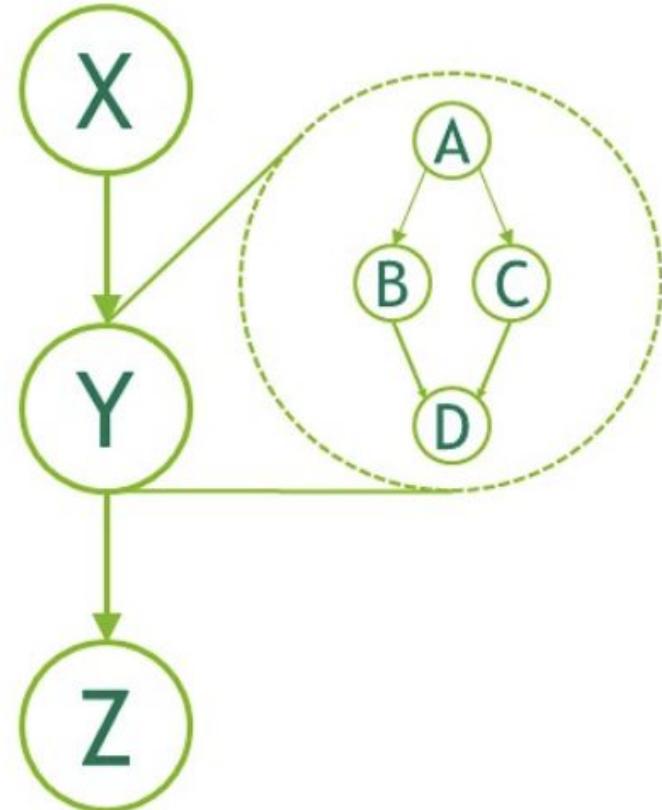
int main() {
    int *d_out;
    cudaMalloc(&d_out, 8 * sizeof(int));

    parentKernel<<<1, 1>>>(d_out);
    cudaDeviceSynchronize();
}
```



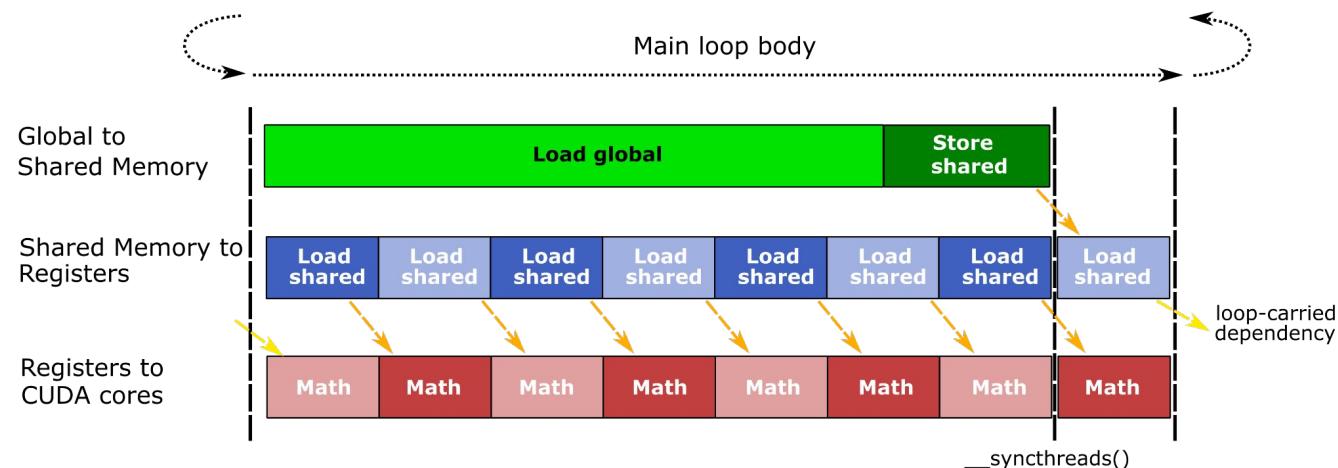
# CUDA Graphs

- ❑ Allows for defining structured computation which will likely be repeated
- ❑ Useful for long running applications, or where kernels are invoked many times, but do not individually take long
- ❑ In these cases, the cost of starting up the kernel can begin to accumulate
- ❑ CUDA graphs allow for amortization of this cost only once at program start



# Pipelining

- ❑ CUDA pipelining overlaps data transfer and kernel execution.
- ❑ Requires multiple non-default streams.
- ❑ Use cudaMemcpyAsync + kernel launches in different streams.
- ❑ Typical pattern: H2D → Compute → D2H done in staggered streams.
- ❑ Improves throughput when kernel time  $\approx$  memcpy time.



# Pipelining

Chunk 0: H2D → Kernel → D2H  
Chunk 1:      H2D → Kernel → D2H  
Chunk 2:      H2D → Kernel → D2H

```
for (int i = 0; i < N; i += CHUNK) {  
    int off = i;  
  
    cudaMemcpyAsync(d_in + off, h_in + off, CHUNK*sizeof(float),  
                   cudaMemcpyHostToDevice, s1);  
    kernel<<<grid, block, 0, s1>>>(d_in + off, d_out + off);  
    cudaMemcpyAsync(h_out + off, d_out + off, CHUNK*sizeof(float),  
                   cudaMemcpyDeviceToHost, s1);  
  
    cudaMemcpyAsync(d_in + off + CHUNK, h_in + off + CHUNK,  
                   CHUNK*sizeof(float), cudaMemcpyHostToDevice, s2);  
    kernel<<<grid, block, 0, s2>>>(d_in + off + CHUNK,  
                                         d_out + off + CHUNK);  
    cudaMemcpyAsync(h_out + off + CHUNK, d_out + off + CHUNK,  
                   CHUNK*sizeof(float), cudaMemcpyDeviceToHost, s2);  
}  
  
cudaDeviceSynchronize();
```

# PTX/SASS

- ❑ PTX = NVIDIA's virtual ISA (IR between CUDA and hardware).
- ❑ SASS = actual hardware instructions executed by the GPU.
- ❑ PTX is stable across architectures; SASS is architecture specific.
- ❑ Inline PTX allows precise control over instructions inside CUDA C++.
- ❑ Use `nvcc -ptx` to inspect compiler-generated PTX.
- ❑ Use `cuobjdump --dump-sass` (or `nvdism`) to inspect SASS.
- ❑ Performance debugging often requires looking at SASS, not PTX.
- ❑ Register pressure, memory ops, and instruction mix visible in SASS.

# PTX/SASS

## PTX

```
ld.global.f32 %f1, [x];
ld.global.f32 %f2, [y];
add.f32    %f3, %f1, %f2;
st.global.f32 [z], %f3;
```

## SASS

```
LDG.E.SYS R4, [R2];
LDG.E.SYS R5, [R3];
FADD R6, R4, R5;
STG.E.SYS [R7], R6;
```

## C

```
// compile with: nvcc -arch=sm_80 ptx_inline.cu -o ptx_inline
__global__ void add_ptx(float *x, float *y, float *z) {
    int i = threadIdx.x;
    float a = x[i], b = y[i], c;

    // inline PTX: c = a + b
    asm("add.f32 %0, %1, %2;" : "=f"(c) : "f"(a), "f"(b));

    z[i] = c;
}

int main() {
    add_ptx<<<1, 32>>>(x, y, z);
}
```



# Overview

- ❑ Advanced CUDA Constructs
- ❑ **Libraries**
- ❑ Examples of Parallelism in Practice
- ❑ Tooling
- ❑ Alternative Hardware Architectures



# cuBLAS

- ❑ cuBLAS is NVIDIA's optimized BLAS library for GPUs.
- ❑ General library for performing vector, matrix-vector, matrix-matrix operations
- ❑ cuBLAS achieves near-peak FLOP/s using tensor cores when allowed (TF32/FP16/etc.).
- ❑ `cublasCreate()` allocates internal context; reuse the handle for performance.
- ❑ High level of abstraction - if you just need primitives, this library is a good place to start



## cuBLAS

GPU-accelerated basic linear algebra  
(BLAS) library

# cuBLAS

```
// assume A,B,C allocated on device; A(m×k), B(k×n), C(m×n)
cublasHandle_t handle;
cublasCreate(&handle);

const float alpha = 1.0f, beta = 0.0f;

// C = alpha * A * B + beta * C  (column-major convention)
cublasSgemm(handle,
             CUBLAS_OP_N, CUBLAS_OP_N,
             n, m, k,           // note argument order: col-major
             &alpha,
             B, n,               // B is on the left in column-major
             A, k,
             &beta,
             C, n);

cublasDestroy(handle);
```



# cuDNN

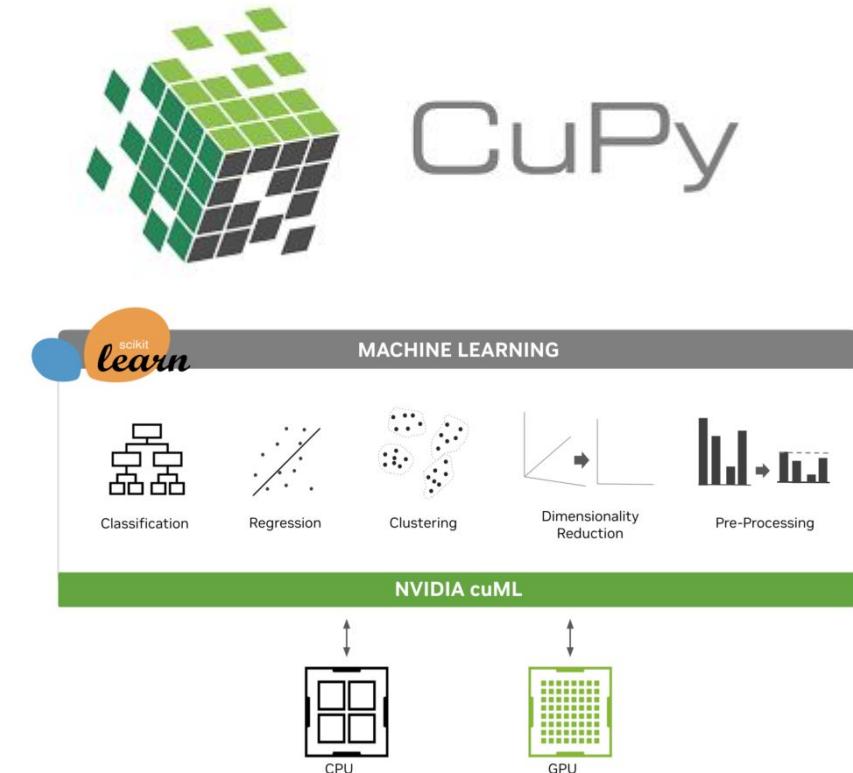
- ❑ cuDNN = NVIDIA's high-performance deep-learning primitives (conv, RNN, norm, activations).
- ❑ Provides highly-optimized GPU kernels (including Tensor Core usage) for DL frameworks.
- ❑ You configure operations using descriptors (tensor, filter, convolution, activation).
- ❑ Runtime selects fastest algorithm (FFT, Winograd, Implicit GEMM, Tensor Op).
- ❑ cuDNN performs: convolution, pooling, softmax, batchnorm, RNN/LSTM, attention (newer).
- ❑ Used in PyTorch, TensorFlow, JAX, MXNet, ONNX Runtime, Triton, TensorRT.
- ❑ Typical workflow: create handle → set descriptors → choose algorithm → allocate workspace → run op.



# Other useful, high level cuda libraries

Each of the following gives high level kernels that can be called “out-of-the-box”.  
No configuration is required outside of making the calls with the necessary inputs for your given problem

- cuFFT
- cuSparse
- cuSolver
- CuPy
- cuML



# CUTLASS C++

- ❑ CUTLASS is a CUDA C++ template library for high-performance kernels.
- ❑ Provides building blocks for tensor cores, tiling, and epilogue fusion.
- ❑ Backbone of PyTorch, TensorRT, cuBLAS, and many NVIDIA ML kernels.
- ❑ Supports FP16, BF16, TF32, FP32, INT8, and Hopper FP8 pipelines.
- ❑ Used to write custom fused kernels beyond what cuBLAS exposes.
- ❑ If you need lower level than cuBLAS, but want higher abstractions than kernels themselves, then use CUTLASS

```
int main() {  
    using Gemm = cutlass::gemm::device::Gemm<  
        float, cutlass::layout::RowMajor,  
        float, cutlass::layout::RowMajor,  
        float, cutlass::layout::RowMajor>;  
  
    Gemm gemm_op;  
  
    float *A, *B, *C; // assume allocated on device  
    int M = 1024, N = 1024, K = 1024;  
  
    Gemm::Arguments args({M, N, K}, {A, K}, {B, N},  
    {C, N}, {C, N});  
  
    gemm_op(args);  
}
```



# CuTe DSL

- ❑ Essentially the python version of CUTLASS C++
- ❑ Python API wraps tuned kernels without requiring C++ compilation.

```
import cutlass
import torch

A = torch.randn(1024, 1024, device="cuda", dtype=torch.float16)
B = torch.randn(1024, 1024, device="cuda", dtype=torch.float16)
C = torch.zeros(1024, 1024, device="cuda", dtype=torch.float16)

gemm = cutlass.op.gemm(
    A.shape[0], A.shape[1], B.shape[1],
    element=A.dtype, accumulate=torch.float32
)

gemm(A, B, C)
torch.cuda.synchronize()
print("Done")
```

# OpenAI Triton

- ❑ Triton enables writing custom GPU kernels in Python-like syntax.
- ❑ Kernels operate over program\_id blocks instead of CUDA thread IDs.
- ❑ Compiler performs automatic tiling, vectorization, and memory coalescing.
- ❑ Supports NVIDIA GPUs (Ampere, Hopper, Ada) with fast kernel generation

```
@triton.jit
def add_kernel(x_ptr, y_ptr, out_ptr, n, BLOCK: tl.constexpr):
    pid = tl.program_id(0)
    offs = pid * BLOCK + tl.arange(0, BLOCK)
    mask = offs < n
    x = tl.load(x_ptr + offs, mask=mask)
    y = tl.load(y_ptr + offs, mask=mask)
    tl.store(out_ptr + offs, x + y, mask=mask)

def add(x, y):
    n = x.numel()
    out = torch.empty_like(x)
    BLOCK = 128
    grid = lambda meta: (triton.cdiv(n, meta['BLOCK']),)
    add_kernel[grid](x, y, out, n, BLOCK=BLOCK)
    return out

x = torch.randn(1024, device='cuda')
y = torch.randn(1024, device='cuda')
z = add(x, y)
print(z[:5])
```



# Overview

- ❑ Advanced CUDA Constructs
- ❑ Libraries
- ❑ **Examples of Parallelism in Practice**
- ❑ Tooling
- ❑ Alternative Hardware Architectures



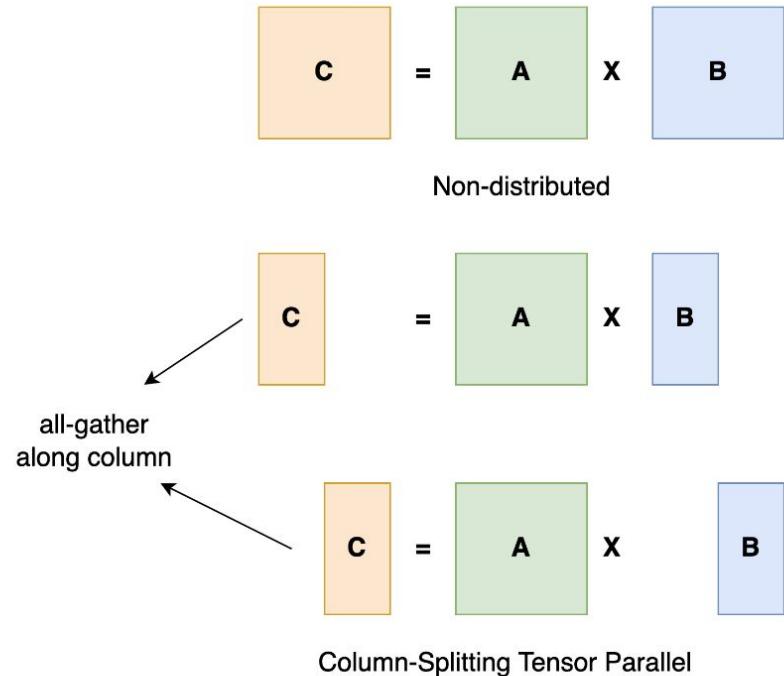
# LLM Parallelism

- ❑ Tensor Parallelism
- ❑ Data Parallelism
- ❑ Context Parallelism
- ❑ Pipeline Parallelism



# LLM Parallelism

- ❑ Tensor Parallelism
- ❑ Data Parallelism
- ❑ Context Parallelism
- ❑ Pipeline Parallelism

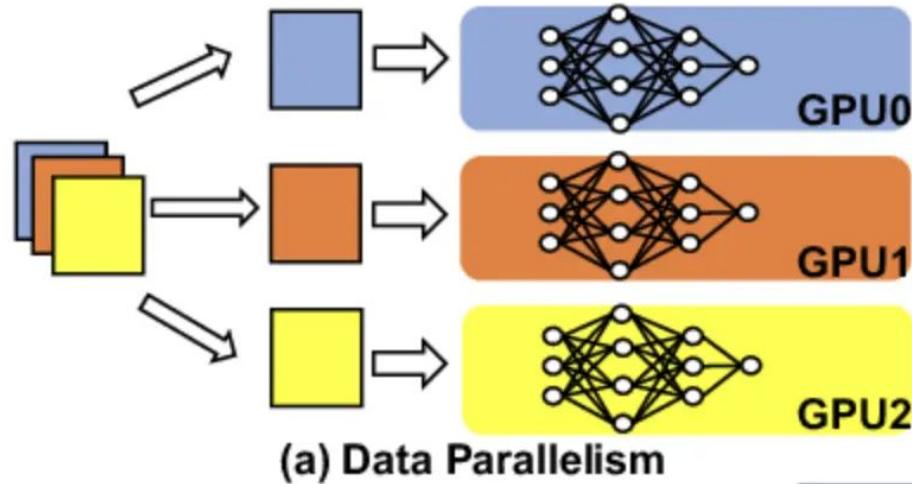


Tensor parallel illustration



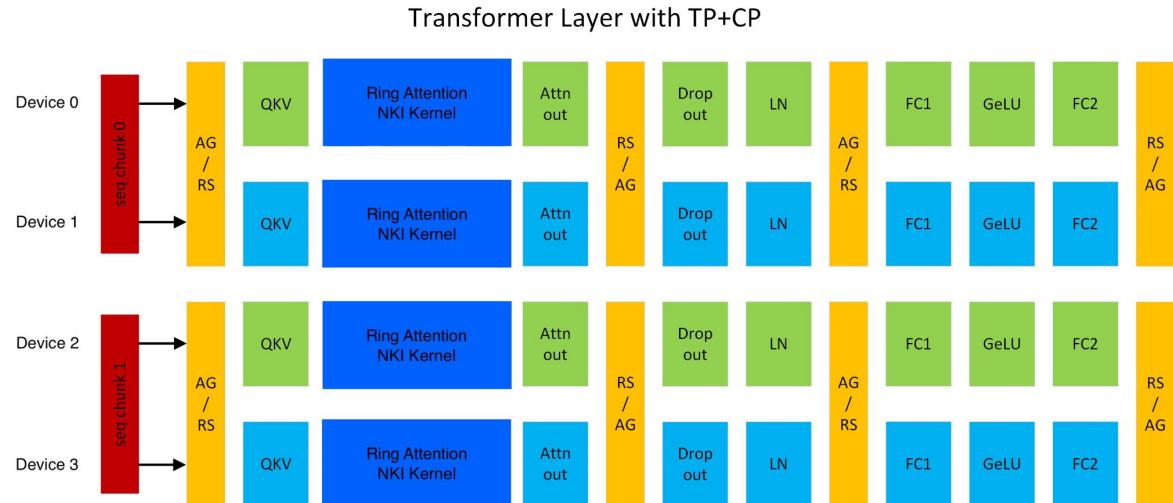
# LLM Parallelism

- ❑ Tensor Parallelism
- ❑ **Data Parallelism**
- ❑ Context Parallelism
- ❑ Pipeline Parallelism



# LLM Parallelism

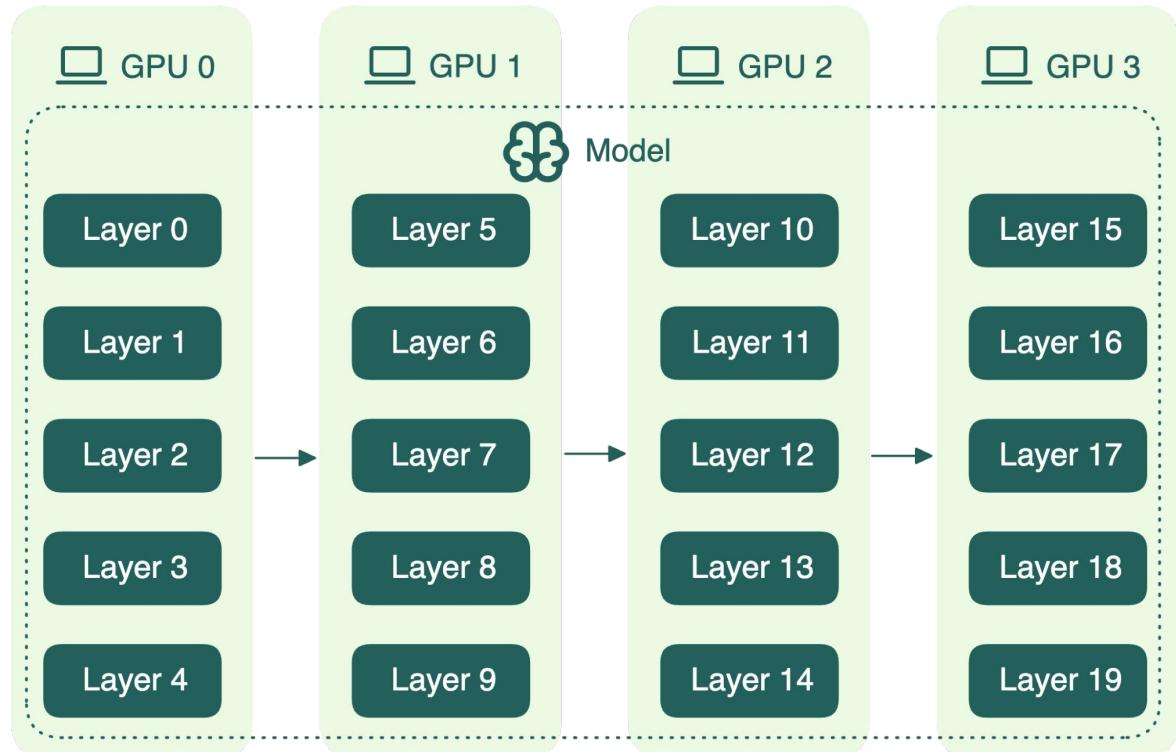
- ❑ Tensor Parallelism
- ❑ Data Parallelism
- ❑ **Context Parallelism**
- ❑ Pipeline Parallelism



# LLM Parallelism

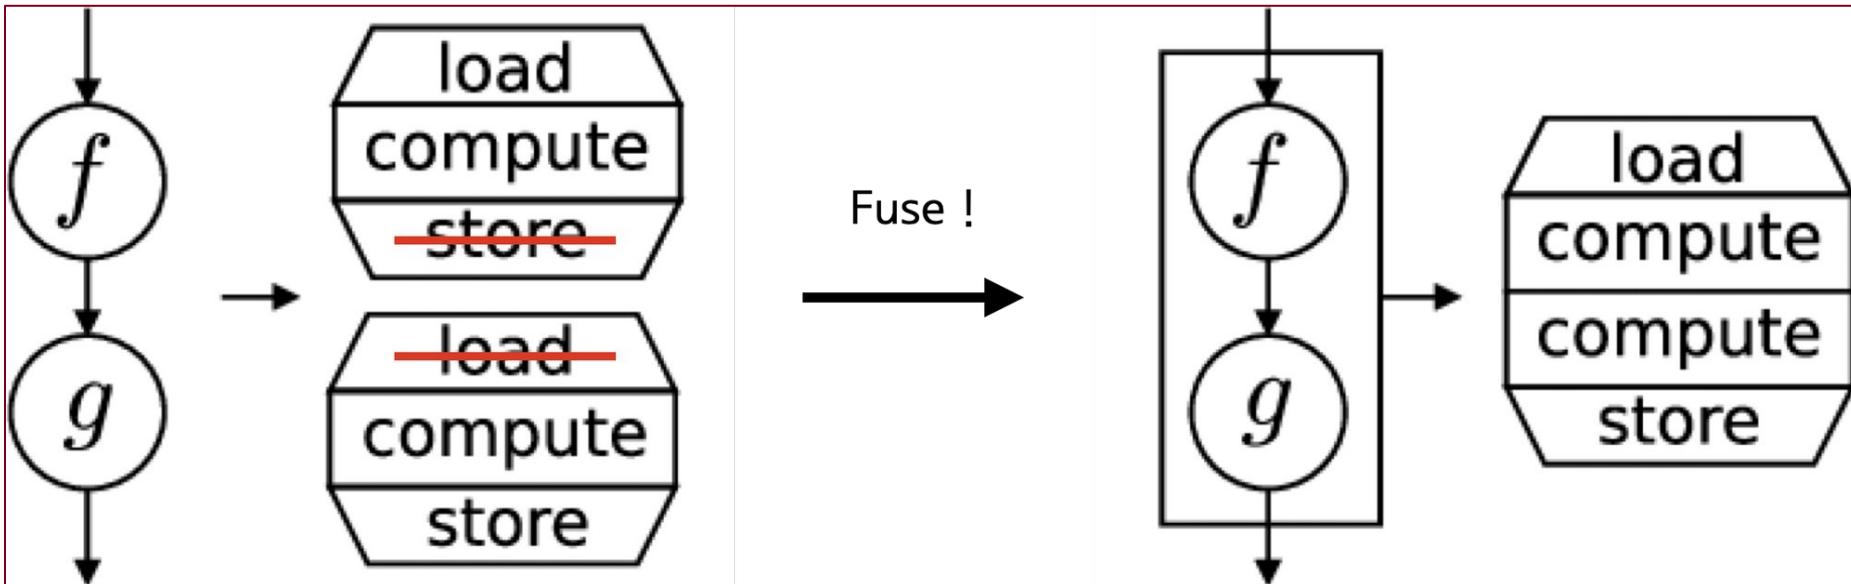
- ❑ Tensor Parallelism
- ❑ Data Parallelism
- ❑ Context Parallelism
- ❑ **Pipeline Parallelism**

Pipeline Parallelism



# Kernel Fusion

Combine multiple, separate kernels,  
into a single kernel

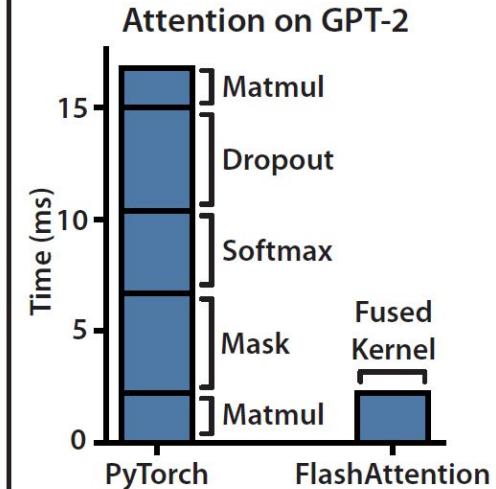
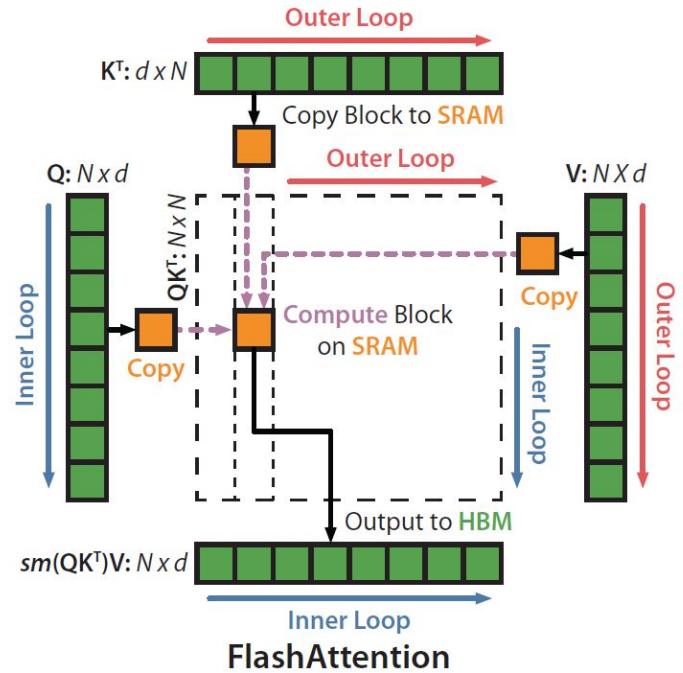
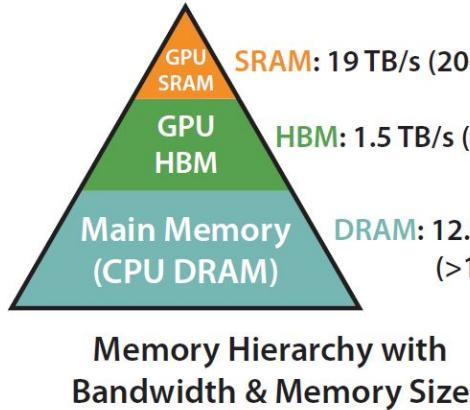


# Deep Learning Compilers

- ❑ Deep learning compilers optimize the entire computation graph, not just individual kernels.
- ❑ Kernel fusion reduces memory traffic (HBM bottleneck) and greatly improves performance.
- ❑ Compilers select and autotune the best kernels (CUTLASS, cuDNN, Triton, custom fused ops).
- ❑ Generate device-specific optimized code (PTX/SASS, tensor-core pipelines, tiling).
- ❑ Manage memory, scheduling, and parallelism across GPU(s) for training and inference.



# Attention



# Worklogs + Reference Lectures for Advanced CUDA

- ❑ [Inside NVIDIA GPUs: Anatomy of high performance matmul kernel](#) (Blog)
- ❑ [How to Optimize a CUDA Matmul Kernel for cuBLAS-like Performance: a Worklog](#) (Blog)
- ❑ [Learning CUDA by optimizing softmax: A worklog](#) (Blog)
- ❑ [Normalization layer optimization : Worklog](#) (Blog)
- ❑ GPU-MODE (lectures on more cuda internals)



# Overview

- ❑ Advanced CUDA Constructs
- ❑ Libraries
- ❑ Examples of Parallelism in Practice
- ❑ **Tooling**
- ❑ Alternative Hardware Architectures

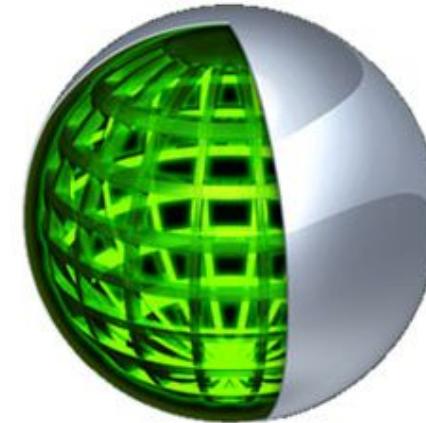


# Profiling (Nvidia Nsight Compute)

- ❑ Determines how the kernel itself performs
- ❑ Nsight Compute provides fine-grained GPU metrics such as cache hit rates, occupancy, warp stalls, and memory throughput.
- ❑ Kernel-level GPU profiler focused on per-kernel performance metrics.
- ❑ Reports warp occupancy, memory throughput, tensor core usage, and stall reasons.
- ❑ Allows section-based analysis (Memory, Scheduler, Roofline, Warp State).
- ❑ Supports detailed profiling on modern NVIDIA chips (Volta, Turing, Ampere, Hopper, Blackwell).

# Profiling (Nvidia Nsight Systems)

- ❑ Determines end-end performance of kernel
- ❑ System-wide profiler for CPU–GPU interactions and timeline analysis.
- ❑ Identifies kernel launch bottlenecks, synchronization stalls, and pipeline gaps.
- ❑ Provides multi-process and multi-node tracing for HPC workloads.
- ❑ Supports profiling across NVIDIA GPU generations (Kepler → Ada/Lovelace → Hopper/Blackwell).



NVIDIA®  
Nsight™

# Godbolt

- ❑ Useful site for examining PTX/SASS in real-time
- ❑ Interactive compiler to examine how your work compiles to PTX, then into SASS
- ❑ For very high performance kernels, examining this code is crucial



The image shows a screenshot of the Godbolt compiler interface. On the left, there is a code editor window displaying C++ code for a CUDA kernel. The code includes includes for `<algorithm>`, `<cassert>`, `<cstdio>`, `<cstdlib>`, `<cublas_v2.h>`, and `<cuda_runtime.h>`. It defines a macro `CEIL_DIV(M, N)` and contains a `__global__ void sgemmVectorize` function. The function takes parameters `int M, int N, int K, float alpha, float *B, float beta, float *C`. Inside the function, it declares constants `BM = 128`, `BN = 128`, `BK = 8`, `TM = 8`, and `TN = 8`. It also declares `cRow` and `cCol` as `const uint` variables based on `blockIdx.y` and `blockIdx.x` respectively. It calculates `totalResultsBlocktile` as `BM * BN` and `numThreadsBlocktile` as `totalResultsBlocktile / (TM * TN)`. A note in the code states: `// A thread is responsible for calculating TM*TN elements in the k`.

On the right, the assembly output is shown in three columns. The first column contains line numbers from 1 to 345. The second column contains the assembly instructions, which are color-coded by operation type: `add.s64` (green), `ld.global.v4.f32` (light green), `mul.f32` (purple), `mul.s32` (light purple), `cvt.u64.u32` (yellow), `add.s64` (green), `shl.b64` (light green), `add.s64` (green), `ld.global.v4.f32` (light green), `mul.f32` (purple), `mul.s32` (light purple), `mul.f32` (purple), `mul.s32` (light purple), `fma.rn.f32` (yellow), `fma.rn.f32` (purple), `fma.rn.f32` (purple), `fma.rn.f32` (purple), `st.global.v4.f32` (yellow), `add.s32` (light green), `cvt.u64.u32` (yellow), `add.s64` (green), `shl.b64` (light green), `add.s64` (green), `ld.global.v4.f32` (light green), `mul.f32` (purple), `mul.s32` (light purple), `mul.f32` (purple), `mul.s32` (light purple), `fma.rn.f32` (yellow), and `fma.rn.f32` (purple). The third column contains the memory addresses and register references for each instruction.



# Overview

- ❑ Advanced CUDA Constructs
- ❑ Libraries
- ❑ Examples of Parallelism in Practice
- ❑ Tooling
- ❑ **Alternative Hardware Architectures**

# MLX (Apple)

- ❑ MLX is Apple's lightweight array framework providing NumPy-like APIs with JIT compilation for CPU/GPU acceleration on Apple Silicon.
- ❑ Uses Metal Performance Shaders (MPS) and Metal compute kernels under the hood, mapping ops to tile-optimized GPU pipelines rather than CUDA-style SMs.
- ❑ Supports automatic differentiation, fused kernels, and graph-level optimizations to reduce memory traffic and launch overhead.



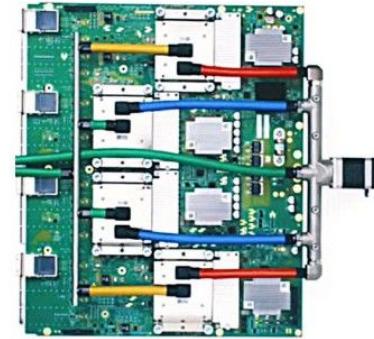
# MI-Series (AMD)

- ❑ AMD Instinct accelerators optimized for HPC and AI workloads.
- ❑ Provide high HBM bandwidth and large memory capacity for large-scale parallel computation.
- ❑ Support ROCm software stack with HIP for CUDA-portable GPU programming.
- ❑ Current flagship models include MI210, MI250X, and MI300 series accelerators used in large HPC systems.



# TPU (Google)

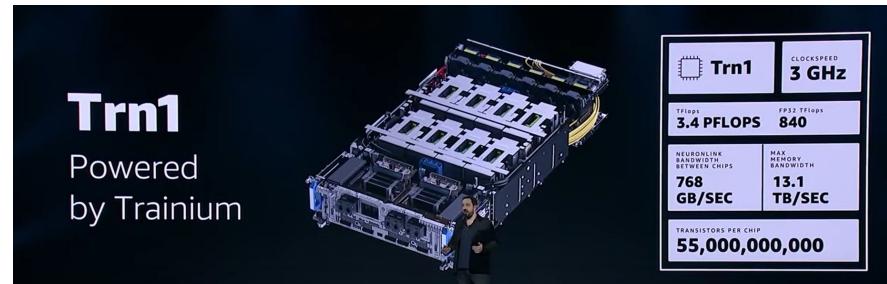
- TPUs are Google's domain-specific accelerators optimized for large-scale matrix multiplication.
- They use systolic arrays to maximize data reuse and throughput for ML workloads.
- TPU memory hierarchy emphasizes high-bandwidth on-chip buffer access over large off-chip DRAM.
- Production generations include TPU v1–v5p, with v5e/v5p powering modern large-model training clusters.



TPU<sup>V4</sup>

# Trainium (Amazon)

- ❑ AWS Trainium is a custom accelerator optimized for large-scale deep learning training.
- ❑ Uses NeuronCore v2 architecture with high memory bandwidth and dataflow-style execution.
- ❑ Integrates tightly with AWS EC2 Trn1/Trn2 instances for scalable distributed training.
- ❑ Trainium family includes Trainium (training) and Inferentia/Inferentia2 (inference) chips.



# Gaudi (Intel)

- ❑ Intel Gaudi is a purpose-built AI accelerator optimized for deep learning training.
- ❑ Architecture emphasizes high-bandwidth on-chip SRAM and integrated RoCE NICs for scale-out.
- ❑ Supports BF16, FP8, and mixed-precision training with high throughput.
- ❑ Product line includes Gaudi, Gaudi2, and Gaudi3 accelerators used in commercial AI clusters.



# Other Potential Hardware in the Future

- Qualcomm
- Broadcom/OpenAI
- Cerebras
- ...

Qualcomm

