

# CSCI 5451: Final Project

Fall 2025

This document describes the requirements, expectations, and grading scheme for the final project in CSCI 5451 (Parallel Algorithms). The project is worth a total of **24 points**, split across two major components:

- **Section 1: Project Planning Meeting** – 4 points
- **Section 2: Final Project & Report** – 20 points

## 1 Section 1: Project Planning Meeting (4 points)

### Scheduling and Deadline

Each group must schedule a **one-on-one final project planning meeting** with James Mooney using the Google Calendar booking link:

[Google Calendar Booking Link](#)

This meeting must take place **on or before Wednesday, November 26, 2025**. Scheduling and attending this meeting is worth **4 points**.

To receive credit:

- a) Use the booking link above to schedule a time.
- b) Ensure the meeting occurs by the deadline.
- c) The corresponding Canvas item for this requirement is available at:

[Canvas: Project Planning Meeting Assignment](#)

### Meeting Preparation

You **must** come to the meeting with concrete ideas for your final project. At a minimum, you should:

- Identify which of the project categories in Section 2.1 you are most interested in.
- Bring at least one specific problem, dataset, algorithm, or codebase you might work with.
- Be prepared to discuss:
  - Why the problem is interesting or challenging from a *parallel algorithms* perspective.

- What hardware (CPUs/GPUs, cluster nodes, etc.) and software stack (MPI, OpenMP, CUDA, etc.) you might use.
- Any initial ideas you already have for parallelization, optimization, or experimentation.

The goal of this meeting is to make sure your project is ambitious, well-scoped, and feasible given the course timeline and available resources. You will be referred back to Section 2.1 during the meeting as needed.

## 2 Section 2: Final Project and Report (20 points)

This section constitutes the main body of your final project. You will:

- Choose a project direction from the categories below.
- Implement and evaluate one or more parallel solutions.
- Design and run a thorough experimental evaluation.
- Submit:
  - (a) A written report.
  - (b) A GitHub repository with your code and tests.

The entirety of Section 2 is graded out of **20 points**, broken down in the grading rubric at the end of this section.

### 2.1 Project Categories

Your project can fall broadly into *one* (or a well-justified combination) of the following categories (this list is not exhaustive - if you have an alternative idea, it should be thoroughly developed).

#### 1) Parallelizing a difficult problem from another course.

- Choose a computationally intensive or conceptually challenging problem from a different course (e.g., machine learning, numerical methods, graphics, vision, data mining).
- Design and implement one or more parallel algorithms/implementations for this problem.

#### 2) Applying kernel fusion to a larger composition of kernels.

- Identify a workload composed of multiple kernels (e.g., a sequence of linear algebra operations, image-processing stages, or neural network components).
- Apply kernel fusion techniques to improve performance.
- For reference see:

## Kernel Fusion Article

### 3) Implementing an attention variant in CUDA.

- Implement a variant of attention (e.g., scaled dot-product attention, flash-style attention, sparse attention) in CUDA.
- Explore and compare different kernel designs, memory layouts, tiling strategies, etc.

### 4) In-depth performance analysis across hardware/software targets.

- Choose an algorithm or kernel (e.g., FFT, matrix multiplication, graph algorithm).
- Profile and compare performance across multiple hardware and/or software environments:
  - Different CPU architectures, GPU models, or node types.
  - Different compiler flags, libraries, or parallel programming models.
- Your emphasis is on deep analysis and explanation, not just raw benchmarks.

### 5) Worklog: iterative performance tuning on a target problem.

- Start from a baseline implementation of a nontrivial problem.
- Maintain a “worklog” documenting your optimization steps, design choices, and performance changes.
- For an example of the style of worklog (not to copy, but to emulate the spirit), see:

#### [CUDA MMM Worklog Example](#)

## 2.2 Testing Requirements

Regardless of which path you choose, your project must include:

### • At least 25 unit tests for correctness.

- These tests should validate that your implementation(s) produce correct or acceptable results across a variety of inputs.
- Where appropriate, tests should cover edge cases, varying problem sizes, and stress tests.
- Each test should have a reason for being chosen. Why is this helpful to test on?

### • At least 10 tests that demonstrate speedups.

- These tests should clearly compare performance against a baseline (e.g., sequential or naive parallel implementation).
- You should report timing, speedup, and (where applicable) efficiency metrics.

Your tests should be executable from the GitHub repository with clear instructions in the README. We should be able to reproduce at least a subset of your reported results using these tests.

## 2.3 Expected Deliverables

You must submit the following:

### 1. Written Report

A report that follows and addresses the grading scheme in Section 2.5. At a minimum, your report should include:

#### a) Introduction and Problem Description

- Describe the problem you chose in detail.
- Explain why it is interesting or challenging from a parallel perspective.

#### b) Methods / Parallelization Strategies

- Explain what strategies you used for parallelization and *why* you chose them.
- Discuss relevant theoretical considerations (e.g., work, communication, memory hierarchy).
- Describe different implementations you tried (at least 5 variants; see grading rubric).

#### c) Experimental Setup

- Describe the hardware (CPUs/GPUs, number of cores, memory, etc.) and software (compilers, libraries, frameworks) used.
- Explain how tests were run, input sizes, and any important configuration details.

#### d) Results and Analysis

- Present detailed results, including:
  - Speedup and efficiency metrics for at least one implementation across your test set.
  - Comparisons of at least 5 separate implementations and their speedups on some subset of tests.
- Use tables and/or plots to clearly present data.
- Discuss what worked, what did not, and how performance changed as you tuned the implementation.

#### e) Discussion and Conclusions

- Summarize the main takeaways.
- Reflect on what you would try next with more time (further optimizations, different architectures, larger problem sizes, etc.).

## 2. GitHub Repository

You must provide a link to a public or appropriately accessible **GitHub repository** containing:

- All source code for your implementations.
- Your test suite (a single test can be in both the categories below if you test a large problem both for correctness and speedups):
  - At least 25 correctness tests.
  - At least 10 performance/speedup tests.
- A clear README that includes:
  - How to build and run your code.
  - How to run correctness tests.
  - How to run performance tests and reproduce (at least some of) the reported results.
  - Any required dependencies or environment setup instructions.

We should be able to clone the repository, follow your instructions, and observe speedups similar to those reported in your paper for a subset of your experiments.

### 2.4 Submission and Deadline

All materials for Section 2 (report and GitHub link) must be submitted via the [Canvas Final Project link](#) by:

**December 18, 2025**

Please see the course Canvas site for the [submission link](#) and any additional logistics.

### 2.5 Grading Breakdown for Section 2 (20 points)

Section 2 is graded according to the following rubric:

#### (1) Reproducibility and Replication of Results (8 points)

- We should be able to clone your GitHub repository, follow your README, and run a subset of your tests.
- For that subset, we should see results that are consistent with those reported in your paper, including:
  - Correctness on the tested inputs.
  - Observable speedups relative to your baseline(s).

- Points are awarded based on:
  - Clarity and completeness of instructions.
  - Ease of running tests and scripts.
  - Consistency between reported and reproduced results.

## (2) Problem Description and Parallelization Strategy (4 points)

- Your report should include a clear and detailed description of:
  - The problem you chose (including necessary background).
  - Why it is meaningful in the context of parallel algorithms.
- You should thoroughly describe what strategies you chose for parallelization and *why*:
  - Parallel decomposition, communication patterns, synchronization, etc.
  - Any relevant trade-offs you considered.

## (3) Results, Implementations, and Performance Analysis (8 points)

- Your results section should be **detailed and exhaustive** within the scope of your project.
- At a minimum, it must include:
  - Speedup and efficiency metrics from the test suite for **at least one** implementation.
  - A comparison of **at least 5 separate implementations** and their speedups on some set of tests.
- You should carefully explain:
  - What exactly was done in each implementation (e.g., different tilings, memory layouts, synchronization strategies, kernel fusion variants, etc.).
  - What other things you tried (even if they did not work as well) and what results they produced.
  - How you tuned your implementation(s) to obtain the best possible performance.
- Strong projects will:
  - Provide insightful analysis rather than just tables of numbers.
  - Connect empirical results back to theoretical expectations.
  - Clearly identify bottlenecks and explain improvements across implementations.

**Total:** 4 points (Section 1) + 20 points (Section 2) = **24 points**.