


# CSCI 5541: Natural Language Processing

## Lecture 7: Language Models: RNN, LSTM, and Seq2Seq



UNIVERSITY OF MINNESOTA  
**Driven to Discover®**

# Announcement (0213)

- ❑ Minor HW2 Revisions --> See slack announcement
- ❑ HW3 is released. The due date is due Tue, Feb 25.
- ❑ Project
  - Brainstorming is due next Tuesday, Feb 18
  - Groups have been assigned in slack
  - There are a couple of students not yet in groups. If you have a fully formed group and are willing to take on someone else, let me know.



# Ngram LM



Uni-gram

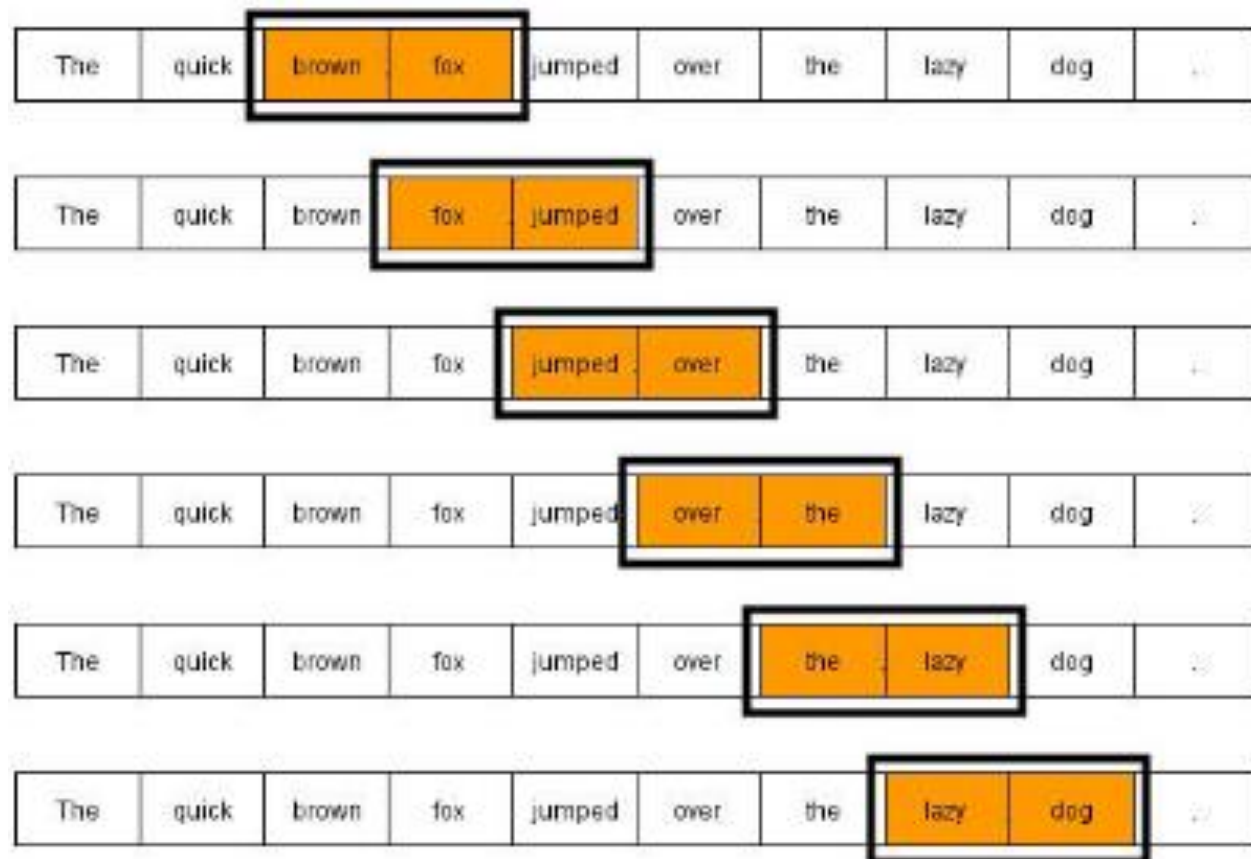
$$\prod_{i=1}^n P(w_i) \\ \times P(STOP)$$

$$\frac{c(w_i)}{N}$$

Bi-gram

$$\prod_{i=1}^n P(w_i | w_{i-1}) \\ \times P(STOP | w_n)$$


$$\frac{c(w_{i-1}, w_i)}{c(w_{i-1})}$$



# Sparsity in Ngram LM

	i	want	to	eat	chinese	food	lunch	spend
i	5	827	0	9	0	0	0	2
want	2	0	608	1	6	6	5	1
to	2	0	4	686	2	0	6	211
eat	0	0	2	0	16	2	42	0
chinese	1	0	0	0	0	82	1	0
food	15	0	15	0	1	4	0	0
lunch	2	0	0	0	0	1	0	0
spend	1	0	1	0	0	0	0	0

**Figure 4.1** Bigram counts for eight of the words (out of  $V = 1446$ ) in the Berkeley Restaurant Project corpus of 9332 sentences. Zero counts are in gray.



$$\frac{c(w_{i-1}, w_i)}{c(w_{i-1})} \rightarrow \frac{c(w_{i-1}, w_i) + \alpha}{c(w_{i-1}) + V\alpha}$$

$$P(w_i | w_{i-2}, w_{i-1}) = \lambda_1 P(w_i | w_{i-2}, w_{i-1}) + \lambda_2 P(w_i | w_{i-1}) + \lambda_3 P(w_i)$$

$$\lambda p + (1 - \lambda)q$$



# Ngram LM vs Neural LM

To avoid the data sparsity  
problem from the ngram LM



# Neural LM



$$x = [v(w_1); \dots v(w_k)]$$


Concatenation ( $k \times V$ )

$w_1 = \text{tried}$


$w_2 = \text{to}$

$w_3 = \text{prepare}$

$w_4 = \text{midterms}$



One-hot encoding



Distributed representation



Multi-class (Vocab)  
classification

Bengio et al. 2003, A Neural Probabilistic Language Model



# Neural LM




$$P(w) = P(w_i | w_{i-k} \dots w_{i-1}) = \text{softmax}(W \cdot h)$$

One-hot encoding  
( $|x| = V$ )

$$W_1 \in \mathbb{R}^{kV \times H} \quad W_2 \in \mathbb{R}^{H \times V}$$
$$b_1 \in \mathbb{R}^H \quad b_2 \in \mathbb{R}^V$$

Output space:  $|y| = V$



$$x = [v(w_1); \dots; v(w_k)]$$

$$\hat{y} = \text{softmax}(hW_2 + b_2)$$



# Neural LM



Represent high-dimensional words (and contexts) as low-dimensional vectors

One-hot encoding  
( $|x| = V$ )

Distributed representation  
( $|y| = H$ )

A large dark red trapezoid, wider on the left and narrower on the right, representing a dimensionality reduction. The text  $V \gg H$  is centered inside it in white.

$V \gg H$







Conditioning context ( $X [k \times V]$ )

tried to prepare midterm but I was too tired of...

Next word to predict ( $Y$ )

Context window size:  $k=4$





Conditioning context ( $X [k \times V]$ )

tried to prepare midterm but I was too tired of...

Next word to predict ( $Y$ )

Context window size:  $k=4$





Conditioning context ( $X [k \times V]$ )

tried to prepare midterm but I was too tired of...

Next word to predict ( $Y$ )

Context window size:  $k=4$



# Neural LM against Ngram LM



## Pros

- ❑ No sparsity problem
- ❑ Don't need to store all observed n-gram counts

## Cons

- ❑ Fixed context window is too small (larger window, larger  $W$ )
  - Windows can never be large enough
- ❑ Different words are multiplied by completely different weights ( $W$ ); no **symmetry** in how the inputs are processed.



# Outline

- ❑ Linearization: A general heuristic for model improvement
- ❑ Recurrent Neural Network (RNN)
- ❑ Long Short-term Memory (LSTM)
- ❑ Implementation of RNN and LSTM using PyTorch
- ❑ Sequence-to-Sequence modeling
- ❑ Teaser: Transformer-based LMs
- ❑ Why language models are useful?



# Outline


- ❑ **Linearization: A general heuristic for model improvement**
- ❑ Recurrent Neural Network (RNN)
- ❑ Long Short-term Memory (LSTM)
- ❑ Implementation of RNN and LSTM using PyTorch
- ❑ Sequence-to-Sequence modeling
- ❑ Teaser: Transformer-based LMs
- ❑ Why language models are useful?



# How do we make a better model?




# More Params are Better






# Increasing depth is more efficient than width

Model 1: ❌



Model 2: ✅



...but very deep models are harder to train




Figure 12.2 Deep network architecture with multiple layers.

# Why is this so challenging?





Figure 12.2 Deep network architecture with multiple layers.




# Backprop Revisited



# Analogy #1: A Game of Telephone




# Analogy #2: A funnel of information



# Linearization Solves This

- ❑ We need a better way to reduce the number of operations performed between our weights and our loss function (Residual connections)
- ❑ We need a better way to ensure we are not bottlenecking any representations into some channel which is too small to contain all the information we need (Attention mechanism → later)



# Outline


- ❑ Linearization: A general heuristic for model improvement
- ❑ **Recurrent Neural Network (RNN)**
- ❑ Long Short-term Memory (LSTM)
- ❑ Implementation of RNN and LSTM using PyTorch
- ❑ Sequence-to-Sequence modeling
- ❑ Teaser: Transformer-based LMs
- ❑ Why language models are useful?






# Recurrent Neural Network (RNN)

RNN allow arbitrarily-sized conditioning contexts;  
condition on the **entire sequence history**.



# Recurrent Neural Network



Neural-LM:

$$P(w) = P(w_i | w_{i-k} \dots w_{i-1}) = \text{softmax}(W \cdot \mathbf{h})$$


RNN:

$$P(w) = P(w_i | \text{context}) \\ = \text{softmax}(W \cdot \mathbf{h}_i)$$




# Recurrent Neural Network

- Each time set has two inputs:
- $X_i$  (the observation at time step  $i$ ):
  - One-hot vector, feature vector, or distributed representation of input token at  $i$  step




# Recurrent Neural Network

- ❑ Each time set has two inputs:
- ❑  $X_i$  (the observation at time step  $i$ ):
  - One-hot vector, feature vector, or distributed representation of input token at  $i$  step
- ❑  $S_{i-1}$  (the output of the previous state):
  - Base case:  $S_0 = 0$  vector



# Recurrent Neural Network


- Each time set has two outputs:
- $S_i = R(X_i, S_{i-1})$ 
  - $R$  computes the **output state** as a function of the *current input* and *previous state*
- $y_i = O(S_i)$ 
  - $O$  computes the **output** as a function of the *current output state*



# RNN Training

output as  
shifted by one


sequence of  
words



# RNN Training

output as  
shifted by one

sequence of  
words

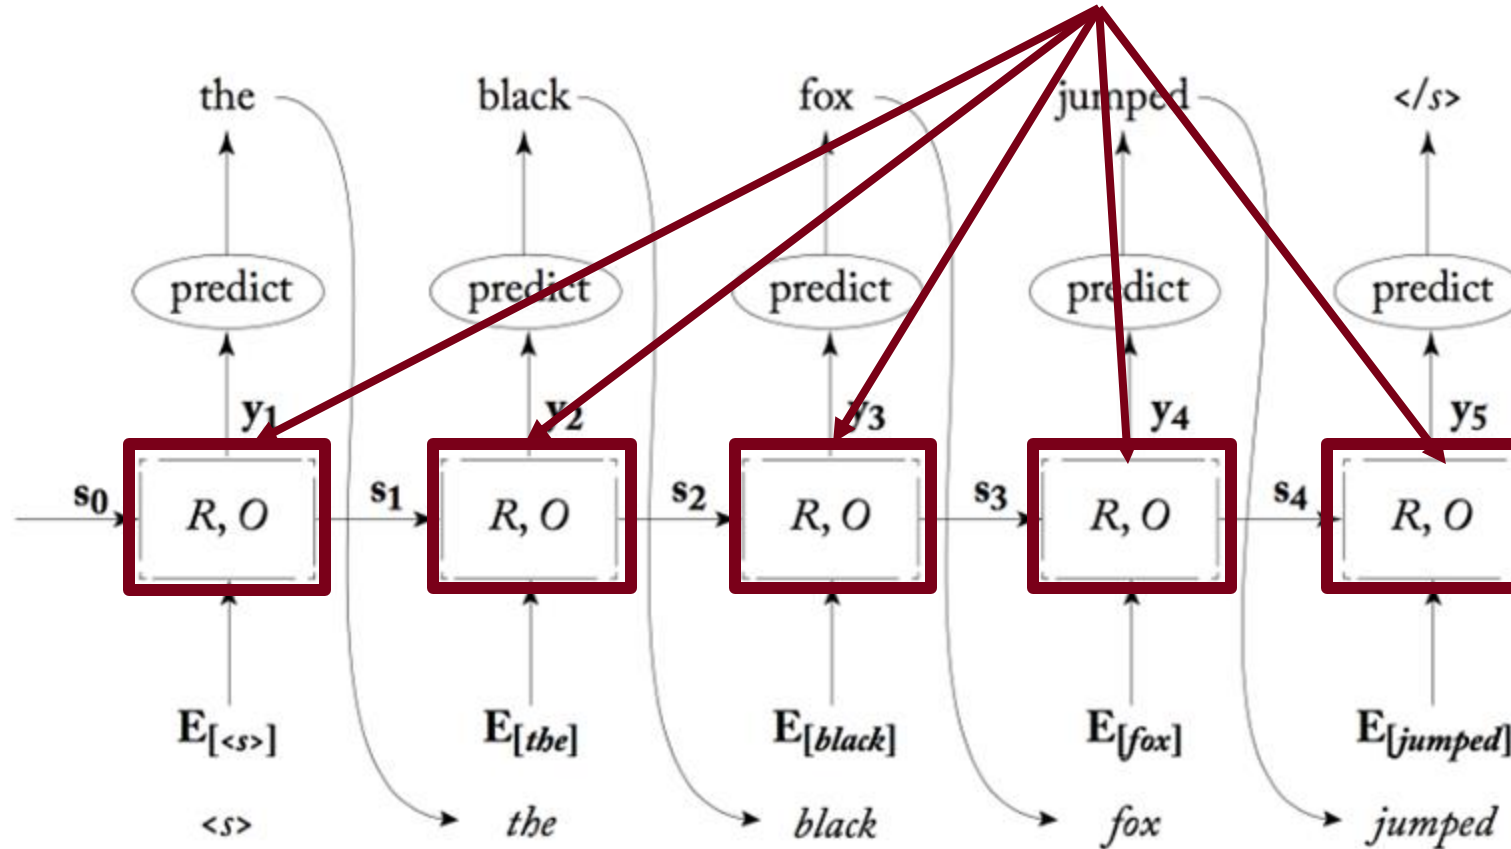


# RNN Training

Parameters are shared!  
Derivatives are accumulated.

output as  
shifted by one

sequence of  
words






# What can RNNs do?

- ❑ Represent a sentence
  - Read whole sentence, make a prediction
- ❑ Represent a context within a sentence
  - Read context up until that point


# Representing Sentences

- Sentence classification
- Conditioned generation




# Representing Context within Sentence

- Tagging
- Language modeling




# e.g., Language Modeling

- Language modeling is like a tagging task, where each tag is the next word!



# e.g., POS Tagging with Bi-RNNs



# Outline

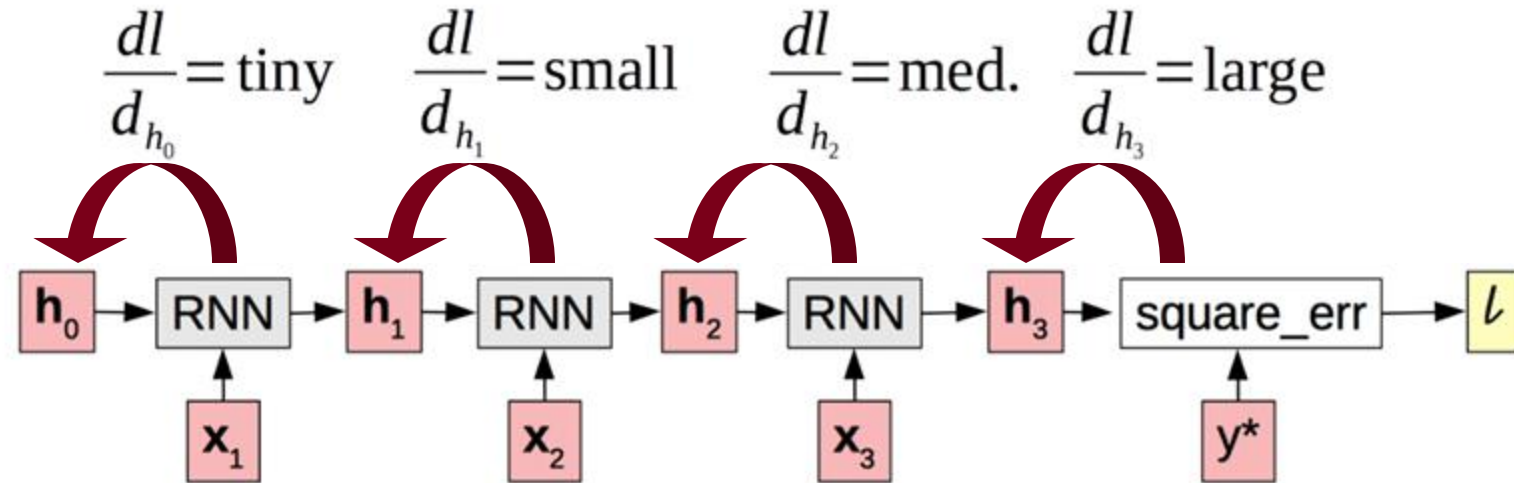
- ❑ Linearization: A general heuristic for model improvement
- ❑ Recurrent Neural Network (RNN)
- ❑ **Long Short-term Memory (LSTM)**
- ❑ Implementation of RNN and LSTM using PyTorch
- ❑ Sequence-to-Sequence modeling
- ❑ Teaser: Transformer-based LMs
- ❑ Why language models are useful?



# Vanishing Gradient



- ❑ Gradients decrease as they get pushed back



- ❑ Why? “Squashed” by non-linearities or small weights in matrices

# A Solution: Long Short-term Memory (LSTM)


(Hochreiter and Schmidhuber 1997)

- ❑ Make **additive connections** between time steps
- ❑ Addition does not modify the gradient, no vanishing
- ❑ **Gates** to control the information flow






# RNN Structure





$$s_i = R(x_i, s_{i-1})$$
$$y_i = O(s_i)$$





<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

# RNN vs LSTM Structure



<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

# RNN vs LSTM Structure



<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

# LSTM Structure

- ❑ **Forget gate:** what value do we try to add/forget to the memory cell?
- ❑ **Input gate:** how much of the update do we allow to go through?
- ❑ **Output gate:** how much of the cell do we reflect in the next state?

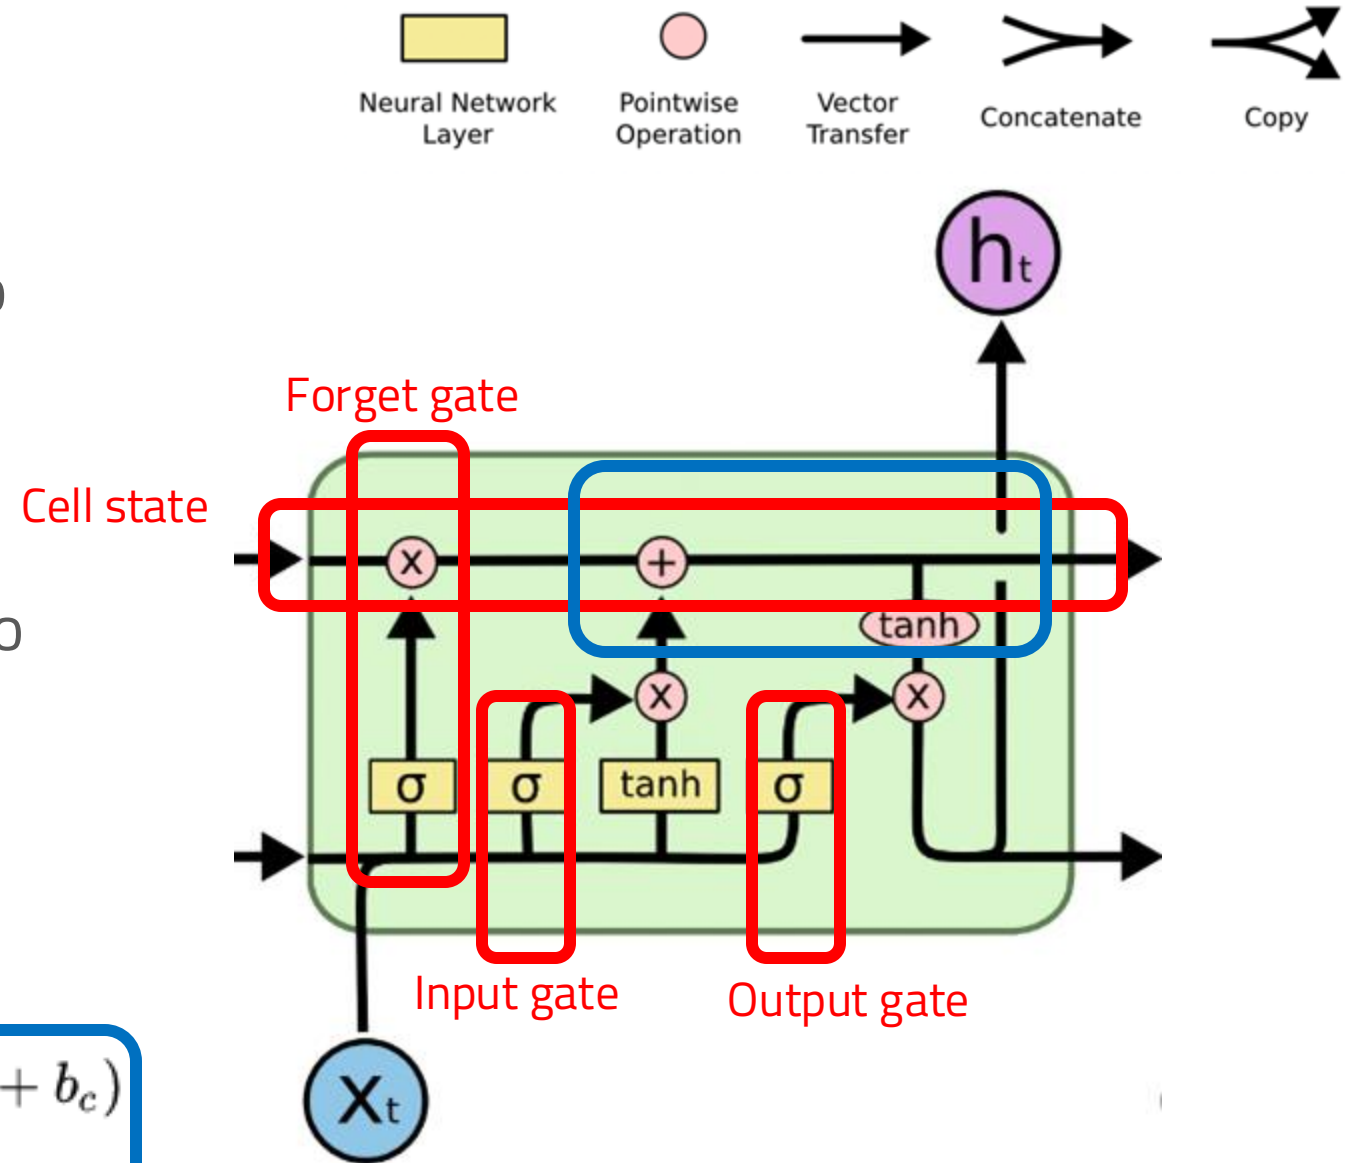
$$f_t = \sigma_g(W_f x_t + U_f h_{t-1} + b_f)$$

$$i_t = \sigma_g(W_i x_t + U_i h_{t-1} + b_i)$$

$$o_t = \sigma_g(W_o x_t + U_o h_{t-1} + b_o)$$

$$c_t = f_t \circ c_{t-1} + i_t \circ \sigma_c(W_c x_t + U_c h_{t-1} + b_c)$$

$$h_t = o_t \circ \sigma_h(c_t)$$



<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

# LSTM variant: Gated Recurrent Unit (GRU)

(Cho et al., 2014)


- Combines the forget and input gates into a single "update gate."
- Merges the cell state and hidden state
- And, other small changes

$$z_t = \sigma_g(W_z x_t + U_z h_{t-1} + b_z)$$

$$r_t = \sigma_g(W_r x_t + U_r h_{t-1} + b_r)$$

$$h_t = (1 - z_t) \circ h_{t-1} + z_t \circ \sigma_h(W_h x_t + U_h(r_t \circ h_{t-1}) + b_h)$$


Additive or Non-linear



<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

# Most Important Takeaway

- ❑ The Cell State is an information highway
- ❑ Gradient can flow over this without nearly as many issues of vanishing/exploding gradients that we saw in RNNs
- ❑ We are doing a better job at reducing the 'distance' between our loss function and each individual parameter



# A Solution: Long Short-term Memory (LSTM)

(Hochreiter and Schmidhuber 1997)

- ❑ Make **additive connections** between time steps
- ❑ Addition does not modify the gradient, no vanishing
- ❑ **Gates** to control the information flow



# Outline

- ❑ Linearization: A general heuristic for model improvement
- ❑ Recurrent Neural Network (RNN)
- ❑ Long Short-term Memory (LSTM)
- ❑ **Implementation of RNN and LSTM using PyTorch**
- ❑ Sequence-to-Sequence modeling
- ❑ Teaser: Transformer-based LMs
- ❑ Why language models are useful?





**class RNN(nn.Module):**

```
def __init__(self, input_size: int, hidden_size: int, output_size: int) -> None:
    super().__init__()
```

```
...
```

```
self.i2h = nn.Linear(input_size, hidden_size, bias=False)
```

```
self.h2h = nn.Linear(hidden_size, hidden_size)
```

```
self.h2o = nn.Linear(hidden_size, output_size)
```

```
def forward(self, x, hidden_state):
```

```
    x = self.i2h(x)
```


```
    hidden_state = self.h2h(hidden_state)
```

```
    hidden_state = torch.tanh(x + hidden_state)
```

```
    out = self.h2o(hidden_state)
```

```
    return out, hidden_state
```

$$\left. \begin{array}{l} s_i = R(x_i, s_{i-1}) \\ y_i = O(s_i) \end{array} \right\}$$



```
class RNN(nn.Module):
```

```
    def __init__(self, input_size, output_size, hidden_dim, n_layers):
        super(RNN, self).__init__()
```

```
    ...
```

```
    self.rnn = nn.RNN(input_size, hidden_dim, n_layers, batch_first=True)
    self.fc = nn.Linear(hidden_dim, output_size)
```

```
    def forward(self, x, hidden):
```


```
        r_out, hidden = self.rnn(x, hidden)
```

```
        r_out = r_out.view(-1, self.hidden_dim)
```

```
        return self.fc(r_out) , hidden
```

$$s_i = R(x_i, s_{i-1})$$

$$y_i = O(s_i)$$



# x (batch\_size, seq\_length, input\_size)  
 # hidden (n\_layers, batch\_size, hidden\_dim)  
 # r\_out (batch\_size, time\_step, hidden\_size)

**class LSTM (nn.Module):**

**def** \_\_init\_\_(self, num\_classes, input\_size, hidden\_size, num\_layers, seq\_length):

    super(LSTM1, self).\_\_init\_\_()

    ...

    self.lstm = **nn.LSTM**(input\_size=input\_size, hidden\_size=hidden\_size, num\_layers=num\_layers, batch\_first=True)

    self.fc = **nn.Linear**(hidden\_size, num\_classes)

    self.relu = **nn.ReLU**()

**def** forward(self,x):

    h\_0 = Variable(torch.zeros(self.num\_layers, x.size(0), self.hidden\_size))


    c\_0 = Variable(torch.zeros(self.num\_layers, x.size(0), self.hidden\_size))

    output, (hn, cn) = **self.lstm**(x, (h\_0, c\_0))

    hn = hn.view(-1, self.hidden\_size)

    return **self.fc** (**self.relu**(hn))

 PyTorch



$$f_t = \sigma_g(W_f x_t + U_f h_{t-1} + b_f)$$

$$i_t = \sigma_g(W_i x_t + U_i h_{t-1} + b_i)$$

$$o_t = \sigma_g(W_o x_t + U_o h_{t-1} + b_o)$$

$$c_t = f_t \circ c_{t-1} + i_t \circ \sigma_c(W_c x_t + U_c h_{t-1} + b_c)$$

$$h_t = o_t \circ \sigma_h(c_t)$$

# Outline


- ❑ Linearization: A general heuristic for model improvement
- ❑ Recurrent Neural Network (RNN)
- ❑ Long Short-term Memory (LSTM)
- ❑ Implementation of RNN and LSTM using PyTorch
- ❑ **Sequence-to-Sequence modeling**
- ❑ Teaser: Transformer-based LMs
- ❑ Why language models are useful?



# Connecting RNN to RNN for sequence-to-sequence (seq2seq) modeling




# RNN (decoder) for language modeling




# RNN (decoder) for language modeling

What if we encode some specific context, instead of random state?



# RNN (encoder) - RNN (decoder) for machine translation

“나는 이 영화가 싫어요”  
“Odio esta película”






# RNN (encoder) - RNN (decoder) for dialogue generation

“나는 이 영화가 싫어요”  
“Odio esta película”

“what do you think about  
*Avengers: Endgame*?”

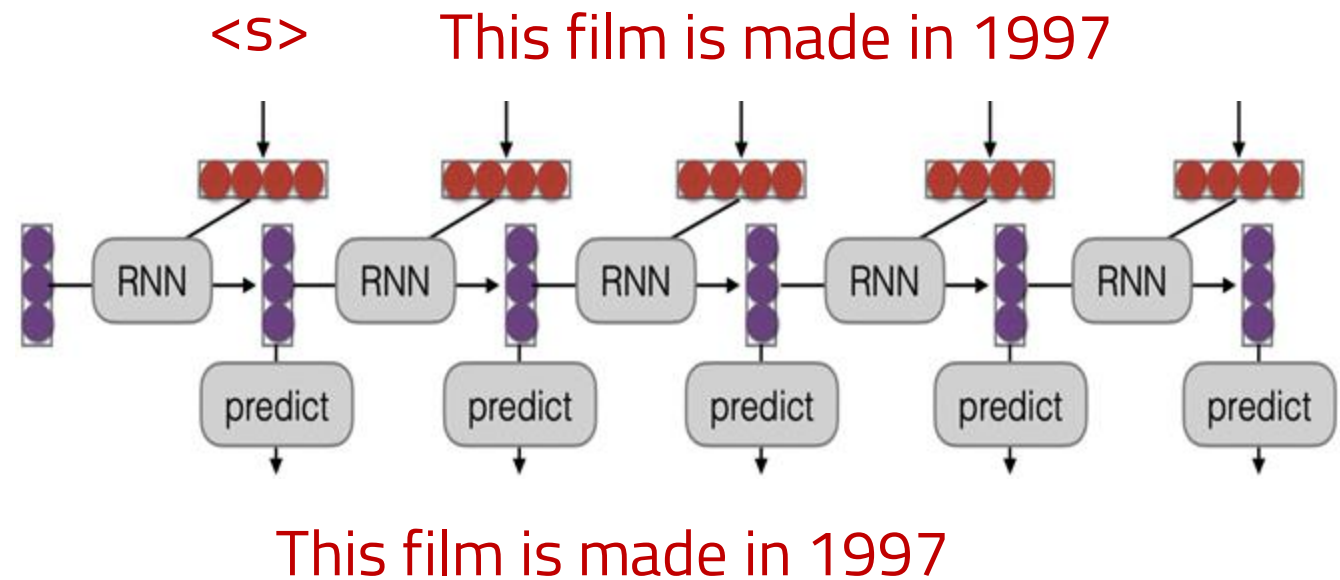


# RNN (encoder) - RNN (decoder) for question answering

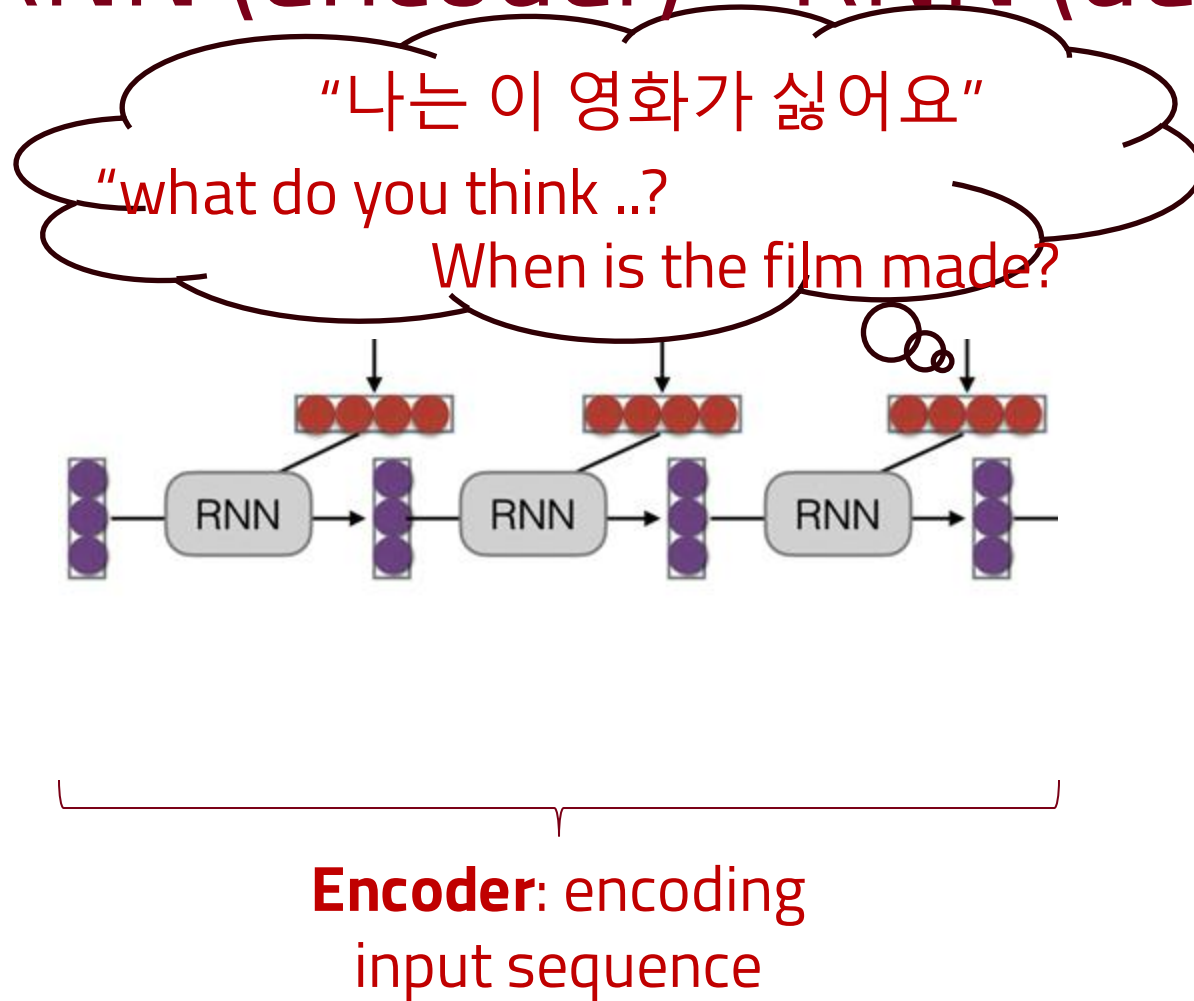
“나는 이 영화가 싫어요”  
“Odio esta película”

“what do you think about  
*Avengers: Endgame*?”


When is the film made?



# Sequence-to-sequence modeling using RNN (encoder) - RNN (decoder)



# Sequence-to-sequence modeling using RNN (encoder) - RNN (decoder)




# Problem: forgetting input context as input gets longer



“나는 이 영화가 싫어요”

“what do you think ..? When is the film made?”





# Solution (teaser): Seq2seq with attention



“나는 이 영화가 싫어요”

“what do you think ..? When is the film made?”



Attention layer = Input context  
attended on all previous context  
(will be covered more in Transformer)



# State-of-the-art Language Models



# Teaser: Transformer-based LMs

- ❑ SOTA LMs: **GPT-2**, Radford et al. 2018; **GPT-3**, Brown et al. 2020

Trigram	LSTM	<b>GPT-2</b>	<b>GPT-3</b>
109	58.3	35.8	20.5

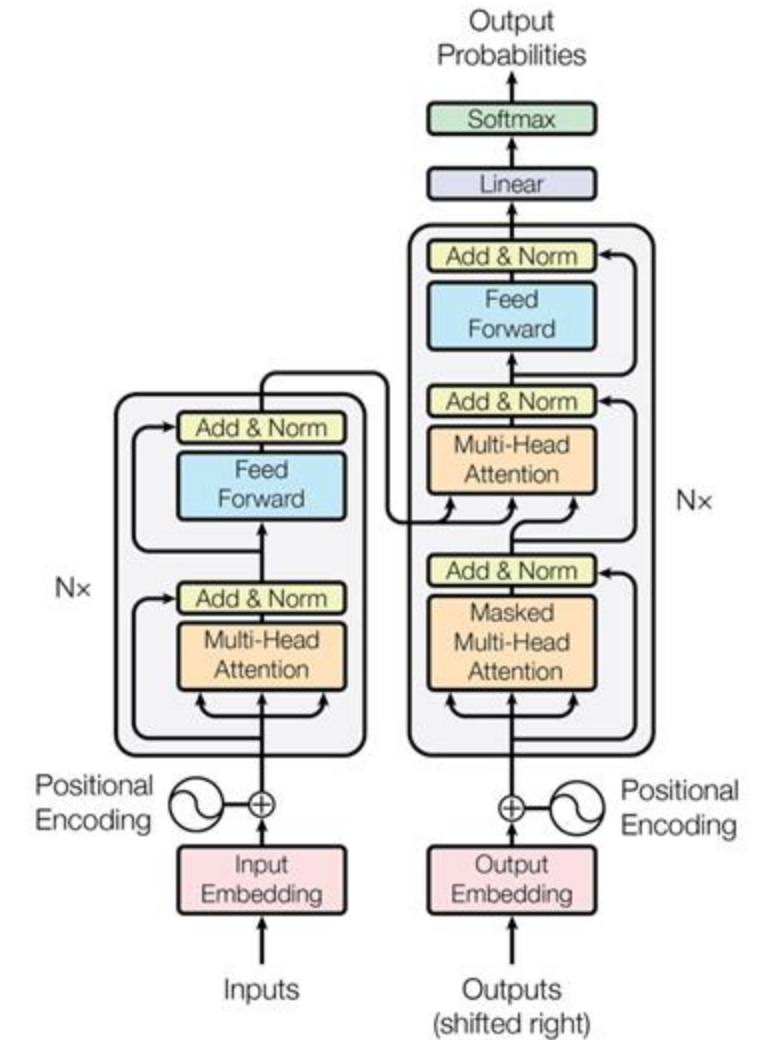




Figure 1: The Transformer - model architecture.







Jürgen Schmidhuber

Pronounce: You again Shmidhoobuh

Technical Report IDSIA-23-23, IDSIA

## AI Blog

Twitter: @SchmidhuberAI

14 December 2023

## How 3 Turing Awardees Republished Key Methods and Ideas Whose Creators They Failed to Credit

This write-up is meant to correct an inaccurate history of Artificial Intelligence (AI) propagated by recent uninformed news articles, posts in social media, and a [large language model](#). Most of its statements are taken from a less streamlined [report](#)<sup>[T722]</sup> that has been reviewed on relevant AI mailing lists, profiting from feedback by many experts and well-known AI pioneers. **The piece is aimed at people who are not aware of the numerous AI priority disputes, but are willing to check the facts.**



Perplexity

Ngram

LSTM

RNN

GRU

ELMo

GPT

BERT

GPT2

GPT3

1990s

1997

2003


2014

2018


2019

2020

2021




# Teaser: Two Objectives for Language Model Pretraining



$$\log p(\mathbf{x}) = \sum_{t=1}^T \log p(x_t | \mathbf{x}_{<t})$$

Next-token prediction



$$\log p(\bar{\mathbf{x}} | \hat{\mathbf{x}}) = \sum_{t=1}^T \text{mask}_t \log p(x_t | \hat{\mathbf{x}})$$

Reconstruct masked tokens



Why better language models are useful?



Language models can directly **encode knowledge** present in the training corpus.

The director of 2001: A Space Odyssey is \_\_\_\_\_



Language models can directly **encode knowledge** present in the training corpus.

Query	Answer	Generation
Francesco Bartolomeo Conti was born in ____.	Florence	Rome [-1.8] , <b>Florence</b> [-1.8] , Naples

Petroni et al. (2019), "Language Models as Knowledge Bases?" (ACL)



Language models can directly **encode knowledge** present in the training corpus.


Query	Answer	Generation
Francesco Bartolomeo Conti was born in ____.	Florence	Rome [-1.8] , <b>Florence</b> [-1.8] , Naples
Adolphe Adam died in ____.	Paris	<b>Paris</b> [-0.5] , London [-3.5] , Vienna
English bulldog is a subclass of ____.	dog	dogs [-0.3] , breeds [-2.2] , <b>dog</b>
The official language of Mauritius is ____.	English	<b>English</b> [-0.6] , French [-0.9] , Arabic
Patrick Oboya plays in ____ position.	midfielder	centre [-2.0] , center [-2.2] , <b>midfielder</b>
Hamburg Airport is named after ____.	Hamburg	Hess [-7.0] , Hermann [-7.1] , Schmidt

Petroni et al. (2019), "Language Models as Knowledge Bases?" (ACL)





Language models can be a **foundation** for various **tasks** across **different modalities**



Language models are stochastic parrots



Bender et al. (2021), "On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?"